

Agenda

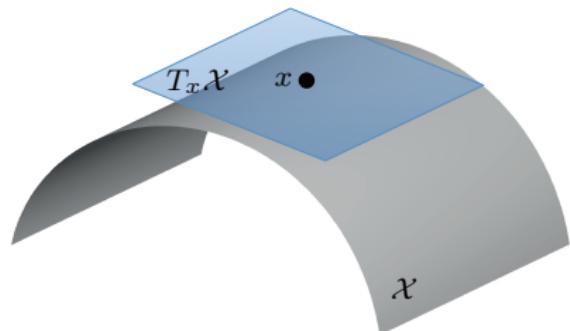
- Deep learning on regular structures
 - Multi-view representation
 - Volumetric Representation
- **Intrinsic deep learning on manifolds**
 - Spectral methods
 - Spatial methods
 - Embedding-based methods
- Deep learning on point cloud and parametric models

Extrinsic vs Intrinsic CNNs

Intrinsic

Riemannian geometry in one minute

- Manifold \mathcal{X} = topological space
- No global Euclidean structure
- **Tangent plane** $T_x\mathcal{X}$ = local Euclidean representation of manifold \mathcal{X} around x

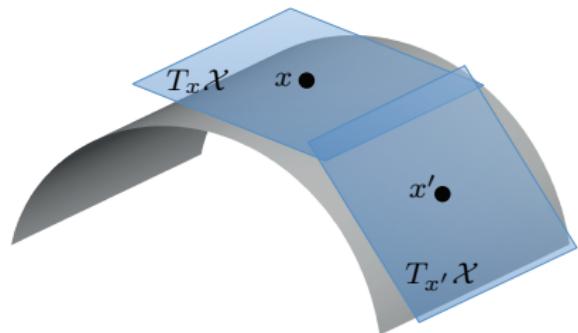


Riemannian geometry in one minute

- Manifold \mathcal{X} = topological space
- No global Euclidean structure
- Tangent plane $T_x \mathcal{X}$ = local Euclidean representation of manifold \mathcal{X} around x
- Riemannian metric

$$\langle \cdot, \cdot \rangle_{T_x \mathcal{X}} : T_x \mathcal{X} \times T_x \mathcal{X} \rightarrow \mathbb{R}$$

depending smoothly on x



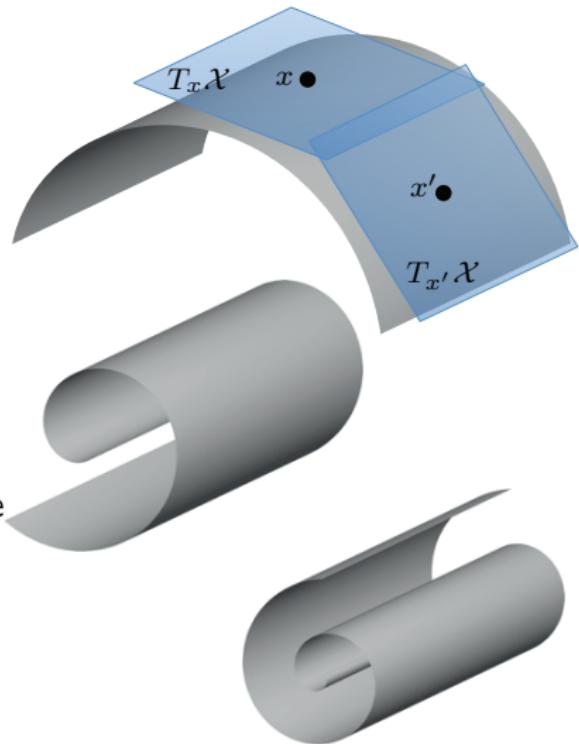
Riemannian geometry in one minute

- Manifold \mathcal{X} = topological space
- No global Euclidean structure
- Tangent plane $T_x \mathcal{X}$ = local Euclidean representation of manifold \mathcal{X} around x
- Riemannian metric

$$\langle \cdot, \cdot \rangle_{T_x \mathcal{X}} : T_x \mathcal{X} \times T_x \mathcal{X} \rightarrow \mathbb{R}$$

depending smoothly on x

Isometry = metric-preserving shape deformation



Riemannian geometry in one minute

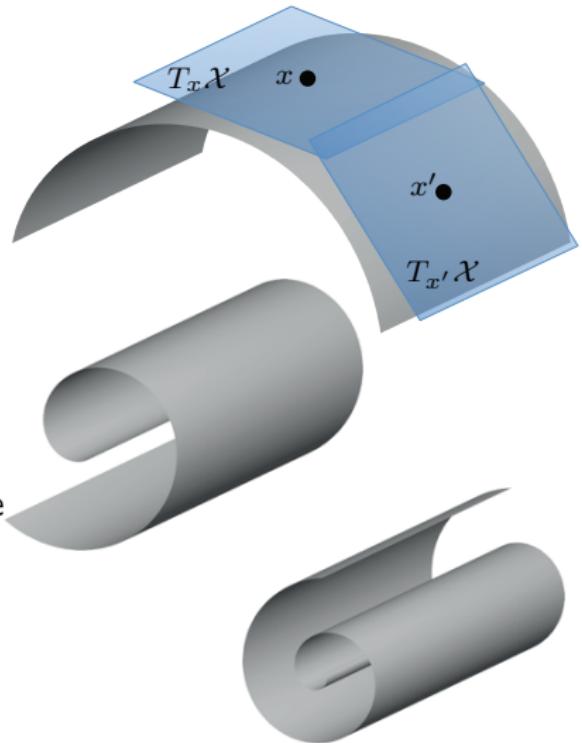
- Manifold \mathcal{X} = topological space
- No global Euclidean structure
- Tangent plane $T_x \mathcal{X}$ = local Euclidean representation of manifold \mathcal{X} around x
- Riemannian metric

$$\langle \cdot, \cdot \rangle_{T_x \mathcal{X}} : T_x \mathcal{X} \times T_x \mathcal{X} \rightarrow \mathbb{R}$$

depending smoothly on x

Isometry = metric-preserving shape deformation

Intrinsic = expressed solely in terms of the Riemannian metric



Riemannian geometry in one minute

- Manifold \mathcal{X} = topological space
- No global Euclidean structure
- Tangent plane $T_x \mathcal{X}$ = local Euclidean representation of manifold \mathcal{X} around x
- Riemannian metric

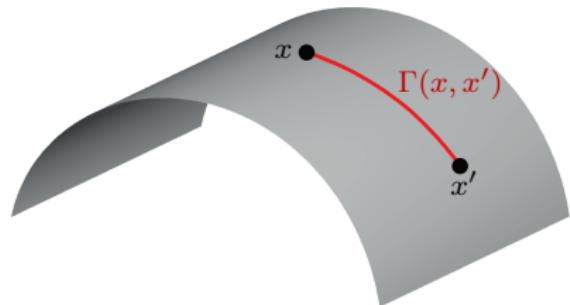
$$\langle \cdot, \cdot \rangle_{T_x \mathcal{X}} : T_x \mathcal{X} \times T_x \mathcal{X} \rightarrow \mathbb{R}$$

depending smoothly on x

Isometry = metric-preserving shape deformation

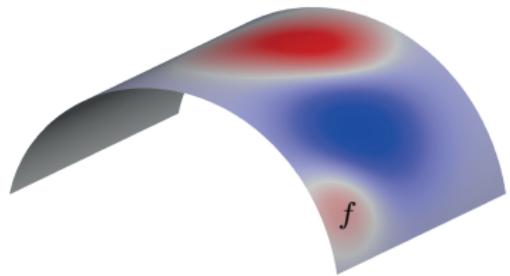
Intrinsic = expressed solely in terms of the Riemannian metric

- Geodesic = shortest path on \mathcal{X} between x and x'



Calculus on manifolds: scalar fields

Scalar field $f : \mathcal{X} \rightarrow \mathbb{R}$

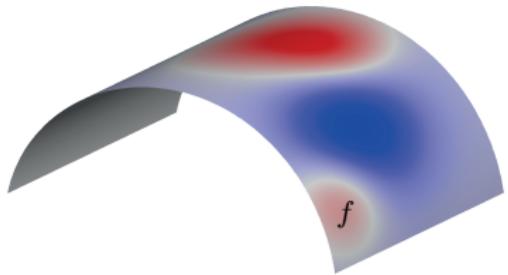


Calculus on manifolds: scalar fields

Scalar field $f : \mathcal{X} \rightarrow \mathbb{R}$

Hilbert space $L^2(\mathcal{X})$ with inner product

$$\langle f, g \rangle_{L^2(\mathcal{X})} = \int_{\mathcal{X}} f(x)g(x)dx$$



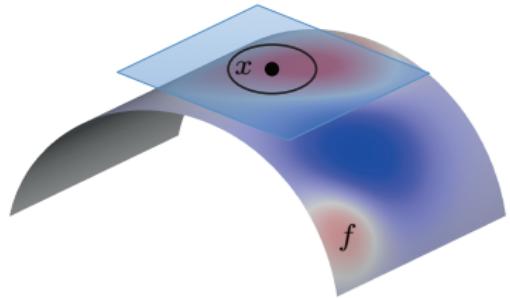
where dx = area element induced by the Riemannian metric

Calculus on manifolds: Laplacian operator

Laplacian $\Delta: L^2(\mathcal{X}) \rightarrow L^2(\mathcal{X})$

$$\Delta f = -\operatorname{div}(\nabla f)$$

“difference between $f(x)$ and average value of f around x ”

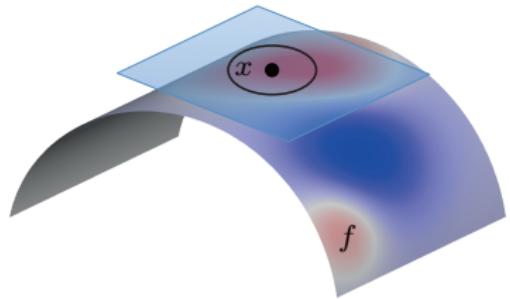


Calculus on manifolds: Laplacian operator

Laplacian $\Delta: L^2(\mathcal{X}) \rightarrow L^2(\mathcal{X})$

$$\Delta f = -\operatorname{div}(\nabla f)$$

“difference between $f(x)$ and average value of f around x ”



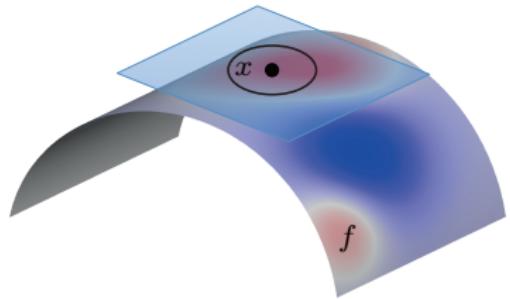
- Intrinsic (expressed solely in terms of the Riemannian metric)

Calculus on manifolds: Laplacian operator

Laplacian $\Delta: L^2(\mathcal{X}) \rightarrow L^2(\mathcal{X})$

$$\Delta f = -\operatorname{div}(\nabla f)$$

“difference between $f(x)$ and average value of f around x ”



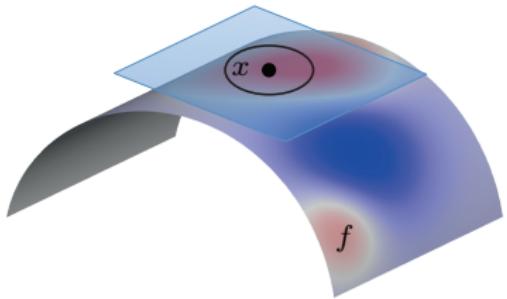
- Intrinsic (expressed solely in terms of the Riemannian metric)
- Isometry-invariant

Calculus on manifolds: Laplacian operator

Laplacian $\Delta: L^2(\mathcal{X}) \rightarrow L^2(\mathcal{X})$

$$\Delta f = -\operatorname{div}(\nabla f)$$

“difference between $f(x)$ and average value of f around x ”



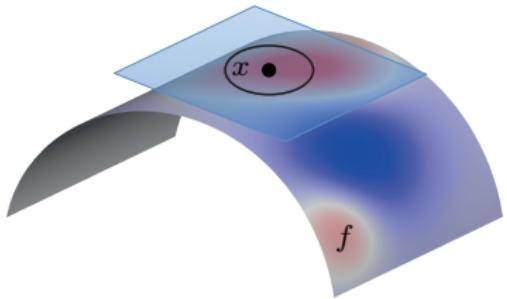
- Intrinsic (expressed solely in terms of the Riemannian metric)
- Isometry-invariant
- Self-adjoint $\langle \Delta f, g \rangle_{L^2(\mathcal{X})} = \langle f, \Delta g \rangle_{L^2(\mathcal{X})}$

Calculus on manifolds: Laplacian operator

Laplacian $\Delta: L^2(\mathcal{X}) \rightarrow L^2(\mathcal{X})$

$$\Delta f = -\operatorname{div}(\nabla f)$$

“difference between $f(x)$ and average value of f around x ”



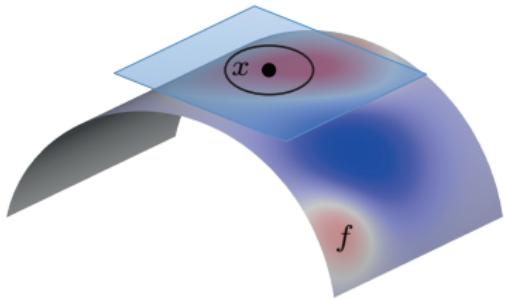
- Intrinsic (expressed solely in terms of the Riemannian metric)
- Isometry-invariant
- Self-adjoint $\langle \Delta f, g \rangle_{L^2(\mathcal{X})} = \langle f, \Delta g \rangle_{L^2(\mathcal{X})} \Rightarrow$ orthogonal eigenfunctions

Calculus on manifolds: Laplacian operator

Laplacian $\Delta: L^2(\mathcal{X}) \rightarrow L^2(\mathcal{X})$

$$\Delta f = -\operatorname{div}(\nabla f)$$

“difference between $f(x)$ and average value of f around x ”



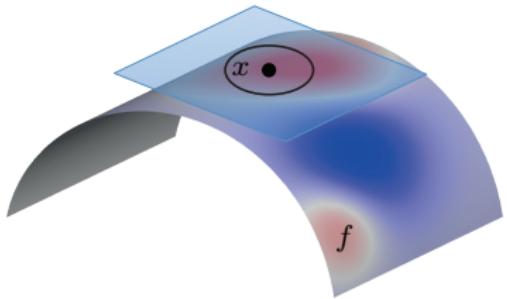
- Intrinsic (expressed solely in terms of the Riemannian metric)
- Isometry-invariant
- Self-adjoint $\langle \Delta f, g \rangle_{L^2(\mathcal{X})} = \langle f, \Delta g \rangle_{L^2(\mathcal{X})} \Rightarrow$ orthogonal eigenfunctions
- Positive semidefinite

Calculus on manifolds: Laplacian operator

Laplacian $\Delta: L^2(\mathcal{X}) \rightarrow L^2(\mathcal{X})$

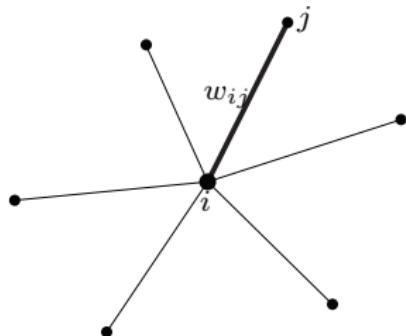
$$\Delta f = -\operatorname{div}(\nabla f)$$

“difference between $f(x)$ and average value of f around x ”



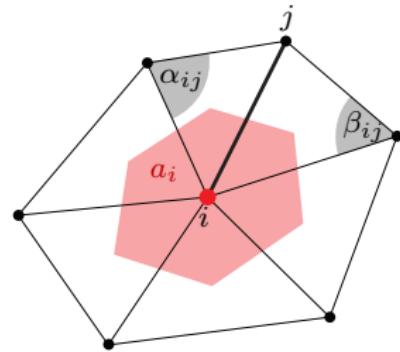
- Intrinsic (expressed solely in terms of the Riemannian metric)
- Isometry-invariant
- Self-adjoint $\langle \Delta f, g \rangle_{L^2(\mathcal{X})} = \langle f, \Delta g \rangle_{L^2(\mathcal{X})} \Rightarrow$ orthogonal eigenfunctions
- Positive semidefinite \Rightarrow non-negative eigenvalues

Discrete Laplacian



Undirected graph $(\mathcal{V}, \mathcal{E})$

$$(\Delta f)_i \approx \sum_{(i,j) \in \mathcal{E}} w_{ij} (f_i - f_j)$$

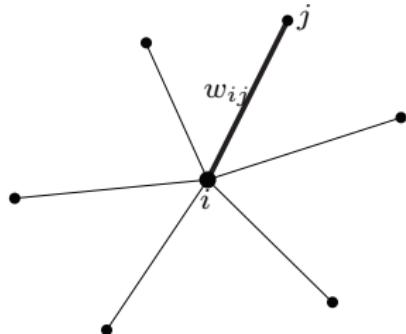


Triangular mesh $(\mathcal{V}, \mathcal{E}, \mathcal{F})$

$$(\Delta f)_i \approx \frac{1}{a_i} \sum_{(i,j) \in \mathcal{E}} \frac{\cot \alpha_{ij} + \cot \beta_{ij}}{2} (f_i - f_j)$$

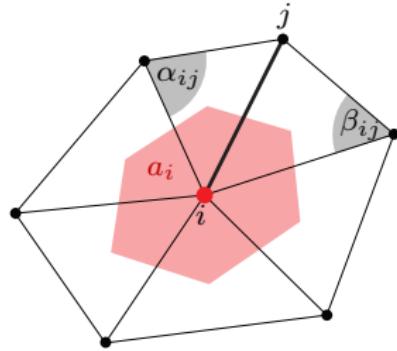
a_i = local area element

Discrete Laplacian



Undirected graph $(\mathcal{V}, \mathcal{E})$

$$(\Delta f)_i \approx \sum_{(i,j) \in \mathcal{E}} w_{ij} (f_i - f_j)$$



Triangular mesh $(\mathcal{V}, \mathcal{E}, \mathcal{F})$

$$(\Delta f)_i \approx \frac{1}{a_i} \sum_{(i,j) \in \mathcal{E}} \frac{\cot \alpha_{ij} + \cot \beta_{ij}}{2} (f_i - f_j)$$

a_i = local area element

In matrix-vector notation

$$\Delta \mathbf{f} = \mathbf{A}^{-1} (\mathbf{D} - \mathbf{W}) \mathbf{f}$$

where $\mathbf{f} = (f_1, \dots, f_n)^\top$, \mathbf{W} is the stiffness matrix, $\mathbf{A} = \text{diag}(a_1, \dots, a_n)$ is the mass matrix, and $\mathbf{D} = \text{diag}(\sum_{j \neq 1} w_{1j}, \dots, \sum_{j \neq n} w_{nj})$

Laplacian eigenfunctions and eigenvalues

Laplacian Δ of a **compact manifold** \mathcal{X} has countably many eigenfunctions

$$\Delta\phi_i(x) = \lambda_i\phi_i(x), \quad i = 1, 2, \dots$$

Laplacian eigenfunctions and eigenvalues

Laplacian Δ of a **compact manifold** \mathcal{X} has countably many eigenfunctions

$$\Delta\phi_i(x) = \lambda_i\phi_i(x), \quad i = 1, 2, \dots$$

- Eigenfunctions are **real** and **orthonormal** $\langle \phi_i, \phi_j \rangle_{L^2(\mathcal{X})} = \delta_{ij}$

Laplacian eigenfunctions and eigenvalues

Laplacian Δ of a **compact manifold** \mathcal{X} has countably many eigenfunctions

$$\Delta\phi_i(x) = \lambda_i\phi_i(x), \quad i = 1, 2, \dots$$

- Eigenfunctions are **real** and **orthonormal** $\langle \phi_i, \phi_j \rangle_{L^2(\mathcal{X})} = \delta_{ij}$
- Eigenvalues are **non-negative** $0 = \lambda_1 \leq \lambda_2 \leq \dots$

Laplacian eigenfunctions and eigenvalues

Laplacian Δ of a **compact manifold \mathcal{X}** has countably many eigenfunctions

$$\Delta\phi_i(x) = \lambda_i\phi_i(x), \quad i = 1, 2, \dots$$

- Eigenfunctions are **real** and **orthonormal** $\langle \phi_i, \phi_j \rangle_{L^2(\mathcal{X})} = \delta_{ij}$
- Eigenvalues are **non-negative** $0 = \lambda_1 \leq \lambda_2 \leq \dots$

Eigendecomposition of a discrete Laplacian matrix Δ

$$\Delta\Phi = \Phi\Lambda$$

where $\Lambda = \text{diag}(\lambda_1, \dots, \lambda_n)$ is a diagonal matrix of eigenvalues and $\Phi = (\phi_1, \dots, \phi_n)$ is a matrix of eigenvectors

Laplacian eigenfunctions and eigenvalues

Laplacian Δ of a **compact manifold** \mathcal{X} has countably many eigenfunctions

$$\Delta\phi_i(x) = \lambda_i\phi_i(x), \quad i = 1, 2, \dots$$

- Eigenfunctions are **real** and **orthonormal** $\langle \phi_i, \phi_j \rangle_{L^2(\mathcal{X})} = \delta_{ij}$
- Eigenvalues are **non-negative** $0 = \lambda_1 \leq \lambda_2 \leq \dots$

Eigendecomposition of a discrete Laplacian matrix Δ

$$\mathbf{A}^{-1}(\mathbf{D} - \mathbf{W})\Phi = \Phi\Lambda$$

where $\Lambda = \text{diag}(\lambda_1, \dots, \lambda_n)$ is a diagonal matrix of eigenvalues and $\Phi = (\phi_1, \dots, \phi_n)$ is a matrix of eigenvectors

Laplacian eigenfunctions and eigenvalues

Laplacian Δ of a **compact manifold** \mathcal{X} has countably many eigenfunctions

$$\Delta\phi_i(x) = \lambda_i\phi_i(x), \quad i = 1, 2, \dots$$

- Eigenfunctions are **real** and **orthonormal** $\langle \phi_i, \phi_j \rangle_{L^2(\mathcal{X})} = \delta_{ij}$
- Eigenvalues are **non-negative** $0 = \lambda_1 \leq \lambda_2 \leq \dots$

Eigendecomposition of a discrete Laplacian matrix Δ

$$(\mathbf{D} - \mathbf{W})\Phi = \mathbf{A}\Phi\Lambda$$

where $\Lambda = \text{diag}(\lambda_1, \dots, \lambda_n)$ is a diagonal matrix of eigenvalues and
 $\Phi = (\phi_1, \dots, \phi_n)$ is an \mathbf{A} -orthonormal matrix of eigenvectors
($\Phi^\top \mathbf{A} \Phi = \mathbf{I}$)

Laplacian eigenfunctions and eigenvalues

Laplacian Δ of a **compact manifold** \mathcal{X} has countably many eigenfunctions

$$\Delta\phi_i(x) = \lambda_i\phi_i(x), \quad i = 1, 2, \dots$$

- Eigenfunctions are **real** and **orthonormal** $\langle\phi_i, \phi_j\rangle_{L^2(\mathcal{X})} = \delta_{ij}$
- Eigenvalues are **non-negative** $0 = \lambda_1 \leq \lambda_2 \leq \dots$

Eigendecomposition of a discrete Laplacian matrix Δ

$$\mathbf{A}^{-1/2}(\mathbf{D} - \mathbf{W})\mathbf{A}^{-1/2}\mathbf{A}^{1/2}\Phi = \mathbf{A}^{1/2}\Phi\Lambda$$

where $\Lambda = \text{diag}(\lambda_1, \dots, \lambda_n)$ is a diagonal matrix of eigenvalues and $\Phi = (\phi_1, \dots, \phi_n)$ is an \mathbf{A} -orthonormal matrix of eigenvectors
 $(\Phi^\top \mathbf{A} \Phi = \mathbf{I})$

Laplacian eigenfunctions and eigenvalues

Laplacian Δ of a **compact manifold** \mathcal{X} has countably many eigenfunctions

$$\Delta\phi_i(x) = \lambda_i\phi_i(x), \quad i = 1, 2, \dots$$

- Eigenfunctions are **real** and **orthonormal** $\langle \phi_i, \phi_j \rangle_{L^2(\mathcal{X})} = \delta_{ij}$
- Eigenvalues are **non-negative** $0 = \lambda_1 \leq \lambda_2 \leq \dots$

Eigendecomposition of a discrete Laplacian matrix Δ

$$\mathbf{A}^{-1/2}(\mathbf{D} - \mathbf{W})\mathbf{A}^{-1/2}\Psi = \Psi\Lambda$$

where $\Lambda = \text{diag}(\lambda_1, \dots, \lambda_n)$ is a diagonal matrix of eigenvalues and $\Psi = (\psi_1, \dots, \psi_n)$ is an orthonormal matrix of eigenvectors ($\Psi^\top \Psi = \mathbf{I}$)

Laplacian eigenfunctions and eigenvalues

Laplacian Δ of a **compact manifold** \mathcal{X} has countably many eigenfunctions

$$\Delta\phi_i(x) = \lambda_i\phi_i(x), \quad i = 1, 2, \dots$$

- Eigenfunctions are **real** and **orthonormal** $\langle \phi_i, \phi_j \rangle_{L^2(\mathcal{X})} = \delta_{ij}$
- Eigenvalues are **non-negative** $0 = \lambda_1 \leq \lambda_2 \leq \dots$

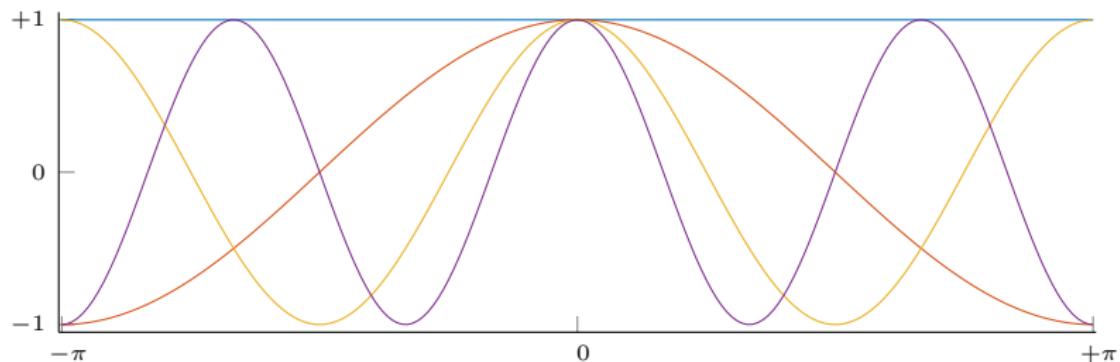
Eigendecomposition of a discrete Laplacian matrix Δ

$$\mathbf{A}^{-1/2}(\mathbf{D} - \mathbf{W})\mathbf{A}^{-1/2}\Psi = \Psi\Lambda$$

where $\Lambda = \text{diag}(\lambda_1, \dots, \lambda_n)$ is a diagonal matrix of eigenvalues and $\Psi = (\psi_1, \dots, \psi_n)$ is an orthonormal matrix of eigenvectors ($\Psi^\top \Psi = \mathbf{I}$)

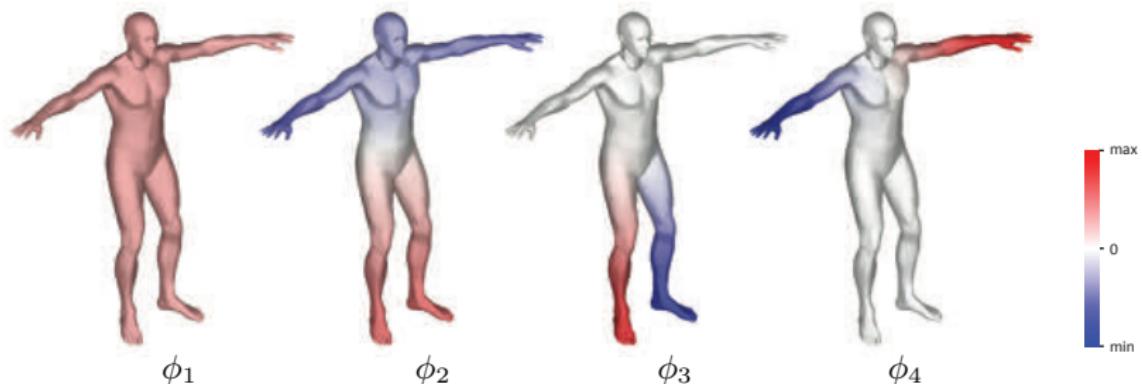
Laplacian eigenvectors = **smoothest orthonormal basis**

Laplacian eigenfunctions: Euclidean



First eigenfunctions of 1D Euclidean Laplacian = standard Fourier basis

Laplacian eigenfunctions: non-Euclidean

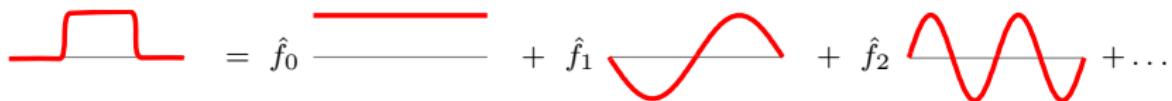


First eigenfunctions of a manifold Laplacian

Fourier analysis: Euclidean

A function $f : [-\pi, \pi] \rightarrow \mathbb{R}$ can be written as Fourier series

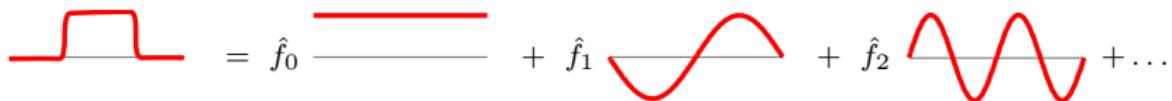
$$f(x) = \sum_{k \geq 0} \frac{1}{2\pi} \int_{-\pi}^{\pi} f(x') e^{-ikx'} dx' e^{ikx}$$



Fourier analysis: Euclidean

A function $f : [-\pi, \pi] \rightarrow \mathbb{R}$ can be written as Fourier series

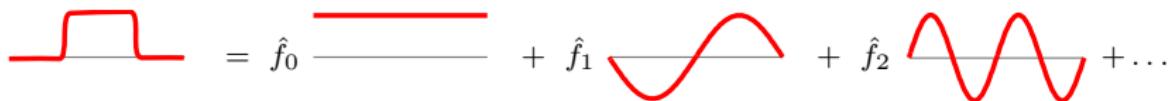
$$f(x) = \sum_{k \geq 0} \frac{1}{2\pi} \underbrace{\int_{-\pi}^{\pi} f(x') e^{-ikx'} dx'}_{\hat{f}_k = \langle f, e^{ikx} \rangle_{L^2([-\pi, \pi])}} e^{ikx}$$



Fourier analysis: Euclidean

A function $f : [-\pi, \pi] \rightarrow \mathbb{R}$ can be written as Fourier series

$$f(x) = \sum_{k \geq 0} \frac{1}{2\pi} \underbrace{\int_{-\pi}^{\pi} f(x') e^{-ikx'} dx'}_{\hat{f}_k = \langle f, e^{ikx} \rangle_{L^2([-\pi, \pi])}} e^{ikx}$$

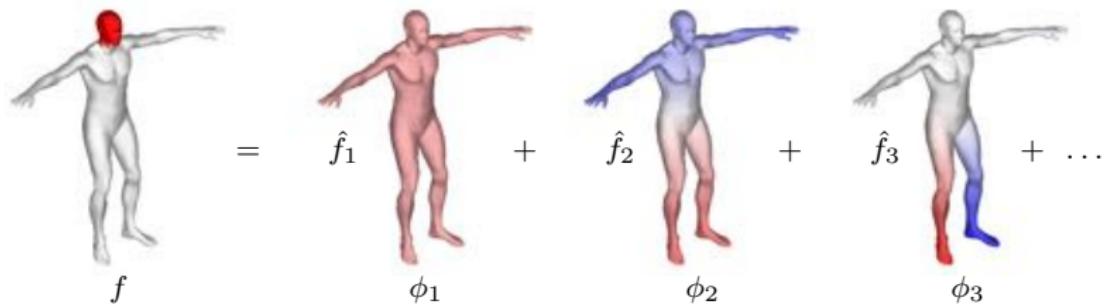


Fourier basis = Laplacian eigenfunctions: $-\frac{d^2}{dx^2} e^{ikx} = k^2 e^{ikx}$

Fourier analysis: non-Euclidean

A function $f : \mathcal{X} \rightarrow \mathbb{R}$ can be written as Fourier series

$$f(x) = \sum_{k \geq 1} \underbrace{\int_{\mathcal{X}} f(x') \phi_k(x') dx'}_{\hat{f}_k = \langle f, \phi_k \rangle_{L^2(\mathcal{X})}} \phi_k(x)$$



Fourier basis = Laplacian eigenfunctions: $\Delta \phi_k(x) = \lambda_k \phi_k(x)$

Physical application: heat equation

$$f_t = -c \Delta f$$

Newton's law of cooling: rate of change of the temperature of an object is proportional to the difference between its own temperature and the temperature of the surrounding

c [m²/sec] = **thermal diffusivity constant**

Heat diffusion on manifolds

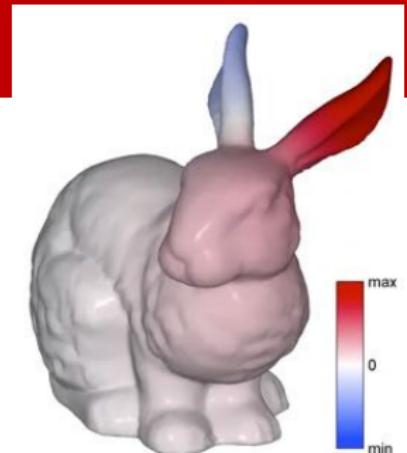
$$\begin{cases} f_t(x, t) = -\Delta f(x, t) \\ f(x, 0) = f_0(x) \end{cases}$$

- $f(x, t)$ = amount of heat at point x at time t
- $f_0(x)$ = initial heat distribution

Heat diffusion on manifolds

$$\begin{cases} f_t(x, t) = -\Delta f(x, t) \\ f(x, 0) = f_0(x) \end{cases}$$

- $f(x, t)$ = amount of heat at point x at time t
- $f_0(x)$ = initial heat distribution



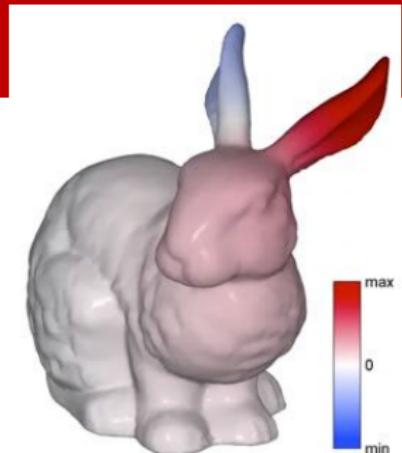
Solution of the heat equation expressed through the **heat operator**

$$f(x, t) = e^{-t\Delta} f_0(x)$$

Heat diffusion on manifolds

$$\begin{cases} f_t(x, t) = -\Delta f(x, t) \\ f(x, 0) = f_0(x) \end{cases}$$

- $f(x, t)$ = amount of heat at point x at time t
- $f_0(x)$ = initial heat distribution



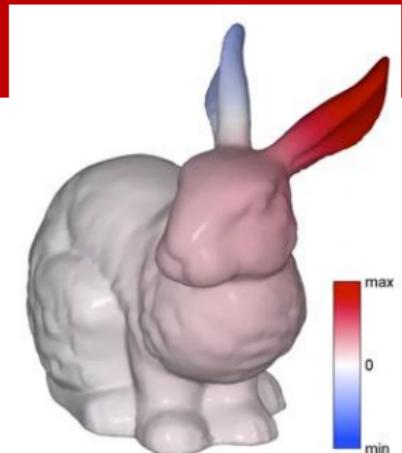
Solution of the heat equation expressed through the **heat operator**

$$f(x, t) = e^{-t\Delta} f_0(x) = \sum_{k \geq 1} \langle f_0, \phi_k \rangle_{L^2(\mathcal{X})} e^{-t\lambda_k} \phi_k(x)$$

Heat diffusion on manifolds

$$\begin{cases} f_t(x, t) = -\Delta f(x, t) \\ f(x, 0) = f_0(x) \end{cases}$$

- $f(x, t)$ = amount of heat at point x at time t
- $f_0(x)$ = initial heat distribution



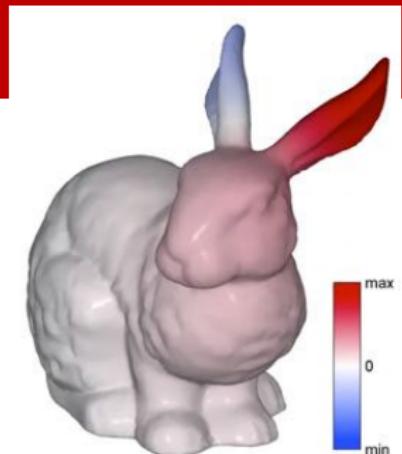
Solution of the heat equation expressed through the **heat operator**

$$\begin{aligned} f(x, t) &= e^{-t\Delta} f_0(x) = \sum_{k \geq 1} \langle f_0, \phi_k \rangle_{L^2(\mathcal{X})} e^{-t\lambda_k} \phi_k(x) \\ &= \int_{\mathcal{X}} f_0(x') \sum_{k \geq 1} e^{-t\lambda_k} \phi_k(x) \phi_k(x') dx' \end{aligned}$$

Heat diffusion on manifolds

$$\begin{cases} f_t(x, t) = -\Delta f(x, t) \\ f(x, 0) = f_0(x) \end{cases}$$

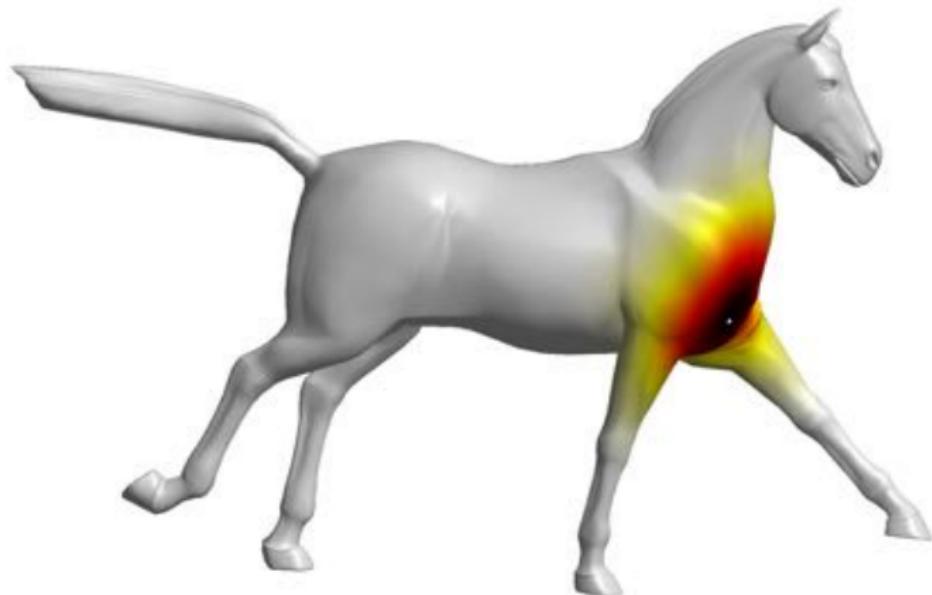
- $f(x, t)$ = amount of heat at point x at time t
- $f_0(x)$ = initial heat distribution



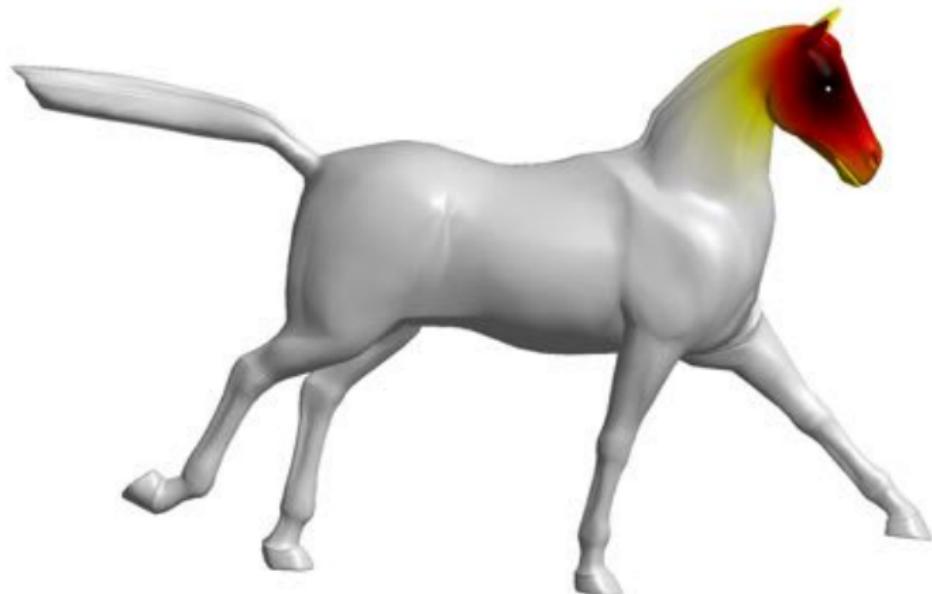
Solution of the heat equation expressed through the **heat operator**

$$\begin{aligned} f(x, t) &= e^{-t\Delta} f_0(x) = \sum_{k \geq 1} \langle f_0, \phi_k \rangle_{L^2(\mathcal{X})} e^{-t\lambda_k} \phi_k(x) \\ &= \int_{\mathcal{X}} f_0(x') \underbrace{\sum_{k \geq 1} e^{-t\lambda_k} \phi_k(x) \phi_k(x')}_{\text{heat kernel } h_t(x, x')} dx' \end{aligned}$$

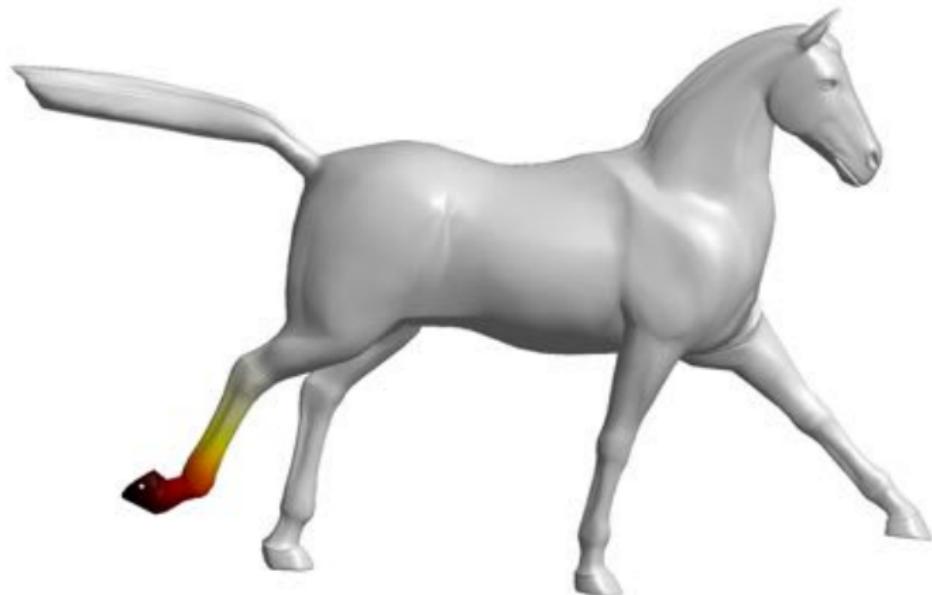
Heat kernels



Heat kernels



Heat kernels



Heat kernels

Interpretation of the heat kernel

Solution of the heat equation on a manifold \mathcal{X} expressed in the Laplacian eigenbasis $\Delta\phi_k(x) = \lambda_k\phi_k(x)$

$$f(x, t) = \int_{\mathcal{X}} f_0(x') \underbrace{\sum_{k \geq 1} e^{-t\lambda_k} \phi_k(x)\phi_k(x')}_{\text{heat kernel } h_t(x, x')}$$

Interpretation of the heat kernel

Solution of the heat equation on a **Euclidean space** $[-\pi, \pi]$ expressed in the Laplacian eigenbasis $-\frac{d^2}{dx^2} e^{ikx} = k^2 e^{ikx}$

$$f(x, t) = \frac{1}{2\pi} \int_{-\pi}^{\pi} f_0(x') \underbrace{\sum_{k \geq 0} e^{-tk^2} e^{ikx} e^{-ikx'} dx'}_{\text{heat kernel } h_t(x, x')}$$

Interpretation of the heat kernel

Solution of the heat equation on a **Euclidean space** $[-\pi, \pi]$ expressed in the Laplacian eigenbasis $-\frac{d^2}{dx^2}e^{ikx} = k^2 e^{ikx}$

$$f(x, t) = \frac{1}{2\pi} \int_{-\pi}^{\pi} f_0(x') \underbrace{\sum_{k \geq 0} e^{-tk^2} e^{ik(x-x')}}_{\text{heat kernel } h_t(x-x')}$$

Interpretation of the heat kernel

Solution of the heat equation on a Euclidean space $[-\pi, \pi]$ expressed in the Laplacian eigenbasis $-\frac{d^2}{dx^2} e^{ikx} = k^2 e^{ikx}$

$$\begin{aligned} f(x, t) &= \frac{1}{2\pi} \int_{-\pi}^{\pi} f_0(x') \underbrace{\sum_{k \geq 0} e^{-tk^2} e^{ik(x-x')}}_{\text{heat kernel } h_t(x-x')} dx' \\ &= (f_0 \star h_t)(x) \end{aligned}$$

Heat kernel = impulse response

Convolution: Euclidean space

Given two functions $f, g : [-\pi, \pi] \rightarrow \mathbb{R}$ their **convolution** is a function

$$(f \star g)(x) = \int_{-\pi}^{\pi} f(x')g(x - x')dx'$$

Convolution: Euclidean space

Given two functions $f, g : [-\pi, \pi] \rightarrow \mathbb{R}$ their **convolution** is a function

$$(f \star g)(x) = \int_{-\pi}^{\pi} f(x')g(x - x')dx'$$

- **Shift-invariance:** $f(x - x_0) \star g(x) = (f \star g)(x - x_0)$

Convolution: Euclidean space

Given two functions $f, g : [-\pi, \pi] \rightarrow \mathbb{R}$ their **convolution** is a function

$$(f \star g)(x) = \int_{-\pi}^{\pi} f(x')g(x - x')dx'$$

- **Shift-invariance:** $f(x - x_0) \star g(x) = (f \star g)(x - x_0)$
- Convolution operator **commutes with Laplacian:** $(\Delta f) \star g = \Delta(f \star g)$

Convolution: Euclidean space

Given two functions $f, g : [-\pi, \pi] \rightarrow \mathbb{R}$ their **convolution** is a function

$$(f \star g)(x) = \int_{-\pi}^{\pi} f(x')g(x - x')dx'$$

- **Shift-invariance:** $f(x - x_0) \star g(x) = (f \star g)(x - x_0)$
- Convolution operator **commutes with Laplacian:** $(\Delta f) \star g = \Delta(f \star g)$
- **Convolution theorem:** Fourier transform diagonalizes the convolution operator

Convolution: Euclidean space

Given two functions $f, g : [-\pi, \pi] \rightarrow \mathbb{R}$ their **convolution** is a function

$$(f \star g)(x) = \int_{-\pi}^{\pi} f(x')g(x - x')dx'$$

- **Shift-invariance:** $f(x - x_0) \star g(x) = (f \star g)(x - x_0)$
- Convolution operator **commutes with Laplacian:** $(\Delta f) \star g = \Delta(f \star g)$
- **Convolution theorem:** Fourier transform diagonalizes the convolution operator \Rightarrow convolution can be computed in the Fourier domain as

$$\widehat{(f \star g)} = \hat{f} \cdot \hat{g}$$

Convolution: Euclidean space

Given two functions $f, g : [-\pi, \pi] \rightarrow \mathbb{R}$ their convolution is a function

$$(f \star g)(x) = \int_{-\pi}^{\pi} f(x')g(x - x')dx'$$

- Shift-invariance: $f(x - x_0) \star g(x) = (f \star g)(x - x_0)$
- Convolution operator commutes with Laplacian: $(\Delta f) \star g = \Delta(f \star g)$
- Convolution theorem: Fourier transform diagonalizes the convolution operator \Rightarrow convolution can be computed in the Fourier domain as

$$\widehat{(f \star g)} = \hat{f} \cdot \hat{g}$$

- Efficient computation using FFT

Convolution Theorem

Convolution of two vectors $\mathbf{f} = (f_1, \dots, f_n)^\top$ and $\mathbf{g} = (g_1, \dots, g_n)^\top$

$$\mathbf{f} * \mathbf{g} = \begin{bmatrix} g_1 & g_2 & \dots & \dots & g_n \\ g_n & g_1 & g_2 & \dots & g_{n-1} \\ \vdots & \vdots & \ddots & \ddots & \vdots \\ g_3 & g_4 & \dots & g_1 & g_2 \\ g_2 & g_3 & \dots & \dots & g_1 \end{bmatrix} \begin{bmatrix} f_1 \\ \vdots \\ f_n \end{bmatrix}$$

Convolution Theorem

Convolution of two vectors $\mathbf{f} = (f_1, \dots, f_n)^\top$ and $\mathbf{g} = (g_1, \dots, g_n)^\top$

$$\mathbf{f} * \mathbf{g} = \underbrace{\begin{bmatrix} g_1 & g_2 & \dots & \dots & g_n \\ g_n & g_1 & g_2 & \dots & g_{n-1} \\ \vdots & \vdots & \ddots & \ddots & \vdots \\ g_3 & g_4 & \dots & g_1 & g_2 \\ g_2 & g_3 & \dots & \dots & g_1 \end{bmatrix}}_{\text{circulant matrix}} \begin{bmatrix} f_1 \\ \vdots \\ f_n \end{bmatrix}$$

Convolution Theorem

Convolution of two vectors $\mathbf{f} = (f_1, \dots, f_n)^\top$ and $\mathbf{g} = (g_1, \dots, g_n)^\top$

$$\mathbf{f} * \mathbf{g} = \underbrace{\begin{bmatrix} g_1 & g_2 & \dots & \dots & g_n \\ g_n & g_1 & g_2 & \dots & g_{n-1} \\ \vdots & \vdots & \ddots & \ddots & \vdots \\ g_3 & g_4 & \dots & g_1 & g_2 \\ g_2 & g_3 & \dots & \dots & g_1 \end{bmatrix}}_{\text{diagonalized by Fourier basis}} \begin{bmatrix} f_1 \\ \vdots \\ f_n \end{bmatrix}$$

Convolution Theorem

Convolution of two vectors $\mathbf{f} = (f_1, \dots, f_n)^\top$ and $\mathbf{g} = (g_1, \dots, g_n)^\top$

$$\mathbf{f} * \mathbf{g} = \begin{bmatrix} g_1 & g_2 & \dots & \dots & g_n \\ g_n & g_1 & g_2 & \dots & g_{n-1} \\ \vdots & \vdots & \ddots & \ddots & \vdots \\ g_3 & g_4 & \dots & g_1 & g_2 \\ g_2 & g_3 & \dots & \dots & g_1 \end{bmatrix} \begin{bmatrix} f_1 \\ \vdots \\ f_n \end{bmatrix}$$

$$= \Phi \begin{bmatrix} \hat{g}_1 \\ \ddots \\ \hat{g}_n \end{bmatrix} \Phi^\top \mathbf{f}$$

Convolution Theorem

Convolution of two vectors $\mathbf{f} = (f_1, \dots, f_n)^\top$ and $\mathbf{g} = (g_1, \dots, g_n)^\top$

$$\mathbf{f} * \mathbf{g} = \begin{bmatrix} g_1 & g_2 & \dots & \dots & g_n \\ g_n & g_1 & g_2 & \dots & g_{n-1} \\ \vdots & \vdots & \ddots & \ddots & \vdots \\ g_3 & g_4 & \dots & g_1 & g_2 \\ g_2 & g_3 & \dots & \dots & g_1 \end{bmatrix} \begin{bmatrix} f_1 \\ \vdots \\ f_n \end{bmatrix}$$

$$= \Phi \begin{bmatrix} \hat{g}_1 \\ \ddots \\ \hat{g}_n \end{bmatrix} \begin{bmatrix} \hat{f}_1 \\ \vdots \\ \hat{f}_n \end{bmatrix}$$

Convolution Theorem

Convolution of two vectors $\mathbf{f} = (f_1, \dots, f_n)^\top$ and $\mathbf{g} = (g_1, \dots, g_n)^\top$

$$\mathbf{f} * \mathbf{g} = \begin{bmatrix} g_1 & g_2 & \dots & \dots & g_n \\ g_n & g_1 & g_2 & \dots & g_{n-1} \\ \vdots & \vdots & \ddots & \ddots & \vdots \\ g_3 & g_4 & \dots & g_1 & g_2 \\ g_2 & g_3 & \dots & \dots & g_1 \end{bmatrix} \begin{bmatrix} f_1 \\ \vdots \\ f_n \end{bmatrix}$$

$$= \Phi \begin{bmatrix} \hat{f}_1 \cdot \hat{g}_1 \\ \vdots \\ \hat{f}_n \cdot \hat{g}_n \end{bmatrix}$$

Spectral convolution

Spectral convolution of $f, g \in L^2(\mathcal{X})$ can be defined by analogy

$$f \star g = \sum_{k \geq 1} \langle f, \phi_k \rangle_{L^2(\mathcal{X})} \langle g, \phi_k \rangle_{L^2(\mathcal{X})} \phi_k$$

Spectral convolution

Spectral convolution of $f, g \in L^2(\mathcal{X})$ can be defined by analogy

$$f \star g = \sum_{k \geq 1} \underbrace{\langle f, \phi_k \rangle_{L^2(\mathcal{X})} \langle g, \phi_k \rangle_{L^2(\mathcal{X})}}_{\text{product in the Fourier domain}} \phi_k$$

Spectral convolution

Spectral convolution of $f, g \in L^2(\mathcal{X})$ can be defined by analogy

$$f \star g = \underbrace{\sum_{k \geq 1} \underbrace{\langle f, \phi_k \rangle_{L^2(\mathcal{X})} \langle g, \phi_k \rangle_{L^2(\mathcal{X})}}_{\text{product in the Fourier domain}} \phi_k}_{\text{inverse Fourier transform}}$$

Spectral convolution

Spectral convolution of $f, g \in L^2(\mathcal{X})$ can be defined by analogy

$$f \star g = \sum_{k \geq 1} \langle f, \phi_k \rangle_{L^2(\mathcal{X})} \langle g, \phi_k \rangle_{L^2(\mathcal{X})} \phi_k$$

In matrix-vector notation

$$\mathbf{f} \star \mathbf{g} = \Phi (\Phi^\top \mathbf{g}) \circ (\Phi^\top \mathbf{f})$$

Spectral convolution

Spectral convolution of $f, g \in L^2(\mathcal{X})$ can be defined by analogy

$$f \star g = \sum_{k \geq 1} \langle f, \phi_k \rangle_{L^2(\mathcal{X})} \langle g, \phi_k \rangle_{L^2(\mathcal{X})} \phi_k$$

In matrix-vector notation

$$\mathbf{f} \star \mathbf{g} = \Phi \operatorname{diag}(\hat{g}_1, \dots, \hat{g}_n) \Phi^\top \mathbf{f}$$

Spectral convolution

Spectral convolution of $f, g \in L^2(\mathcal{X})$ can be defined by analogy

$$f \star g = \sum_{k \geq 1} \langle f, \phi_k \rangle_{L^2(\mathcal{X})} \langle g, \phi_k \rangle_{L^2(\mathcal{X})} \phi_k$$

In matrix-vector notation

$$\mathbf{f} \star \mathbf{g} = \underbrace{\Phi \operatorname{diag}(\hat{g}_1, \dots, \hat{g}_n) \Phi^\top}_{\mathbf{G}} \mathbf{f}$$

Spectral convolution

Spectral convolution of $f, g \in L^2(\mathcal{X})$ can be defined by analogy

$$f \star g = \sum_{k \geq 1} \langle f, \phi_k \rangle_{L^2(\mathcal{X})} \langle g, \phi_k \rangle_{L^2(\mathcal{X})} \phi_k$$

In matrix-vector notation

$$\mathbf{f} \star \mathbf{g} = \underbrace{\Phi \operatorname{diag}(\hat{g}_1, \dots, \hat{g}_n) \Phi^\top}_{\mathbf{G}} \mathbf{f}$$

- Not shift-invariant! (\mathbf{G} has no circulant structure)

Spectral convolution

Spectral convolution of $f, g \in L^2(\mathcal{X})$ can be defined by analogy

$$f \star g = \sum_{k \geq 1} \langle f, \phi_k \rangle_{L^2(\mathcal{X})} \langle g, \phi_k \rangle_{L^2(\mathcal{X})} \phi_k$$

In matrix-vector notation

$$\mathbf{f} \star \mathbf{g} = \underbrace{\Phi \operatorname{diag}(\hat{g}_1, \dots, \hat{g}_n) \Phi^\top}_{\mathbf{G}} \mathbf{f}$$

- Not shift-invariant! (\mathbf{G} has no circulant structure)
- Filter coefficients depend on basis ϕ_1, \dots, ϕ_n

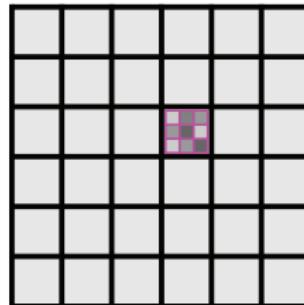
Different formulations of non-Euclidean CNNs



Spectral domain



Spatial domain



Embedding domain

Spectral CNN

Convolutional layer expressed in the **spectral domain**

$$\mathbf{g}_l = \xi \left(\sum_{l'=1}^p \mathbf{\Phi} \hat{\mathbf{W}}_{l,l'} \mathbf{\Phi}^\top \mathbf{f}_{l'} \right) \quad l = 1, \dots, q \\ l' = 1, \dots, p$$

where $\hat{\mathbf{W}}_{l,l} = n \times n$ diagonal matrix of filter coefficients

Spectral CNN

Convolutional layer expressed in the spectral domain

$$\mathbf{g}_l = \xi \left(\sum_{l'=1}^p \mathbf{\Phi} \hat{\mathbf{W}}_{l,l'} \mathbf{\Phi}^\top \mathbf{f}_{l'} \right) \quad l = 1, \dots, q \\ l' = 1, \dots, p$$

where $\hat{\mathbf{W}}_{l,l} = n \times n$ diagonal matrix of filter coefficients

⌚ $\mathcal{O}(n)$ parameters per layer

Spectral CNN

Convolutional layer expressed in the spectral domain

$$\mathbf{g}_l = \xi \left(\sum_{l'=1}^p \mathbf{\Phi} \hat{\mathbf{W}}_{l,l'} \mathbf{\Phi}^\top \mathbf{f}_{l'} \right) \quad l = 1, \dots, q \\ l' = 1, \dots, p$$

where $\hat{\mathbf{W}}_{l,l} = n \times n$ diagonal matrix of filter coefficients

- ⌚ $\mathcal{O}(n)$ parameters per layer
- ⌚ $\mathcal{O}(n^2)$ computation of forward and inverse Fourier transforms $\mathbf{\Phi}^\top, \mathbf{\Phi}$ (no FFT on manifolds or graphs)

Spectral CNN

Convolutional layer expressed in the spectral domain

$$\mathbf{g}_l = \xi \left(\sum_{l'=1}^p \mathbf{\Phi} \hat{\mathbf{W}}_{l,l'} \mathbf{\Phi}^\top \mathbf{f}_{l'} \right) \quad l = 1, \dots, q \\ l' = 1, \dots, p$$

where $\hat{\mathbf{W}}_{l,l} = n \times n$ diagonal matrix of filter coefficients

- ⌚ $\mathcal{O}(n)$ parameters per layer
- ⌚ $\mathcal{O}(n^2)$ computation of forward and inverse Fourier transforms $\mathbf{\Phi}^\top, \mathbf{\Phi}$ (no FFT on manifolds or graphs)
- ⌚ No guarantee of spatial localization of filters

Spectral CNN

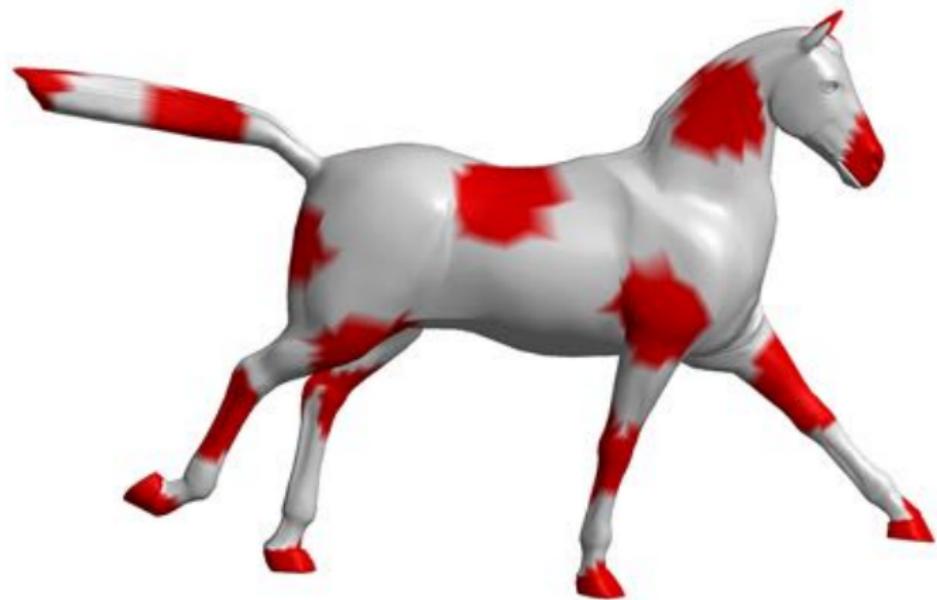
Convolutional layer expressed in the spectral domain

$$\mathbf{g}_l = \xi \left(\sum_{l'=1}^p \mathbf{\Phi} \hat{\mathbf{W}}_{l,l'} \mathbf{\Phi}^\top \mathbf{f}_{l'} \right) \quad l = 1, \dots, q \\ l' = 1, \dots, p$$

where $\hat{\mathbf{W}}_{l,l} = n \times n$ diagonal matrix of filter coefficients

- ⌚ $\mathcal{O}(n)$ parameters per layer
- ⌚ $\mathcal{O}(n^2)$ computation of forward and inverse Fourier transforms $\mathbf{\Phi}^\top, \mathbf{\Phi}$ (no FFT on manifolds or graphs)
- ⌚ No guarantee of spatial localization of filters
- ⌚ Filters are basis-dependent \Rightarrow does not generalize across domains

Basis dependence



Function f

Basis dependence



'Edge detecting' spectral filter $\Phi \hat{W} \Phi^\top f$

Basis dependence

Basis dependence

Localization and Smoothness

In the Euclidean setting (by Parseval's identity)

$$\int_{-\infty}^{+\infty} |x|^{2k} |f(x)|^2 dx = \int_{-\infty}^{+\infty} \left| \frac{\partial^k \hat{f}(\omega)}{\partial \omega^k} \right|^2 d\omega$$

Localization and Smoothness

In the Euclidean setting (by Parseval's identity)

$$\int_{-\infty}^{+\infty} |x|^{2k} |f(x)|^2 dx = \int_{-\infty}^{+\infty} \left| \frac{\partial^k \hat{f}(\omega)}{\partial \omega^k} \right|^2 d\omega$$

⇒ Localization in space = smoothness in frequency domain

Localization and Smoothness

In the Euclidean setting (by Parseval's identity)

$$\int_{-\infty}^{+\infty} |x|^{2k} |f(x)|^2 dx = \int_{-\infty}^{+\infty} \left| \frac{\partial^k \hat{f}(\omega)}{\partial \omega^k} \right|^2 d\omega$$

⇒ Localization in space = smoothness in frequency domain

Parametrize the filter using a smooth spectral transfer function $\tau(\lambda)$

Localization and Smoothness

In the Euclidean setting (by Parseval's identity)

$$\int_{-\infty}^{+\infty} |x|^{2k} |f(x)|^2 dx = \int_{-\infty}^{+\infty} \left| \frac{\partial^k \hat{f}(\omega)}{\partial \omega^k} \right|^2 d\omega$$

⇒ Localization in space = smoothness in frequency domain

Parametrize the filter using a smooth spectral transfer function $\tau(\lambda)$

Application of the filter

$$\tau(\Delta)\mathbf{f} = \Phi \tau(\Lambda) \Phi^\top \mathbf{f}$$

Localization and Smoothness

In the Euclidean setting (by Parseval's identity)

$$\int_{-\infty}^{+\infty} |x|^{2k} |f(x)|^2 dx = \int_{-\infty}^{+\infty} \left| \frac{\partial^k \hat{f}(\omega)}{\partial \omega^k} \right|^2 d\omega$$

⇒ Localization in space = smoothness in frequency domain

Parametrize the filter using a smooth spectral transfer function $\tau(\lambda)$

Application of the filter

$$\tau(\Delta)\mathbf{f} = \Phi \begin{pmatrix} \tau(\lambda_1) & & \\ & \ddots & \\ & & \tau(\lambda_n) \end{pmatrix} \Phi^\top \mathbf{f}$$

Localization and Smoothness

In the Euclidean setting (by Parseval's identity)

$$\int_{-\infty}^{+\infty} |x|^{2k} |f(x)|^2 dx = \int_{-\infty}^{+\infty} \left| \frac{\partial^k \hat{f}(\omega)}{\partial \omega^k} \right|^2 d\omega$$

⇒ Localization in space = smoothness in frequency domain

Parametrize the filter using a smooth spectral transfer function $\tau(\lambda)$

Application of the parametric filter with learnable parameters α

$$\tau_\alpha(\Delta) \mathbf{f} = \Phi \begin{pmatrix} \tau_\alpha(\lambda_1) & & \\ & \ddots & \\ & & \tau_\alpha(\lambda_n) \end{pmatrix} \Phi^\top \mathbf{f}$$

Spectral CNN with polynomial filters

Represent spectral transfer function as a **polynomial** or order r

$$\tau_{\alpha}(\lambda) = \sum_{j=0}^r \alpha_j \lambda^j$$

where $\alpha = (\alpha_0, \dots, \alpha_r)^{\top}$ is the vector of filter parameters

Spectral CNN with polynomial filters

Represent spectral transfer function as a **polynomial** or order r

$$\tau_{\alpha}(\lambda) = \sum_{j=0}^r \alpha_j \lambda^j$$

where $\alpha = (\alpha_0, \dots, \alpha_r)^{\top}$ is the vector of filter parameters

☺ $\mathcal{O}(1)$ parameters per layer

Spectral CNN with polynomial filters

Represent spectral transfer function as a **polynomial** or order r

$$\tau_{\alpha}(\lambda) = \sum_{j=0}^r \alpha_j \lambda^j$$

where $\alpha = (\alpha_0, \dots, \alpha_r)^{\top}$ is the vector of filter parameters

- ☺ $\mathcal{O}(1)$ parameters per layer
- ☺ Filters have guaranteed r -hops support

Spectral CNN with polynomial filters

Represent spectral transfer function as a **polynomial** or order r

$$\tau_{\alpha}(\Delta) = \sum_{j=0}^r \alpha_j \Delta^j$$

where $\alpha = (\alpha_0, \dots, \alpha_r)^{\top}$ is the vector of filter parameters

- ☺ $\mathcal{O}(1)$ parameters per layer
- ☺ Filters have guaranteed r -hops support
- ☺ No explicit computation of $\Phi^{\top}, \Phi \Rightarrow \mathcal{O}(nr)$ complexity

Spectral CNN with polynomial filters

Represent spectral transfer function as a **polynomial** or order r

$$\tau_{\alpha}(\Delta) = \sum_{j=0}^r \alpha_j \Delta^j$$

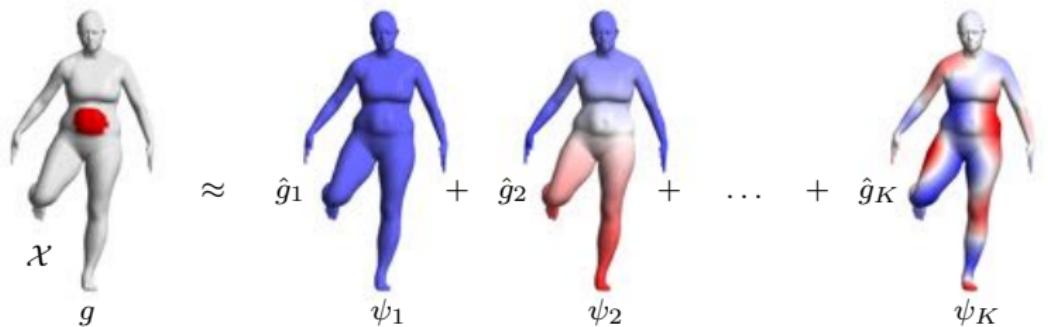
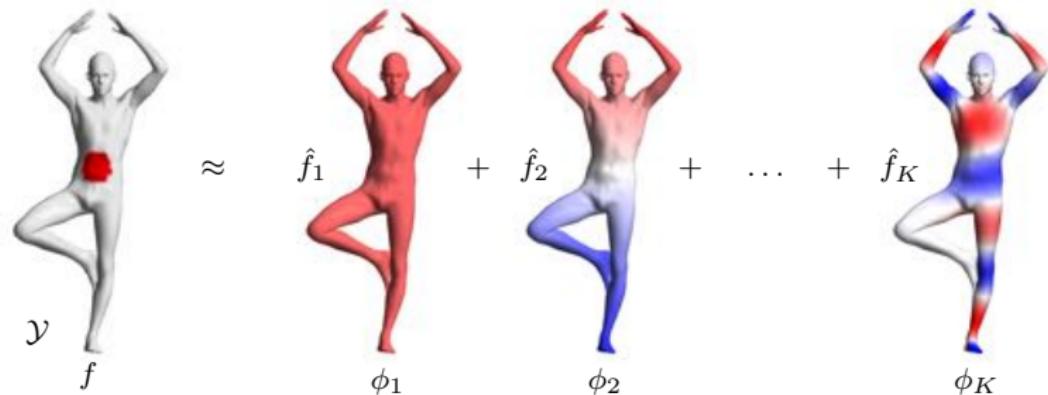
where $\alpha = (\alpha_0, \dots, \alpha_r)^{\top}$ is the vector of filter parameters

- ☺ $\mathcal{O}(1)$ parameters per layer
- ☺ Filters have guaranteed r -hops support
- ☺ No explicit computation of $\Phi^{\top}, \Phi \Rightarrow \mathcal{O}(nr)$ complexity
- ☹ Does not generalize across domains

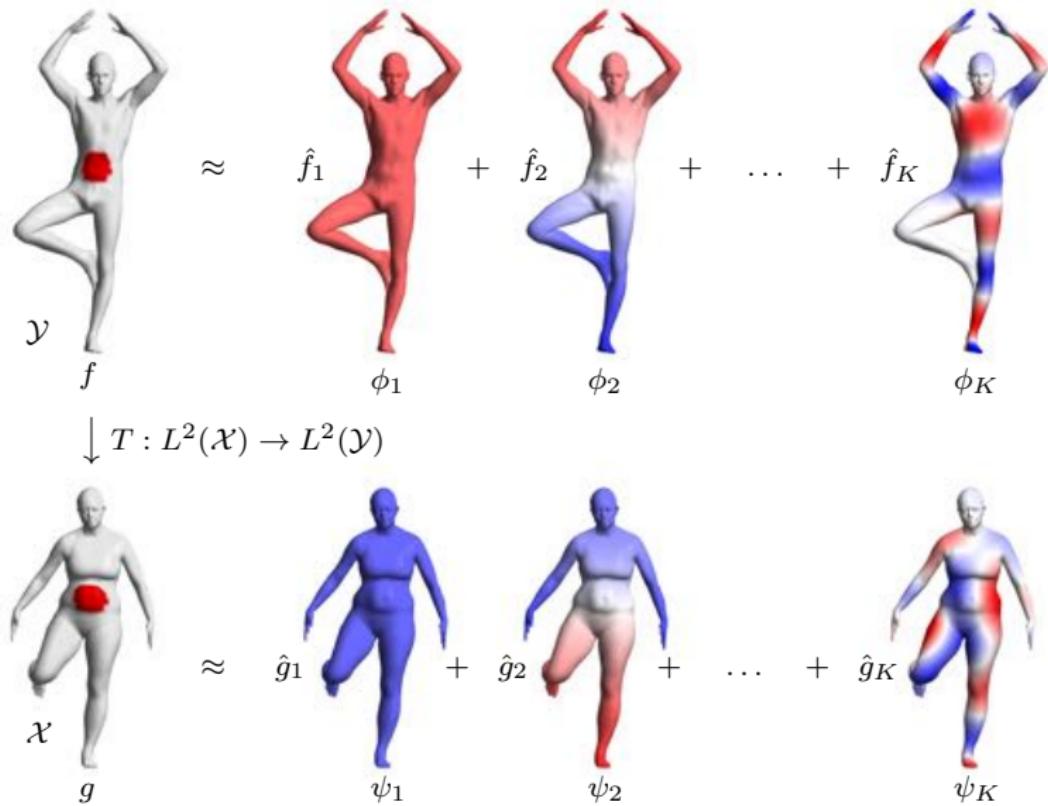
Laplacian eigenbases on non-isometric domains

 ϕ_2  ϕ_3  ϕ_{10}  ϕ_{15}  ϕ_{20}  ψ_2  ψ_3  ψ_{10}  ψ_{15}  ψ_{20}

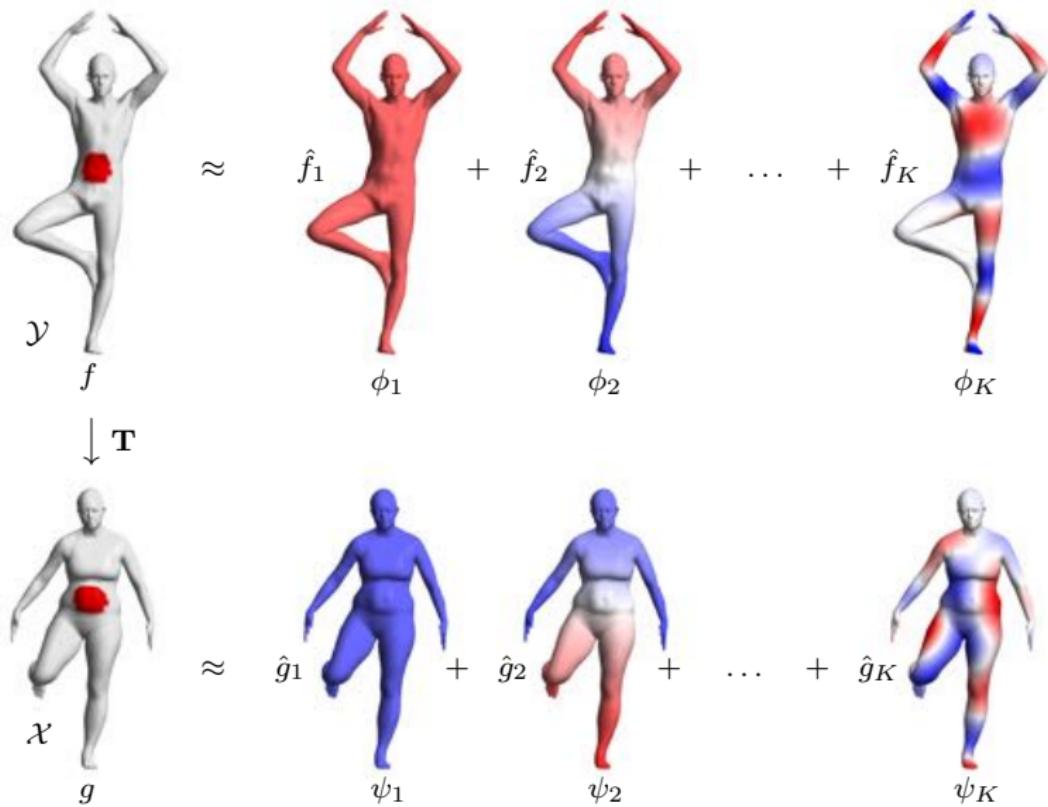
Functional maps



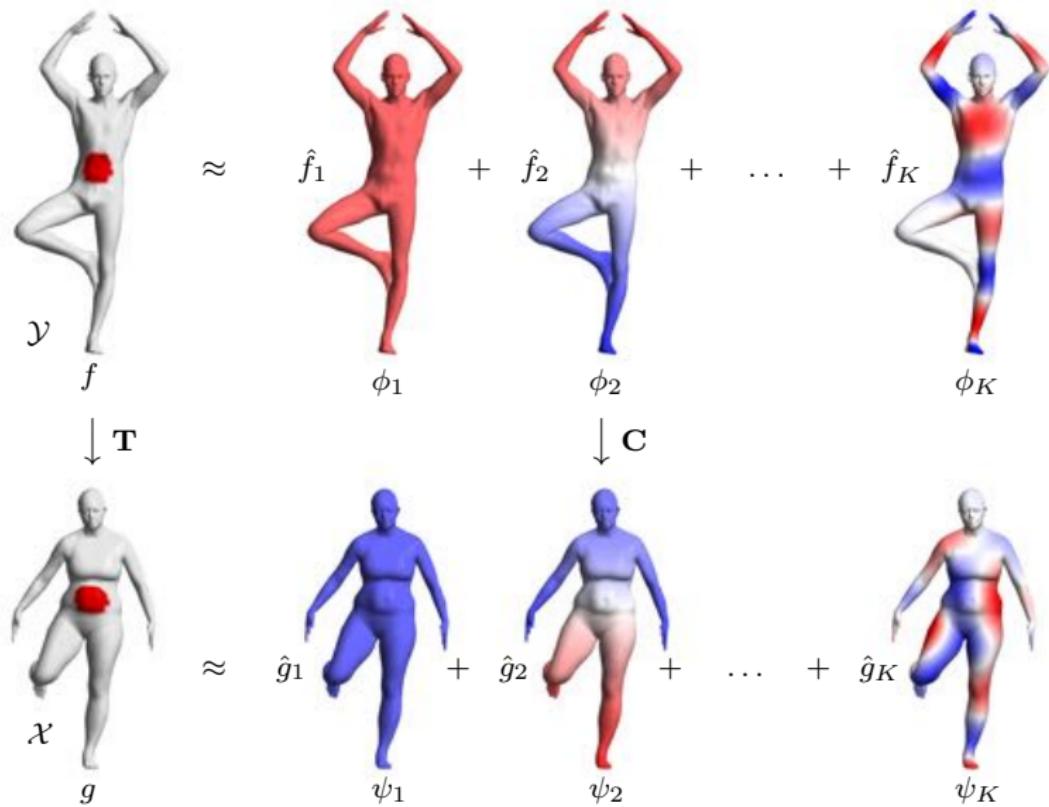
Functional maps



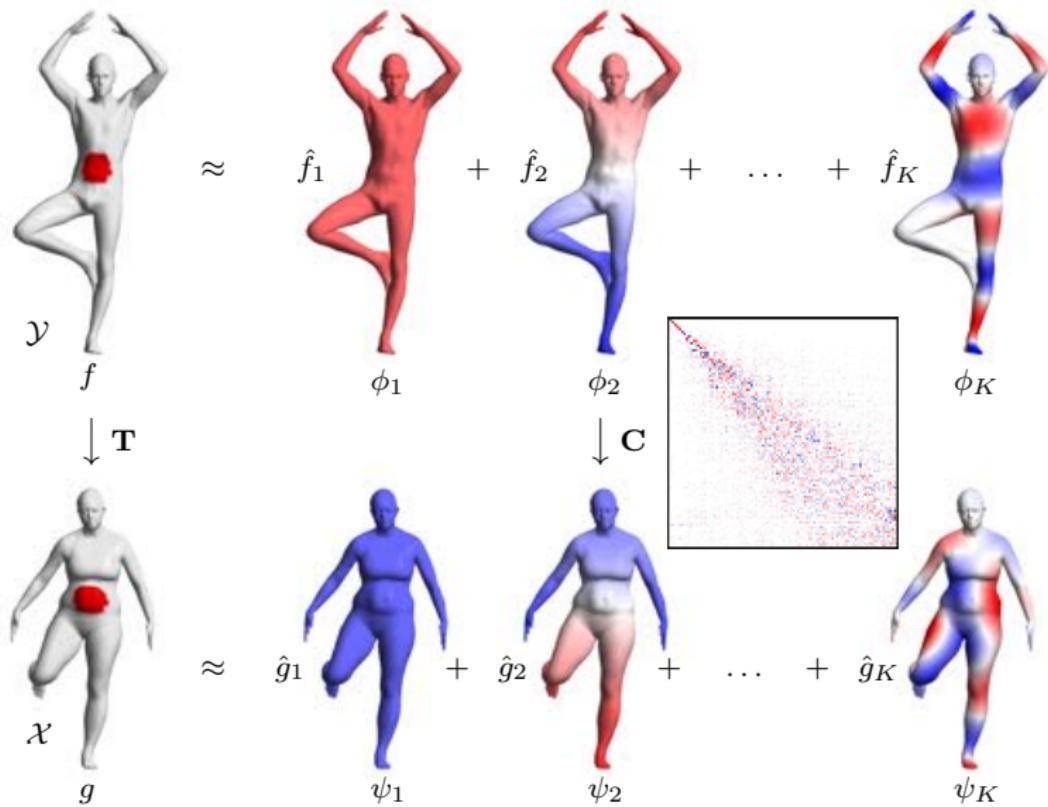
Functional maps



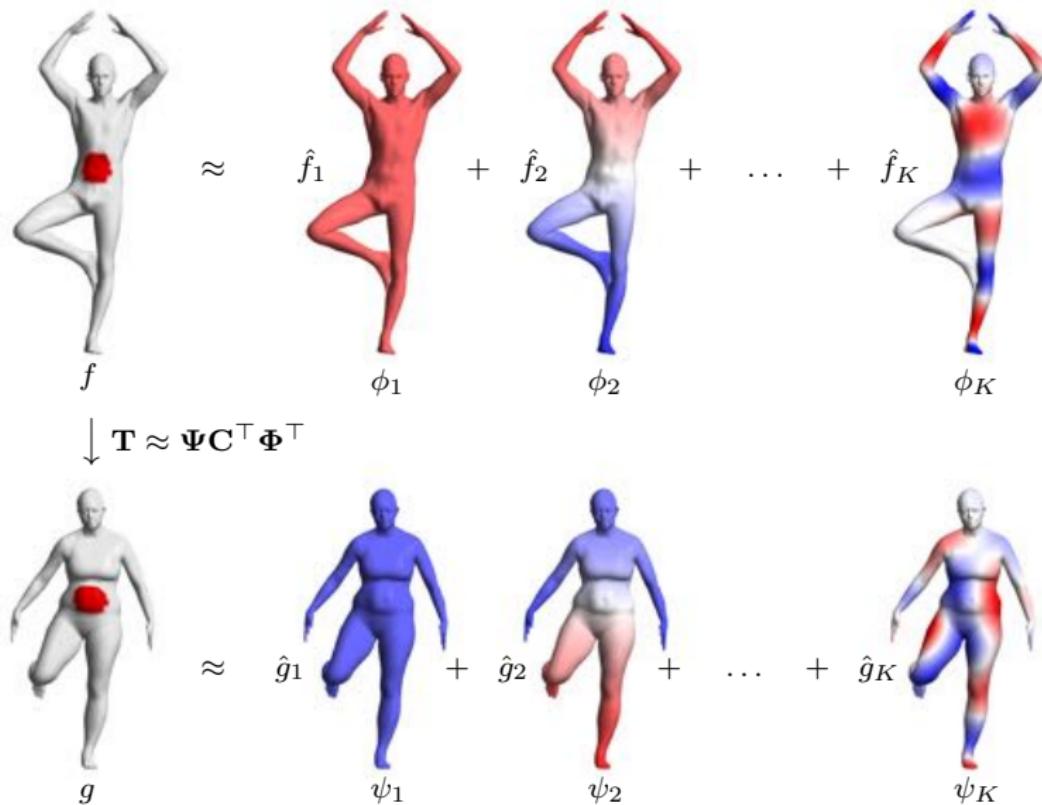
Functional maps



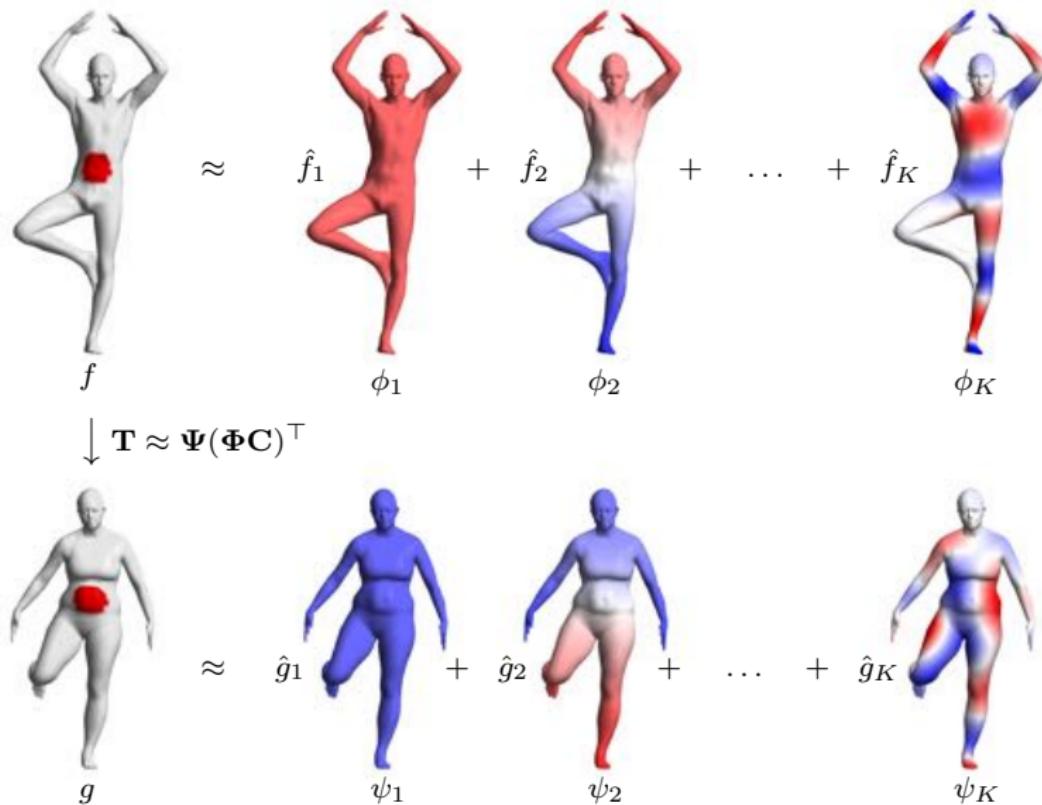
Functional maps



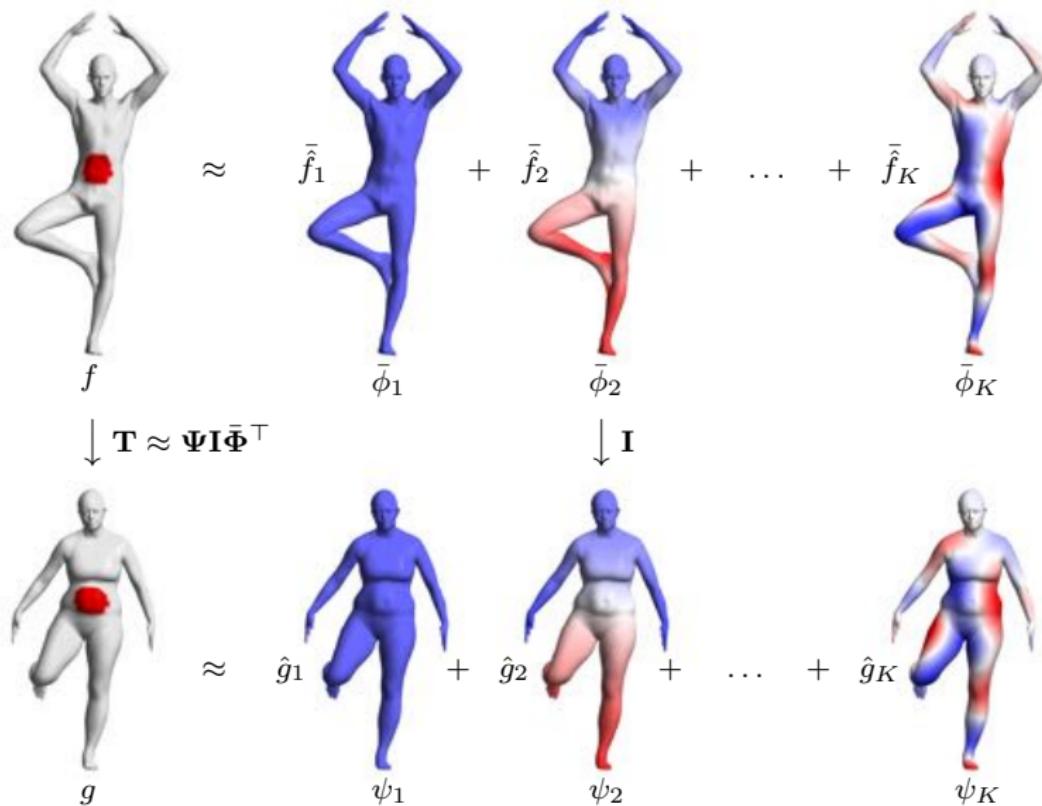
Basis synchronization with functional maps



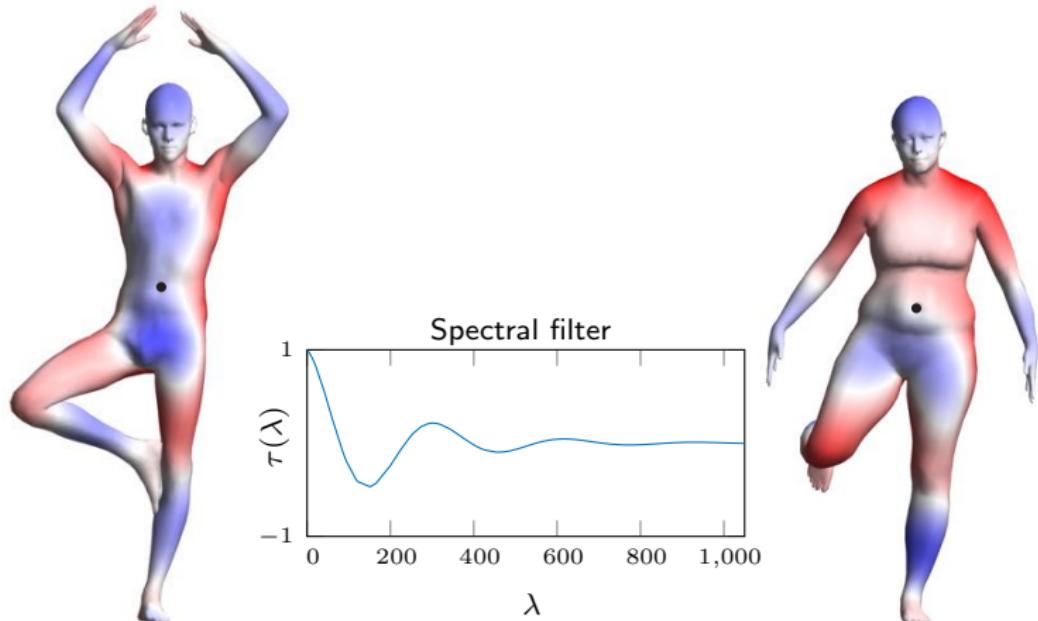
Basis synchronization with functional maps



Basis synchronization with functional maps



Filtering in different bases



$$\Phi \tau(\Lambda_\Phi) \Phi^\top \delta_0$$

$$\Psi \tau(\Lambda_\Psi) \Psi^\top \delta_0$$

Apply spectral filter $\tau(\lambda)$ in **different bases** Φ and Ψ
⇒ **different results!**

Filtering in different bases



$$\Phi \tau(\Lambda_\Phi) \Phi^\top \delta_0$$

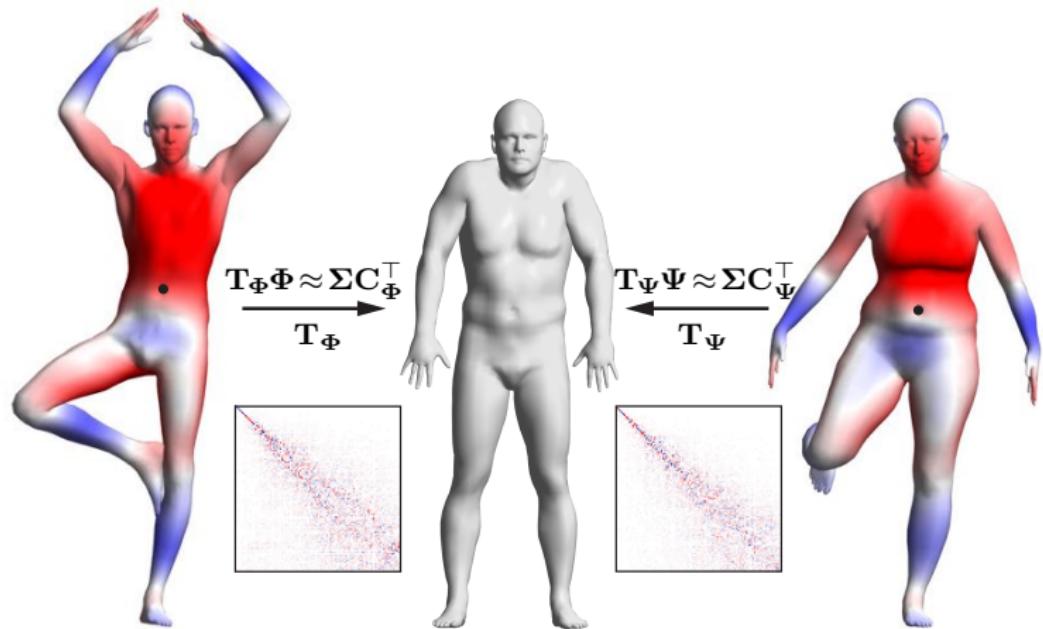
Canonical shape
with basis Σ, Λ



$$\Psi \tau(\Lambda_\Psi) \Psi^\top \delta_0$$

Apply spectral filter $\tau(\lambda)$ in **different bases** Φ and Ψ
 \Rightarrow **different results!**

Filtering in synchronized bases



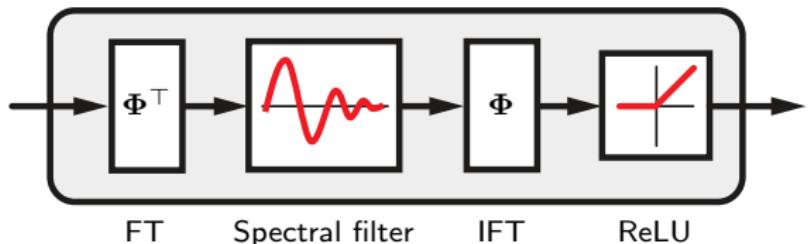
$$\Phi C_\Phi \tau(\Lambda) C_\Phi^\top \Phi^\top \delta_0$$

Canonical shape
with basis Σ, Λ

$$\Psi C_\Psi \tau(\Lambda) C_\Psi^\top \Psi^\top \delta_0$$

Apply spectral filter $\tau(\lambda)$ in **synchronized bases** ΦC_Φ and ΨC_Ψ
 \Rightarrow **similar results!**

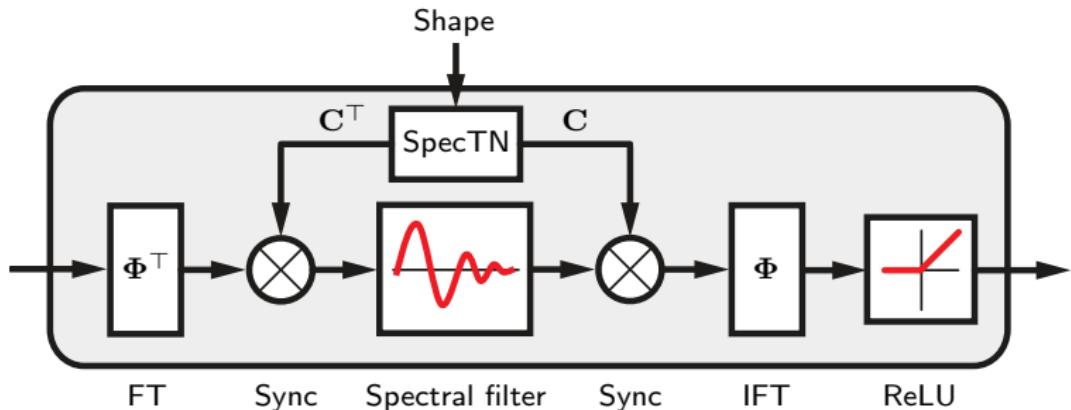
Spectral CNN



Convolutional filter of a Spectral CNN

- ⌚ Fixed basis \Rightarrow Does not generalize across domains
- ⌚ Possible $\mathcal{O}(n)$ complexity avoiding explicit FT and IFT

Spectral Transformer Network



Convolutional filter of a Spectral Transformer Network

- ☺ Basis synchronization allows generalization across domains
- ☹ Explicit FT and IFT

Example: normal prediction with SpecTN

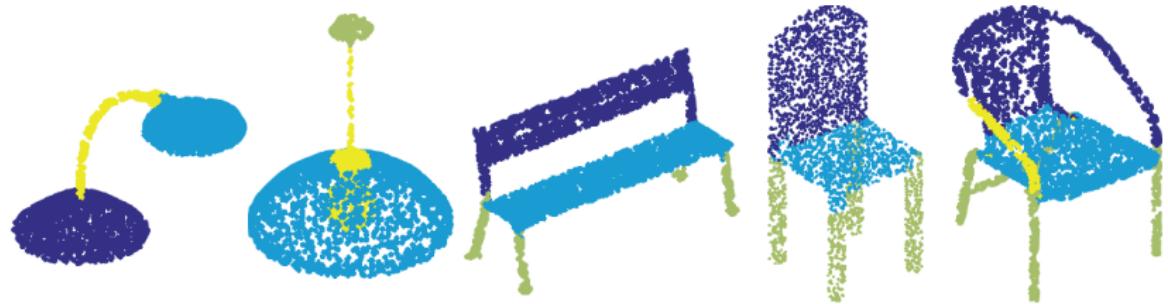


Predicted

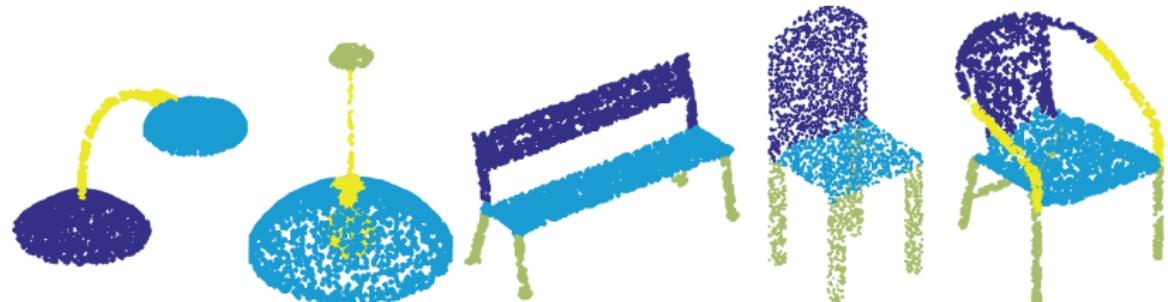


Groundtruth

Example: shape segmentation with SpecTN



Predicted



Groundtruth

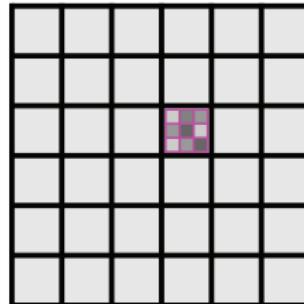
Different formulations of non-Euclidean CNNs



Spectral domain



Spatial domain



Embedding domain

Convolution

Euclidean

Spatial domain

$$(f \star g)(x) = \int_{-\pi}^{\pi} f(x')g(x-x')dx'$$

Non-Euclidean

?

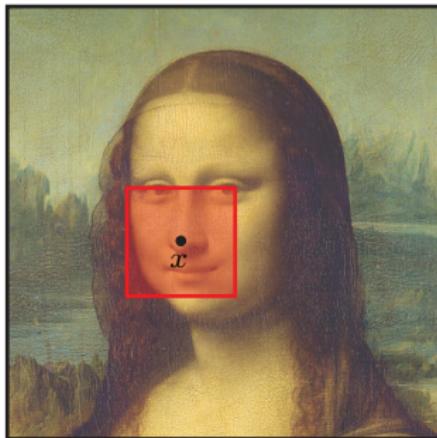
Spectral domain

$$\widehat{(f \star g)}(\omega) = \hat{f}(\omega) \cdot \hat{g}(\omega)$$

$$\widehat{(f \star g)}_k = \langle f, \phi_k \rangle_{L^2(\mathcal{X})} \langle g, \phi_k \rangle_{L^2(\mathcal{X})}$$

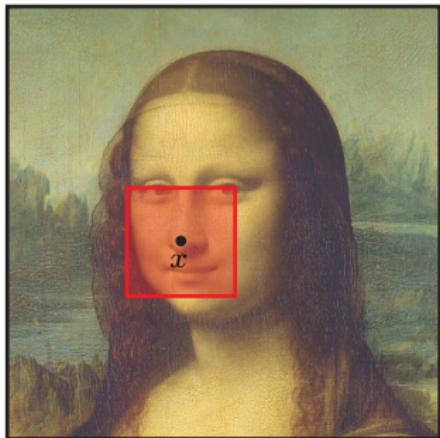
'Convolution Theorem'

Patch operator

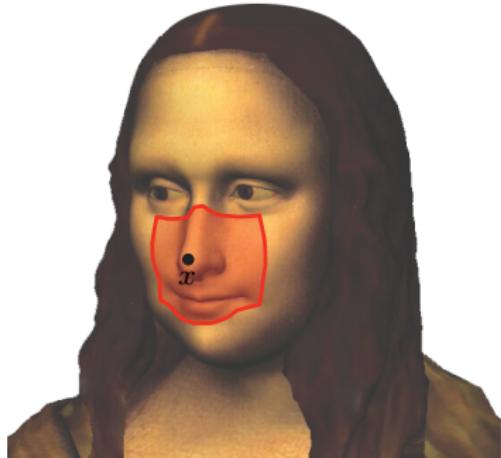


Image

Patch operator

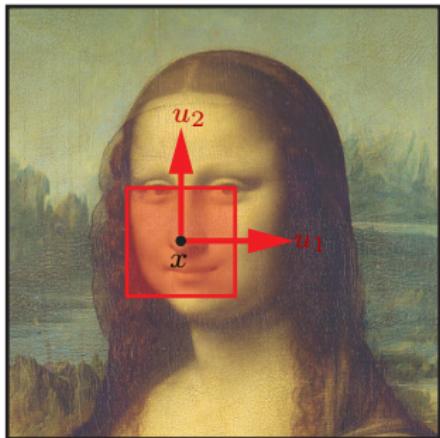


Image

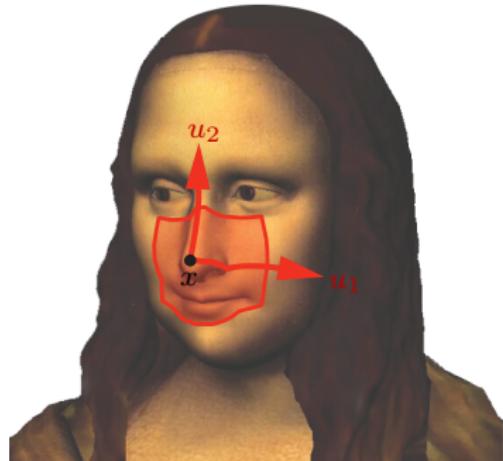


Manifold

Patch operator



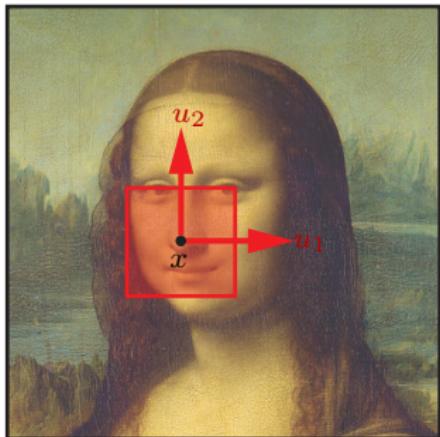
Image



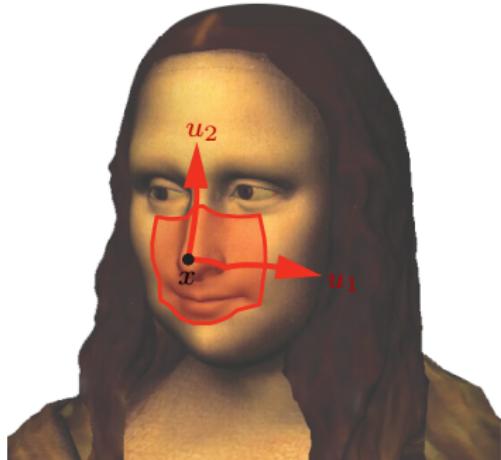
Manifold

- Local system of coordinates: bijection $\varsigma_x : B_{\rho_0}(x) \rightarrow [0, 1]^2$

Patch operator



Image

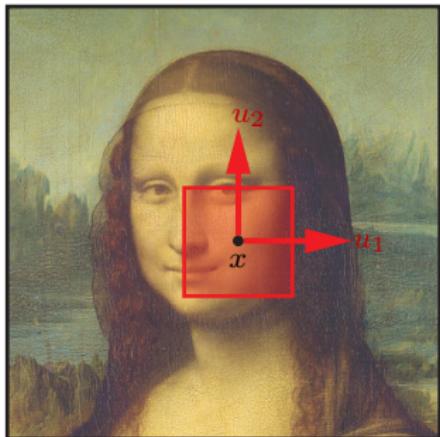


Manifold

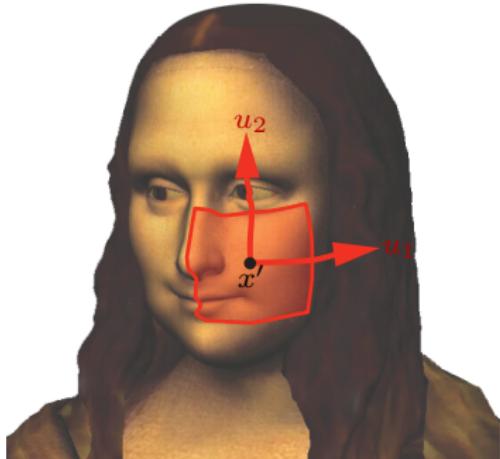
- Local system of coordinates: bijection $\varsigma_x : B_{\rho_0}(x) \rightarrow [0, 1]^2$
- **Patch operator** $\mathcal{D} : L^2(\mathcal{X}) \rightarrow L^2([0, 1]^2)$ mapping f around x

$$(\mathcal{D}(x)f)(\mathbf{u}) = (f \circ \varsigma_x^{-1})(\mathbf{u})$$

Patch operator



Image

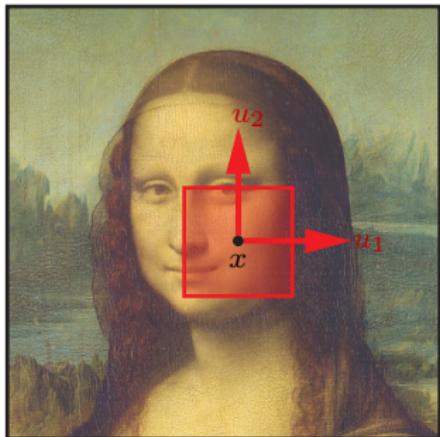


Manifold

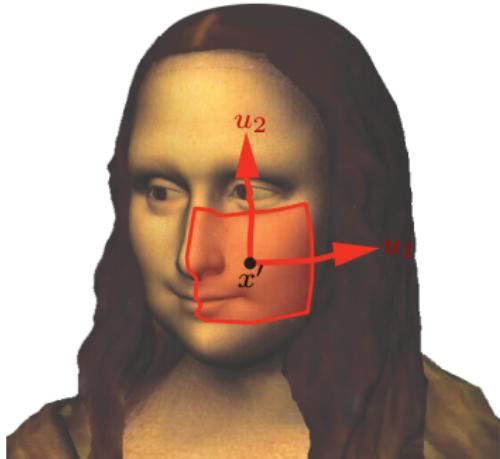
- Local system of coordinates: bijection $\varsigma_x : B_{\rho_0}(x) \rightarrow [0, 1]^2$
- **Patch operator** $\mathcal{D} : L^2(\mathcal{X}) \rightarrow L^2([0, 1]^2)$ mapping f around x

$$(\mathcal{D}(x)f)(\mathbf{u}) = (f \circ \varsigma_x^{-1})(\mathbf{u})$$

Patch operator



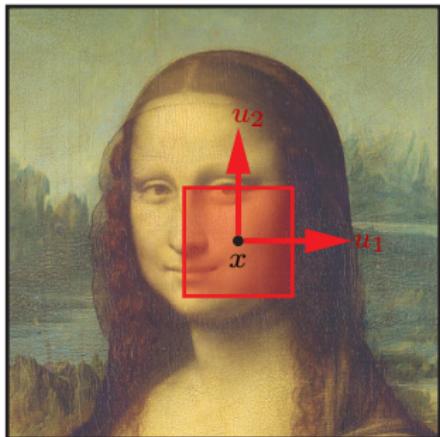
Image



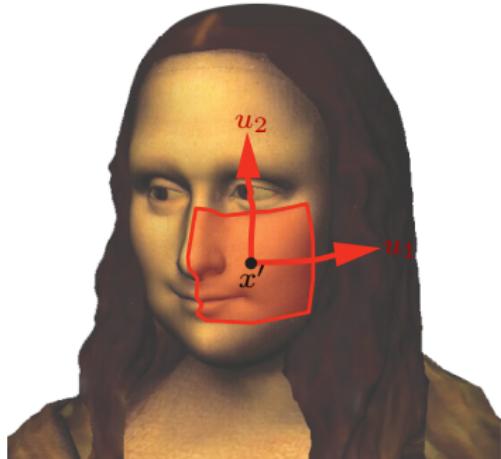
Manifold

- Local system of coordinates: bijection $\varsigma_x : B_{\rho_0}(x) \rightarrow [0, 1]^2$
- **Patch operator** applying weighting function $w_{\mathbf{u}}(x, x') = \delta_{\varsigma_x^{-1}(\mathbf{u})}(x')$
$$(\mathcal{D}(x)f)(\mathbf{u}) = \langle f, w_{\mathbf{u}}(x, x') \rangle_{L^2(\mathcal{X})}$$

Patch operator



Image



Manifold

- Local system of coordinates: bijection $\varsigma_x : B_{\rho_0}(x) \rightarrow [0, 1]^2$
- **Patch operator** applying weighting function $w_{\mathbf{u}}(x, x') = \delta_{\varsigma_x^{-1}(\mathbf{u})}(x')$
$$(\mathcal{D}(x)f)(\mathbf{u}) = \langle f, w_{\mathbf{u}}(x, x') \rangle_{L^2(\mathcal{X})}$$

In Euclidean case, $w(x, x')$ is shift-invariant

Spatial convolution

Spatial convolution of $f \in L^2(\mathcal{X})$ with continuous filter $g \in L^2([0, 1]^2)$

$$(f \star g)(x) = \int_{[0,1]^2} g(\mathbf{u})(\mathcal{D}(x)f)(\mathbf{u}) d\mathbf{u}$$

Spatial convolution

Spatial convolution of $f \in L^2(\mathcal{X})$ with discrete filter $g = (g_1, \dots, g_J)$

$$(f \star g)(x) = \sum_{j=1}^J g_j (\mathcal{D}(x)f)_j$$

Geodesic polar patch operator

Patch expressed in local **geodesic polar coordinates**

$$(\mathcal{D}(x)f)(\rho, \theta) = \int_{\mathcal{X}} \underbrace{w_\rho(x, x') w_\theta(x, x')}_{w_{\rho\theta}(x, x')} f(x') dx'$$



Radial weight

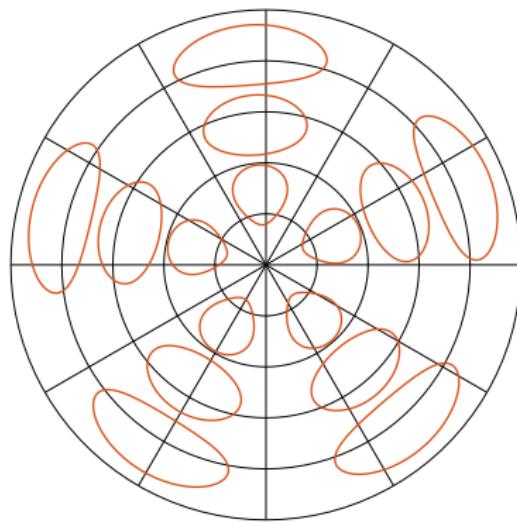
$$w_\rho(x, x') \propto e^{-(d_{\mathcal{X}}(x, x') - \rho)^2 / \sigma_\rho^2}$$



Angular weight

$$w_\theta(x, x') \propto e^{-d_{\mathcal{X}}^2(\Gamma_\theta(x), x') / \sigma_\theta^2}$$

Geodesic polar patch operator construction



Weighting functions of the geodesic polar patch operator shown in (ρ, θ) coordinates (contours mark the $\frac{1}{2}$ -level set)

Geodesic convolution

$$(f \star g)(x) = \int_0^{\rho_0} \int_0^{2\pi} (\mathcal{D}(x)f)(\rho, \theta) g(\theta, \rho) d\rho d\theta$$

Geodesic convolution

$$(f \star g)(x) = \int_0^{\rho_0} \int_0^{2\pi} \underbrace{(\mathcal{D}(x)f)(\rho, \theta)}_{(f \star g)(\rho, \theta)} \underbrace{g(\theta, \rho)}_{g(\rho, \theta)} d\rho d\theta$$



Geodesic convolution

$$(f \star g)(x) = \int_0^{\rho_0} \int_0^{2\pi} \underbrace{(\mathcal{D}(x)f)(\rho, \theta)}_{(f)(\rho, \theta)} \underbrace{g(\theta + \Delta\theta, \rho)}_{(g)(\theta + \Delta\theta, \rho)} d\rho d\theta$$



Angular coordinate origin is arbitrary = **rotation ambiguity!**

Geodesic convolution

$$(f \star g)(x) = \int_0^{\rho_0} \int_0^{2\pi} \underbrace{(\mathcal{D}(x)f)(\rho, \theta)}_{(f)(\rho, \theta)} \underbrace{g(\theta + \Delta\theta, \rho)}_{(g)(\rho, \theta + \Delta\theta)} d\rho d\theta$$



Angular coordinate origin is arbitrary = **rotation ambiguity!**

- Select reference direction, e.g. maximum curvature vector

Geodesic convolution

$$(f \star g)(x) = \int_0^{\rho_0} \int_0^{2\pi} \underbrace{(\mathcal{D}(x)f)(\rho, \theta)}_{(f \star g)(x)} \underbrace{g(\theta + \Delta\theta, \rho)}_{(f \star g)(x)} d\rho d\theta$$



Angular coordinate origin is arbitrary = **rotation ambiguity!**

- Select reference direction, e.g. maximum curvature vector
- Take Fourier transform w.r.t. θ

Geodesic convolution

$$(f \star g)(x) = \int_0^{\rho_0} \int_0^{2\pi} \underbrace{(\mathcal{D}(x)f)(\rho, \theta)}_{(\mathcal{D}(x)f)(\rho, \theta)} \underbrace{g(\theta + \Delta\theta, \rho)}_{g(\theta + \Delta\theta, \rho)} d\rho d\theta$$



Angular coordinate origin is arbitrary = **rotation ambiguity!**

- Select reference direction, e.g. maximum curvature vector
- Take Fourier transform w.r.t. θ

$$(\mathcal{D}(x)f)(\rho, \theta) \quad \longleftrightarrow \quad (\widehat{\mathcal{D}(x)f})(\rho, \omega)$$

Geodesic convolution

$$(f \star g)(x) = \int_0^{\rho_0} \int_0^{2\pi} \underbrace{(\mathcal{D}(x)f)(\rho, \theta)}_{(\mathcal{D}(x)f)(\rho, \theta + \Delta\theta)} \underbrace{g(\theta + \Delta\theta, \rho)}_{g(\theta, \rho)} d\rho d\theta$$



Angular coordinate origin is arbitrary = **rotation ambiguity!**

- Select reference direction, e.g. maximum curvature vector
- Take Fourier transform w.r.t. θ

$$(\mathcal{D}(x)f)(\rho, \theta + \Delta\theta) \quad \longleftrightarrow \quad e^{-i\omega\theta} (\widehat{\mathcal{D}(x)f})(\rho, \omega)$$

Geodesic convolution

$$(f \star g)(x) = \int_0^{\rho_0} \int_0^{2\pi} \underbrace{(\mathcal{D}(x)f)(\rho, \theta)}_{(\mathcal{D}(x)f)(\rho, \theta + \Delta\theta)} \underbrace{g(\theta + \Delta\theta, \rho)}_{g(\theta, \rho)} d\rho d\theta$$



Angular coordinate origin is arbitrary = **rotation ambiguity!**

- Select reference direction, e.g. maximum curvature vector
- Take Fourier transform magnitude w.r.t. θ

$$(\mathcal{D}(x)f)(\rho, \theta + \Delta\theta) \quad \longleftrightarrow \quad |e^{-i\omega\theta} (\widehat{\mathcal{D}(x)f})(\rho, \omega)|$$

Geodesic convolution

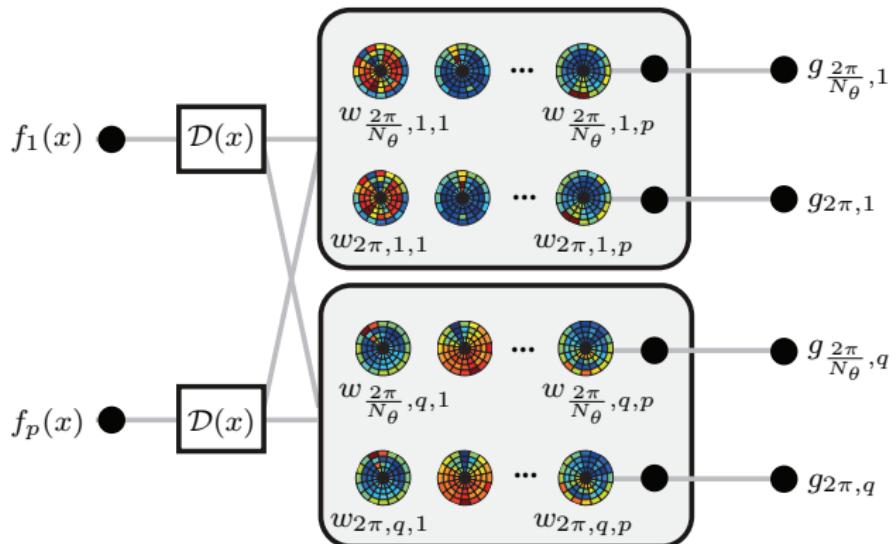
$$(f \star g)(x) = \int_0^{\rho_0} \int_0^{2\pi} \underbrace{(\mathcal{D}(x)f)(\rho, \theta)}_{(f \star g)(x)} \underbrace{g(\theta + \Delta\theta, \rho)}_{(f \star g)(x)} d\rho d\theta$$



Angular coordinate origin is arbitrary = **rotation ambiguity!**

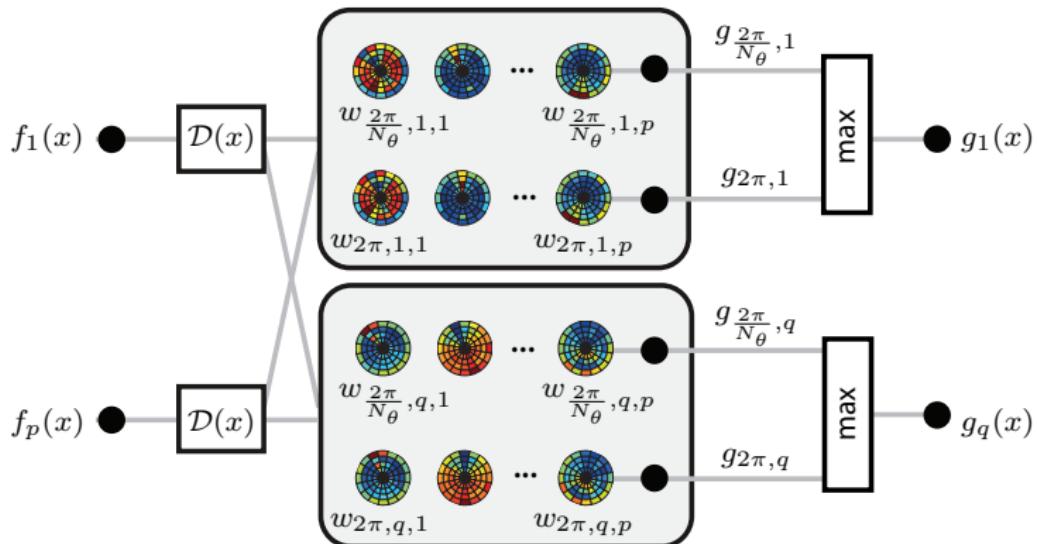
- Select reference direction, e.g. maximum curvature vector
- Take Fourier transform magnitude w.r.t. θ
- Keep all possible rotations

Geodesic convolution layer



$$\text{Conv. layer} \quad g_{\Delta\theta, l}(x) = \xi \left(\sum_{l'=1}^p (f_{l'} \star w_{\Delta\theta, l, l'})(x) \right) \quad \begin{array}{l} l = 1, \dots, q \\ l' = 1, \dots, p \\ \Delta\theta = \frac{2\pi}{N_\theta}, \dots, 2\pi \end{array}$$

Geodesic convolution layer



Conv. layer
$$g_{\Delta\theta, l}(x) = \xi \left(\sum_{l'=1}^p (f_{l'} \star w_{\Delta\theta, l, l'})(x) \right) \quad \begin{aligned} l &= 1, \dots, q \\ l' &= 1, \dots, p \\ \Delta\theta &= \frac{2\pi}{N_\theta}, \dots, 2\pi \end{aligned}$$

Angular
max pooling
$$g_l(x) = \max_{\Delta\theta} g_{\Delta\theta, l}(x)$$

Geodesic CNN (GCNN)

Convolutional layer expressed in the **spatial domain** using geodesic polar patch operator + angular max pooling to solve rotational ambiguity

$$g_l(x) = \max_{\Delta\theta} \xi \left(\sum_{l'=1}^p \underbrace{\int_0^{\rho_0} \int_0^{2\pi} w_{l,l'}(\rho, \theta + \Delta\theta) (\mathcal{D}(x)f_{l'})(\rho, \theta) d\rho d\theta}_{(f_{l'} \star w_{\Delta\theta, l, l'})(x)} \right)$$

$$l = 1, \dots, q$$

$$l' = 1, \dots, p$$

where $\mathbf{w}_{l,l'} = (w_{l,l',1}, \dots, w_{l,l',J})$ are spatial filter coefficients

Geodesic CNN (GCNN)

Convolutional layer expressed in the **spatial domain** using geodesic polar patch operator + angular max pooling to solve rotational ambiguity

$$g_l(x) = \max_{\Delta\theta} \xi \left(\sum_{l'=1}^p \underbrace{\int_0^{\rho_0} \int_0^{2\pi} w_{l,l'}(\rho, \theta + \Delta\theta) (\mathcal{D}(x)f_{l'})(\rho, \theta) d\rho d\theta}_{(f_{l'} \star w_{\Delta\theta, l, l'})(x)} \right)$$

$$l = 1, \dots, q$$

$$l' = 1, \dots, p$$

where $\mathbf{w}_{l,l'} = (w_{l,l',1}, \dots, w_{l,l',J})$ are spatial filter coefficients

⌚ Directional filters

Geodesic CNN (GCNN)

Convolutional layer expressed in the **spatial domain** using geodesic polar patch operator + angular max pooling to solve rotational ambiguity

$$g_l(x) = \max_{\Delta\theta} \xi \left(\sum_{l'=1}^p \underbrace{\int_0^{\rho_0} \int_0^{2\pi} w_{l,l'}(\rho, \theta + \Delta\theta) (\mathcal{D}(x)f_{l'})(\rho, \theta) d\rho d\theta}_{(f_{l'} \star w_{\Delta\theta, l, l'})(x)} \right)$$

$$l = 1, \dots, q$$

$$l' = 1, \dots, p$$

where $\mathbf{w}_{l,l'} = (w_{l,l',1}, \dots, w_{l,l',J})$ are spatial filter coefficients

- ☺ Directional filters
- ☺ Spatially-localized filters

Geodesic CNN (GCNN)

Convolutional layer expressed in the **spatial domain** using geodesic polar patch operator + angular max pooling to solve rotational ambiguity

$$g_l(x) = \max_{\Delta\theta} \xi \left(\sum_{l'=1}^p \underbrace{\int_0^{\rho_0} \int_0^{2\pi} w_{l,l'}(\rho, \theta + \Delta\theta) (\mathcal{D}(x)f_{l'})(\rho, \theta) d\rho d\theta}_{(f_{l'} \star w_{\Delta\theta, l, l'})(x)} \right)$$
$$l = 1, \dots, q$$
$$l' = 1, \dots, p$$

where $\mathbf{w}_{l,l'} = (w_{l,l',1}, \dots, w_{l,l',J})$ are spatial filter coefficients

- ☺ Directional filters
- ☺ Spatially-localized filters
- ☺ $\mathcal{O}(1)$ parameters per layer

Geodesic CNN (GCNN)

Convolutional layer expressed in the **spatial domain** using geodesic polar patch operator + angular max pooling to solve rotational ambiguity

$$g_l(x) = \max_{\Delta\theta} \xi \left(\sum_{l'=1}^p \underbrace{\int_0^{\rho_0} \int_0^{2\pi} w_{l,l'}(\rho, \theta + \Delta\theta) (\mathcal{D}(x)f_{l'})(\rho, \theta) d\rho d\theta}_{(f_{l'} \star w_{\Delta\theta, l, l'})(x)} \right)$$

$$l = 1, \dots, q$$

$$l' = 1, \dots, p$$

where $\mathbf{w}_{l,l'} = (w_{l,l',1}, \dots, w_{l,l',J})$ are spatial filter coefficients

- ☺ Directional filters
- ☺ Spatially-localized filters
- ☺ $\mathcal{O}(1)$ parameters per layer
- ☺ All operations are local $\Rightarrow \mathcal{O}(n)$ computational complexity

Geodesic CNN (GCNN)

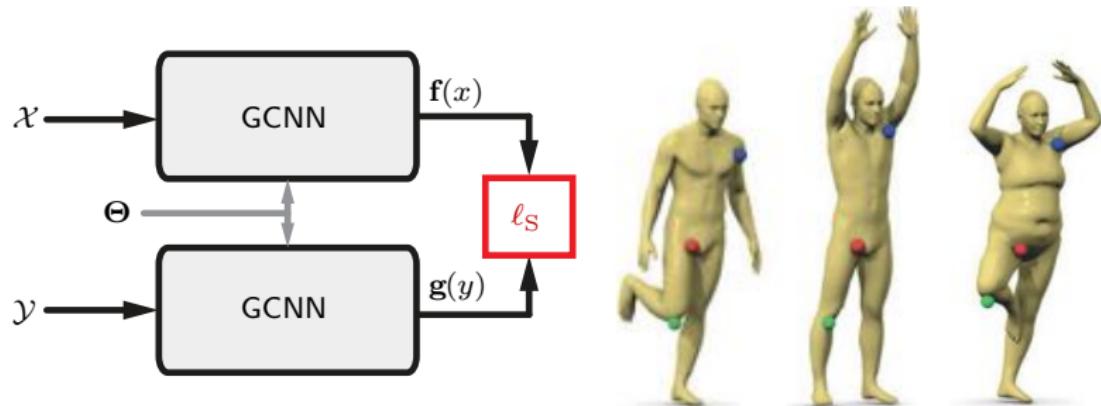
Convolutional layer expressed in the spatial domain using geodesic polar patch operator + angular max pooling to solve rotational ambiguity

$$g_l(x) = \max_{\Delta\theta} \xi \left(\sum_{l'=1}^p \underbrace{\int_0^{\rho_0} \int_0^{2\pi} w_{l,l'}(\rho, \theta + \Delta\theta) (\mathcal{D}(x)f_{l'})(\rho, \theta) d\rho d\theta}_{(f_{l'} \star w_{\Delta\theta, l, l'})(x)} \right)$$
$$l = 1, \dots, q$$
$$l' = 1, \dots, p$$

where $\mathbf{w}_{l,l'} = (w_{l,l',1}, \dots, w_{l,l',J})$ are spatial filter coefficients

- ☺ Directional filters
- ☺ Spatially-localized filters
- ☺ $\mathcal{O}(1)$ parameters per layer
- ☺ All operations are local $\Rightarrow \mathcal{O}(n)$ computational complexity
- ☺ Angular max pooling potentially reduces discriminativity

Example: Learning local descriptors with GCNN



Training set positive (x, x^+) and negative (x, x^-) pairs of points

Siamese net two net instances with shared parameters Θ

$$\begin{aligned} \ell_S(\Theta) = & \gamma \sum_{x, x^+} \|\mathbf{f}_\Theta(x) - \mathbf{f}_\Theta(x^+)\|_2^2 \\ & + (1 - \gamma) \sum_{x, x^-} [\mu - \|\mathbf{f}_\Theta(x) - \mathbf{f}_\Theta(x^-)\|_2^2]_+ \end{aligned}$$

Example: HKS descriptor



Distance in the space of local Heat Kernel Signature (HKS) features
(shown is distance from a point on the shoulder marked in white)

Example: WKS descriptor



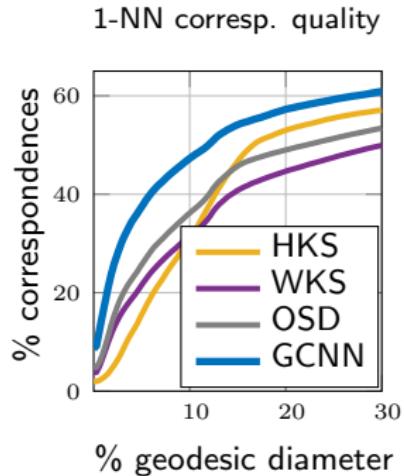
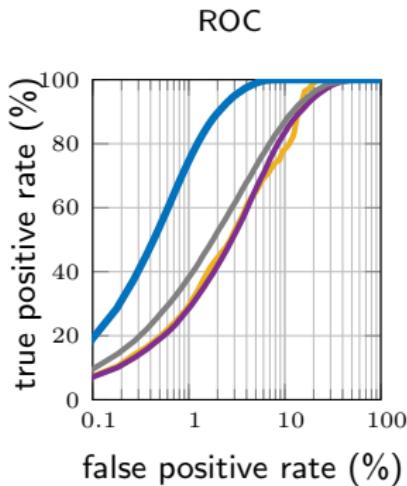
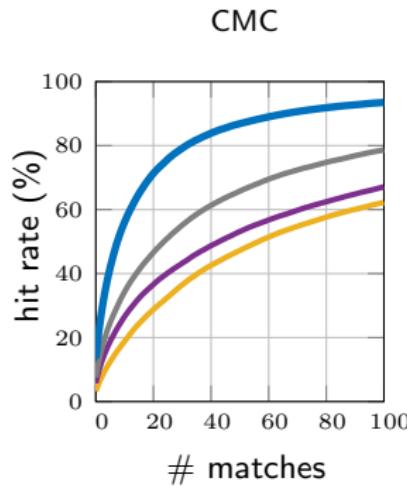
Distance in the space of local Wave Kernel Signature (WKS) features
(shown is distance from a point on the shoulder marked in white)

Example: descriptor learning with GCNN



Distance in the space of local GCNN features
(shown is distance from a point on the shoulder marked in white)

Descriptor quality comparison



Descriptor performance using symmetric Princeton benchmark
(training and testing: disjoint subsets of FAUST)

Homogeneous diffusion

$$f_t(x) = -c\Delta f(x)$$

c = [thermal diffusivity constant](#) describing heat conduction properties of the material (diffusion speed is equal everywhere)

Anisotropic diffusion

$$f_t(x) = -\operatorname{div}(c \nabla f(x))$$

c = **thermal diffusivity constant** describing heat conduction properties of the material (diffusion speed is equal everywhere)

diffusion

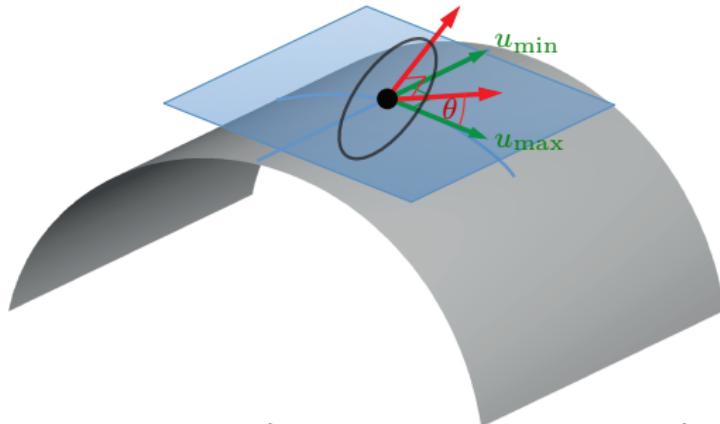
$$f_t(x) = -\operatorname{div}(\mathbf{A}(x)\nabla f(x))$$

$\mathbf{A}(x)$ = heat conductivity tensor describing heat conduction properties of the material (diffusion speed is position + direction dependent)

Anisotropic diffusion

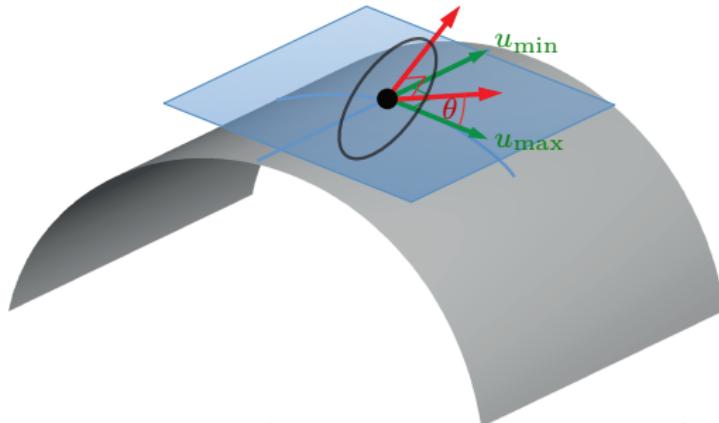
Anisotropic

Anisotropic diffusion on manifolds



$$f_t(x) = -\operatorname{div} \left(\mathbf{R}_\theta \begin{pmatrix} \alpha & \\ & 1 \end{pmatrix} \mathbf{R}_\theta^\top \nabla f(x) \right)$$

Anisotropic diffusion on manifolds



$$f_t(x) = -\operatorname{div} \left(\underbrace{\mathbf{R}_\theta \begin{pmatrix} \alpha \\ 1 \end{pmatrix} \mathbf{R}_\theta^\top}_{\mathbf{D}_{\alpha\theta}(x)} \nabla f(x) \right)$$

- Anisotropic Laplacian $\Delta_{\alpha\theta} f(x) = \operatorname{div} (D_{\alpha\theta}(x) \nabla f(x))$
- θ = orientation w.r.t. max curvature direction
- α = 'elongation'

Anisotropic heat kernels

$$h_{\alpha\theta t}(x, x') = \sum_{k \geq 0} e^{-t\lambda_{\alpha\theta k}} \phi_{\alpha\theta k}(x) \phi_{\alpha\theta k}(x')$$

Orientation θ

Elongation α

Anisotropic CNN (ACNN)

Use anisotropic heat kernels as weighting functions of the patch operator

$$(\mathcal{D}(x)f)_j = \int_{\mathcal{X}} f(x') h_{\alpha_i, \theta_i, t_i}(x, x') dx' \quad j = 1, \dots, J$$

for a discrete set of angles/scales/anisotropic constants

Anisotropic CNN (ACNN)

Use anisotropic heat kernels as weighting functions of the patch operator

$$(\mathcal{D}(x)f)_j = \int_{\mathcal{X}} f(x') h_{\alpha_i, \theta_i, t_i}(x, x') dx' \quad j = 1, \dots, J$$

for a discrete set of angles/scales/anisotropic constants

Anisotropic CNN (ACNN)

Use anisotropic heat kernels as weighting functions of the patch operator

$$(\mathcal{D}(x)f)_j = \int_{\mathcal{X}} f(x') h_{\alpha_i, \theta_i, t_i}(x, x') dx' \quad j = 1, \dots, J$$

for a discrete set of angles/scales/anisotropic constants

Convolutional layer expressed in the **spatial domain**

$$g_l(x) = \xi \left(\sum_{l'=1}^p \mathbf{w}_{l,l'}^\top \mathcal{D}(x)f \right) \quad l = 1, \dots, q \\ \quad l' = 1, \dots, p$$

where $\mathbf{w}_{l,l'} = (w_{l,l',1}, \dots, w_{l,l',J})$ are spatial filter coefficients

Anisotropic CNN (ACNN)

Use anisotropic heat kernels as weighting functions of the patch operator

$$(\mathcal{D}(x)f)_j = \int_{\mathcal{X}} f(x') h_{\alpha_i, \theta_i, t_i}(x, x') dx' \quad j = 1, \dots, J$$

for a discrete set of angles/scales/anisotropic constants

Convolutional layer expressed in the spatial domain

$$g_l(x) = \xi \left(\sum_{l'=1}^p \mathbf{w}_{l,l'}^\top \mathcal{D}(x)f \right) \quad l = 1, \dots, q \\ l' = 1, \dots, p$$

where $\mathbf{w}_{l,l'} = (w_{l,l',1}, \dots, w_{l,l',J})$ are spatial filter coefficients

- ☺ Directional filters
- ☺ Spatially-localized filters
- ☺ $\mathcal{O}(1)$ parameters per layer

Anisotropic CNN (ACNN)

Use anisotropic heat kernels as weighting functions of the patch operator

$$(\mathcal{D}(x)f)_j = \int_{\mathcal{X}} f(x') h_{\alpha_i, \theta_i, t_i}(x, x') dx' \quad j = 1, \dots, J$$

for a discrete set of angles/scales/anisotropic constants

Convolutional layer expressed in the spatial domain

$$g_l(x) = \xi \left(\sum_{l'=1}^p \mathbf{w}_{l,l'}^\top \mathcal{D}(x)f \right) \quad l = 1, \dots, q \\ \quad l' = 1, \dots, p$$

where $\mathbf{w}_{l,l'} = (w_{l,l',1}, \dots, w_{l,l',J})$ are spatial filter coefficients

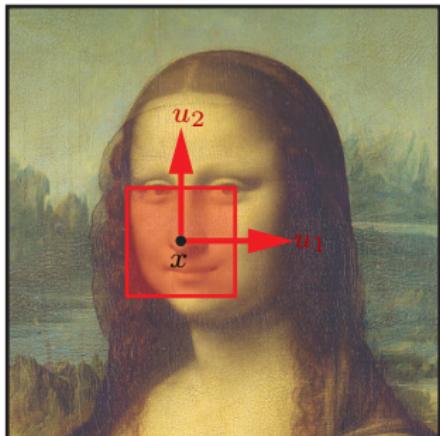
- ☺ Directional filters
- ☺ Spatially-localized filters
- ☺ $\mathcal{O}(1)$ parameters per layer
- ☺ Expensive computation of heat kernels for many orientations

Example: descriptor learning with ACNN

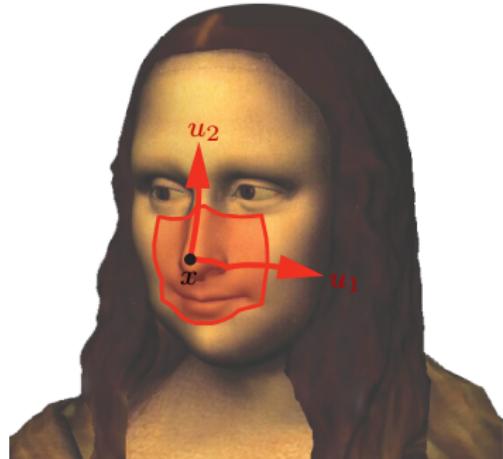


Distance in the space of local ACNN features
(shown is distance from a point on the shoulder marked in white)

Learnable patch operator



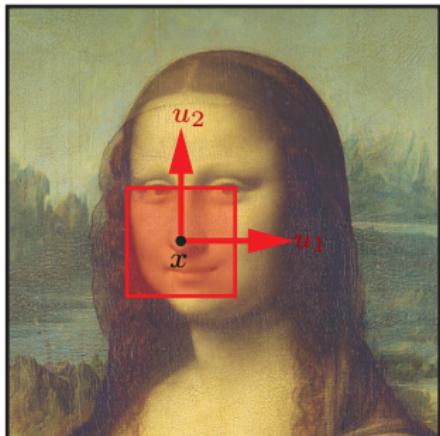
Image



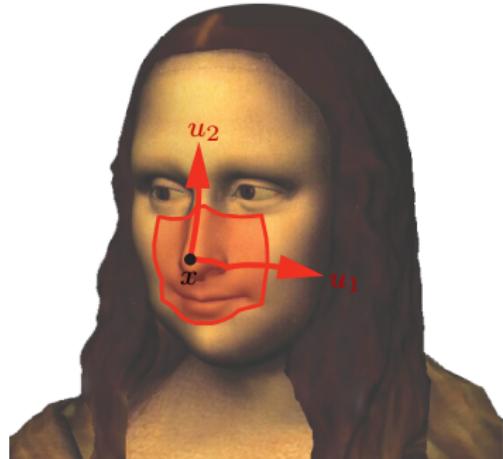
Manifold

- Local system of coordinates $\mathbf{u}(x, x')$ around point x

Learnable patch operator



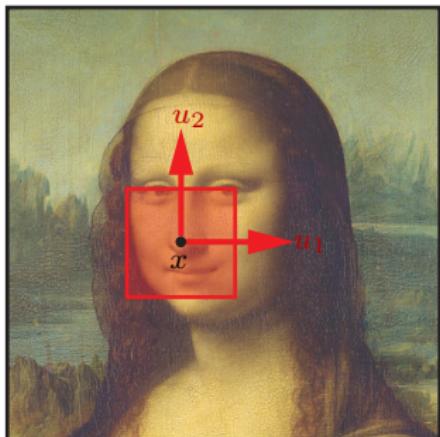
Image



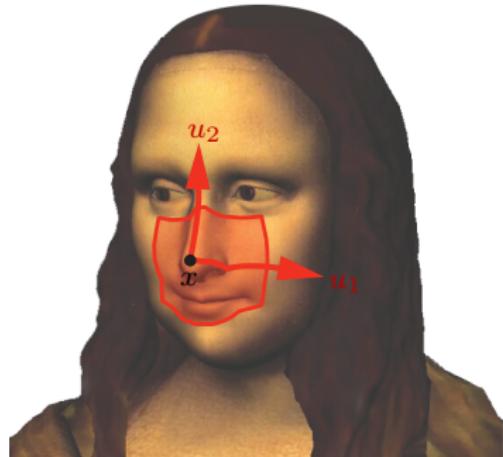
Manifold

- Local system of coordinates $\mathbf{u}(x, x')$ around point x
- Parametric weighting functions $\mathbf{w}_\Theta(\mathbf{u})$

Learnable patch operator



Image

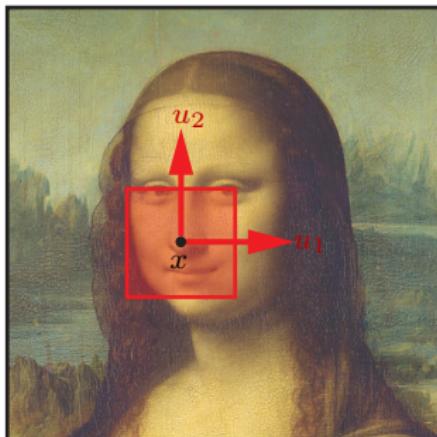


Manifold

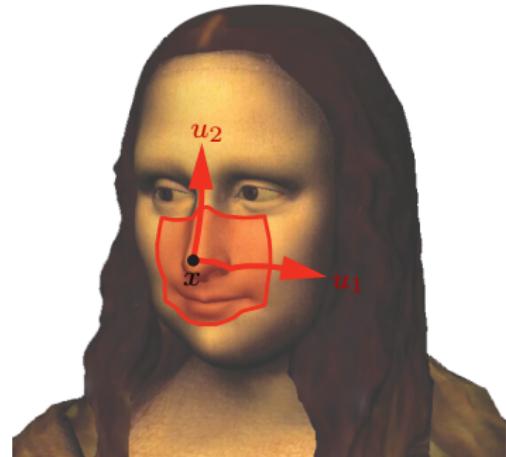
- Local system of coordinates $\mathbf{u}(x, x')$ around point x
- Parametric weighting functions $w_{\Theta}(\mathbf{u})$
- Parametric patch operator applying J such weighting functions

$$(\mathcal{D}_{\Theta_1, \dots, \Theta_J}(x)f)_j = \int_{\mathcal{X}} f(x') w_{\Theta_j}(\mathbf{u}(x, x')) dx' \quad j = 1, \dots, J$$

Learnable patch operator



Image



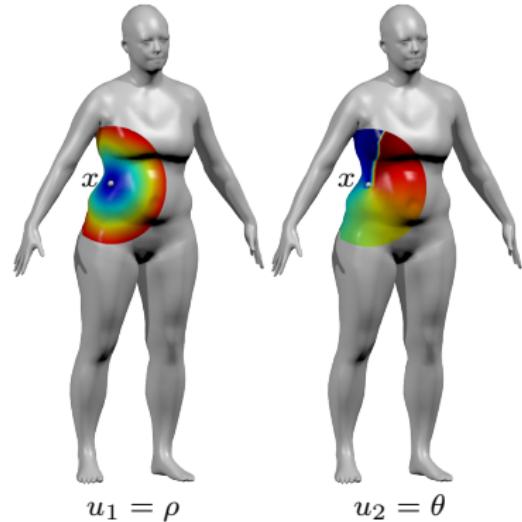
Manifold

- Local system of coordinates $\mathbf{u}(x, x')$ around point x
- Parametric weighting functions, e.g. $w_{\mu, \Sigma}(\mathbf{u}) = e^{-\frac{1}{2}(\mathbf{u}-\boldsymbol{\mu})^\top \boldsymbol{\Sigma}^{-1}(\mathbf{u}-\boldsymbol{\mu})}$
- Parametric patch operator applying J such weighting functions

$$(\mathcal{D}_{\boldsymbol{\mu}_1, \Sigma_1, \dots, \boldsymbol{\mu}_J, \Sigma_J}(x)f)_j = \int_{\mathcal{X}} f(x') w_{\boldsymbol{\mu}_j, \Sigma_j}(\mathbf{u}(x, x')) dx' \quad j = 1, \dots, J$$

Learnable patches on manifolds

- Geodesic polar coordinates
 $\mathbf{u}(x, y) = (\rho(x, y), \theta(x, y))$



Learnable patches on manifolds

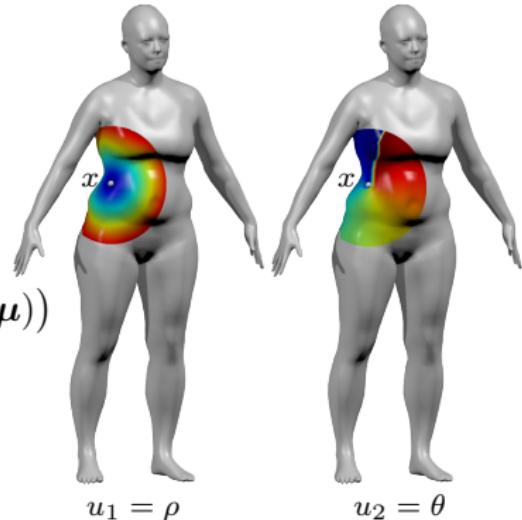
- Geodesic polar coordinates

$$\mathbf{u}(x, y) = (\rho(x, y), \theta(x, y))$$

- Gaussian weighting functions

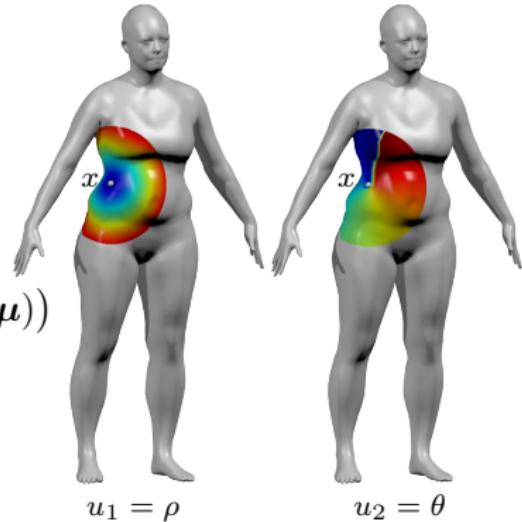
$$w_{\mu, \Sigma}(\mathbf{u}) = \exp\left(-\frac{1}{2}(\mathbf{u} - \mu)^\top \Sigma^{-1}(\mathbf{u} - \mu)\right)$$

with learnable covariance Σ and
mean μ



Learnable patches on manifolds

- Geodesic polar coordinates
 $\mathbf{u}(x, y) = (\rho(x, y), \theta(x, y))$
- Gaussian weighting functions
 $w_{\mu, \Sigma}(\mathbf{u}) = \exp\left(-\frac{1}{2}(\mathbf{u} - \mu)^\top \Sigma^{-1}(\mathbf{u} - \mu)\right)$
with learnable covariance Σ and
mean μ

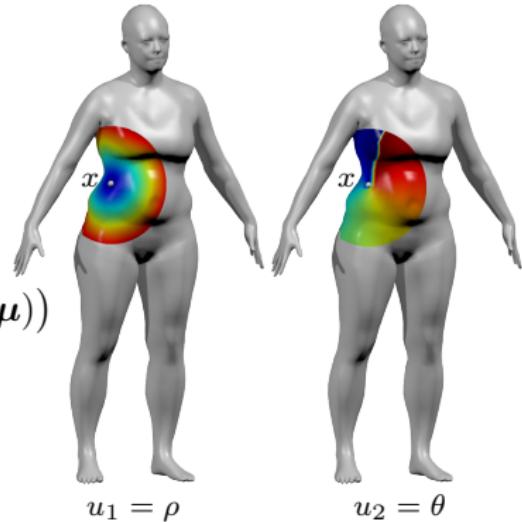


Spatial convolution

$$(f \star g)(x) = \sum_{j=1}^J g_j \int_{\mathcal{X}} w_{\mu_j, \Sigma_j}(\mathbf{u}(x, x')) f(x') dx'$$

Learnable patches on manifolds

- Geodesic polar coordinates
 $\mathbf{u}(x, y) = (\rho(x, y), \theta(x, y))$
- Gaussian weighting functions
 $w_{\mu, \Sigma}(\mathbf{u}) = \exp\left(-\frac{1}{2}(\mathbf{u} - \mu)^\top \Sigma^{-1}(\mathbf{u} - \mu)\right)$
with learnable covariance Σ and
mean μ



Spatial convolution

$$(f \star g)(x) = \int_{\mathcal{X}} \sum_{j=1}^J g_j w_{\mu_j, \Sigma_j}(\mathbf{u}(x, x')) f(x') dx'$$

Learnable patches on manifolds

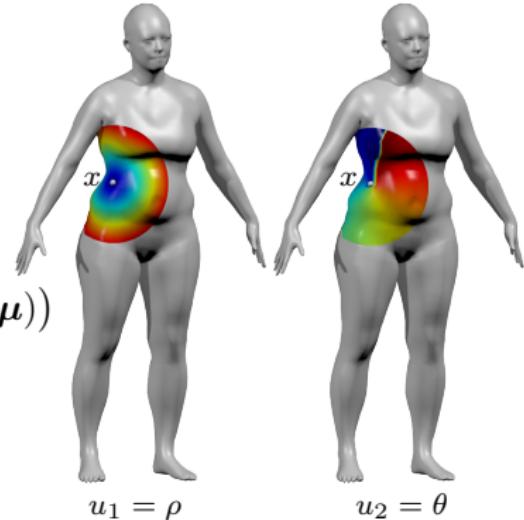
- Geodesic polar coordinates

$$\mathbf{u}(x, y) = (\rho(x, y), \theta(x, y))$$

- Gaussian weighting functions

$$w_{\mu, \Sigma}(\mathbf{u}) = \exp\left(-\frac{1}{2}(\mathbf{u} - \mu)^\top \Sigma^{-1}(\mathbf{u} - \mu)\right)$$

with learnable covariance Σ and
mean μ



Spatial convolution

$$(f \star g)(x) = \int_{\mathcal{X}} \underbrace{\sum_{j=1}^J g_j w_{\mu_j, \Sigma_j}(\mathbf{u}(x, x'))}_{\text{Gaussian mixture } g(\mathbf{u}(x, x'))} f(x') dx'$$

Mixture Model Networks (MoNet)

Convolutional layer expressed in the spatial domain using a learnable patch operator

$$g_l(x) = \xi \left(\sum_{l'=1}^p \mathbf{w}_{l,l'}^\top \mathcal{D}_\Theta(x) f \right) \quad \begin{matrix} l = 1, \dots, q \\ l' = 1, \dots, p \end{matrix}$$

where $\mathbf{w}_{l,l'} = (w_{l,l',1}, \dots, w_{l,l',J})$ are spatial filter coefficients and $\Theta = (\boldsymbol{\mu}_1, \Sigma_1, \dots, \boldsymbol{\mu}_J, \Sigma_J)$ are the patch parameters

Mixture Model Networks (MoNet)

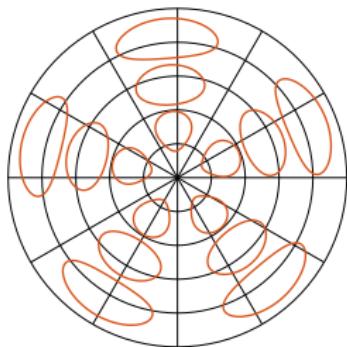
Convolutional layer expressed in the spatial domain using a learnable patch operator

$$g_l(x) = \xi \left(\sum_{l'=1}^p \mathbf{w}_{l,l'}^\top \mathcal{D}_\Theta(x) f \right) \quad \begin{matrix} l = 1, \dots, q \\ l' = 1, \dots, p \end{matrix}$$

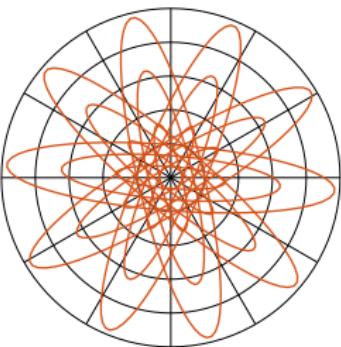
where $\mathbf{w}_{l,l'} = (w_{l,l',1}, \dots, w_{l,l',J})$ are spatial filter coefficients and $\Theta = (\boldsymbol{\mu}_1, \Sigma_1, \dots, \boldsymbol{\mu}_J, \Sigma_J)$ are the patch parameters

- ☺ Directional filters
- ☺ Spatially-localized filters
- ☺ Learnable patch operator
- ☺ $\mathcal{O}(1)$ parameters per layer

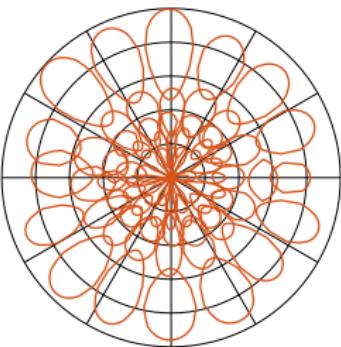
Patch operator weight functions



GCNN



ACNN



MoNet

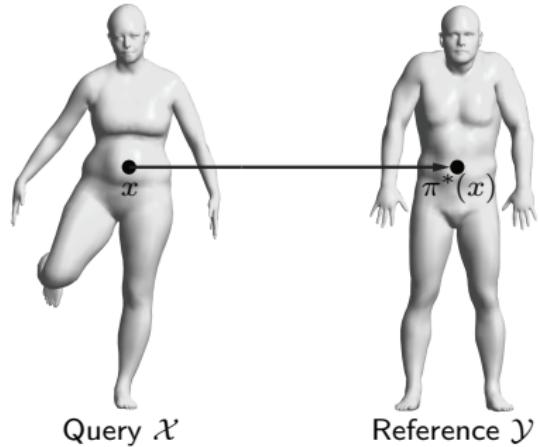
MoNet as generalization of previous methods

Method	Coordinates $\mathbf{u}(x, x')$	Weight function $w_{\Theta}(\mathbf{u})$
CNN ¹	$\mathbf{u}(x') - \mathbf{u}(x)$	$\delta(\mathbf{u} - \mathbf{v})$ fixed parameters $\Theta = \mathbf{v}$
GCNN ²	$\rho(x, x'), \theta(x, x')$	$\exp\left(-\frac{1}{2}(\mathbf{u} - \mathbf{v})^T \begin{pmatrix} \sigma_\rho^2 & \\ & \sigma_\theta^2 \end{pmatrix}^{-1} (\mathbf{u} - \mathbf{v})\right)$ fixed parameters $\Theta = (\mathbf{v}, \sigma_\rho, \sigma_\theta)$
ACNN ³	$\rho(x, x'), \theta(x, x')$	$\exp(-t\mathbf{u}^T \mathbf{R}_\varphi(\alpha_1) \mathbf{R}_\varphi^T \mathbf{u})$ fixed parameters $\Theta = (\alpha, \varphi, t)$
MoNet ⁴	$\rho(x, x'), \theta(x, x')$	$\exp\left(-\frac{1}{2}(\mathbf{u} - \boldsymbol{\mu})^T \boldsymbol{\Sigma}^{-1} (\mathbf{u} - \boldsymbol{\mu})\right)$ learnable parameters $\Theta = (\boldsymbol{\mu}, \boldsymbol{\Sigma})$

Some CNN models can be considered as particular settings of MoNet
with weighting functions of different form

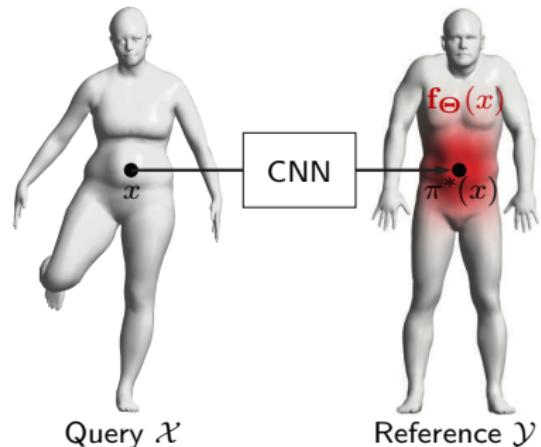
Learning deformation-invariant correspondence

- Groundtruth correspondence
 $\pi^* : \mathcal{X} \rightarrow \mathcal{Y}$ from query shape \mathcal{X}
to some reference shape \mathcal{Y}
(discretized with n vertices)
- Correspondence = **label** each query vertex x as reference vertex y



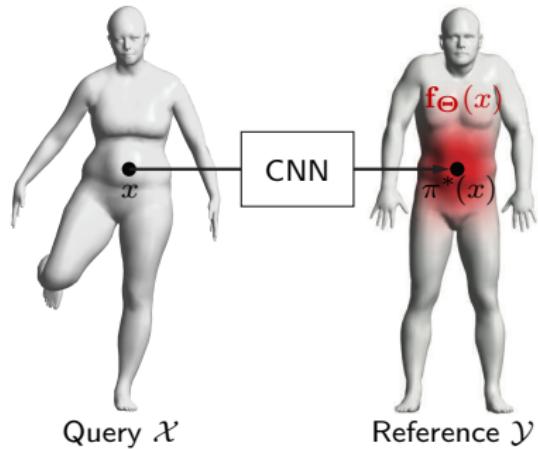
Learning deformation-invariant correspondence

- Groundtruth correspondence
 $\pi^* : \mathcal{X} \rightarrow \mathcal{Y}$ from query shape \mathcal{X}
to some reference shape \mathcal{Y}
(discretized with n vertices)
- Correspondence = **label** each query vertex x as reference vertex y
- Net output at x after softmax layer
 $f_{\Theta}(x) = (f_{\Theta,1}(x), \dots, f_{\Theta,n}(x))$
= probability distribution on \mathcal{Y}



Learning deformation-invariant correspondence

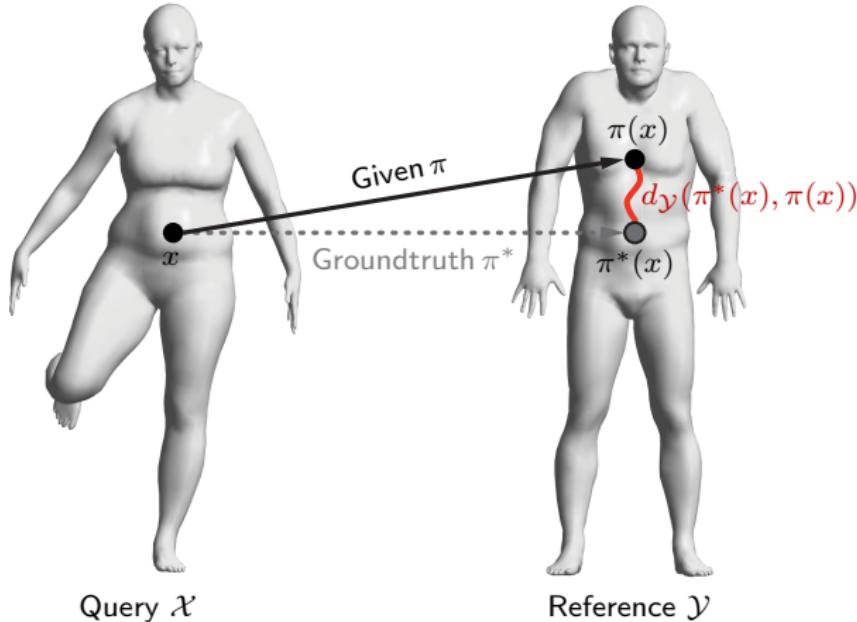
- Groundtruth correspondence
 $\pi^* : \mathcal{X} \rightarrow \mathcal{Y}$ from query shape \mathcal{X}
to some reference shape \mathcal{Y}
(discretized with n vertices)
- Correspondence = **label** each query vertex x as reference vertex y
- Net output at x after softmax layer
 $f_{\Theta}(x) = (f_{\Theta,1}(x), \dots, f_{\Theta,n}(x))$
= probability distribution on \mathcal{Y}



Minimize on training set the **cross entropy** between groundtruth correspondence and output probability distribution w.r.t. net parameters Θ

$$\min_{\Theta} \sum_x H(\delta_{\pi^*(x)}, f_{\Theta}(x))$$

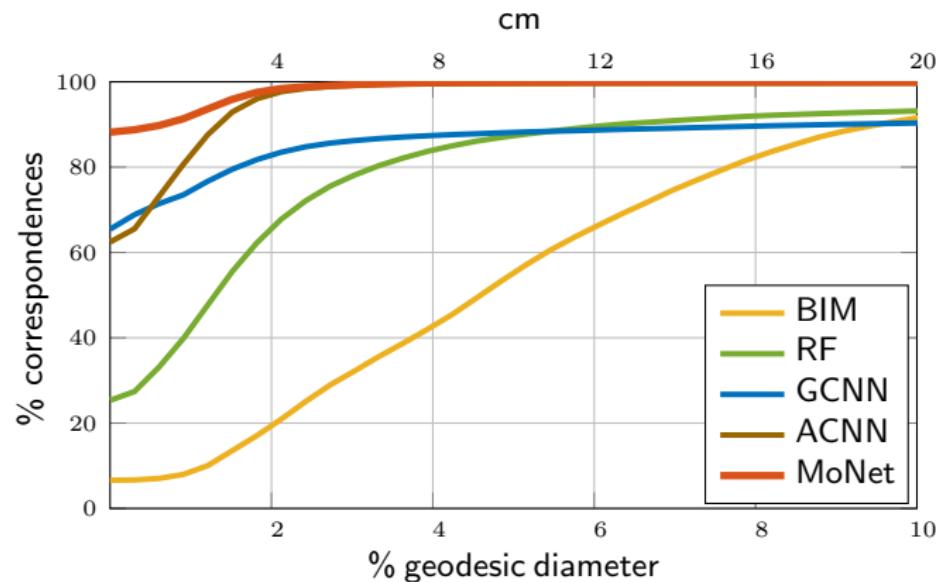
Correspondence evaluation: Princeton benchmark



Pointwise correspondence error = geodesic distance from the groundtruth

$$\epsilon(x) = d_{\mathcal{Y}}(\pi^*(x), \pi(x))$$

Correspondence quality comparison



Correspondence evaluated using asymmetric Princeton benchmark
(training and testing: disjoint subsets of FAUST)

Shape correspondence error: Blended Intrinsic Map



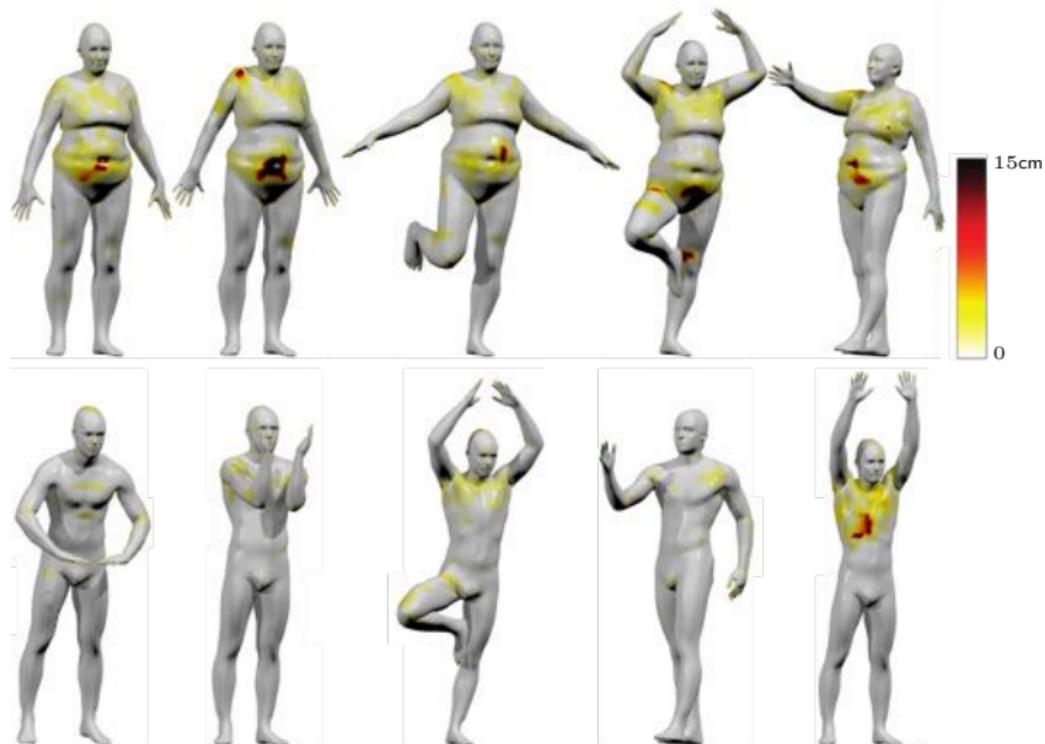
Pointwise correspondence error (geodesic distance from groundtruth)

Shape correspondence error: Geodesic CNN



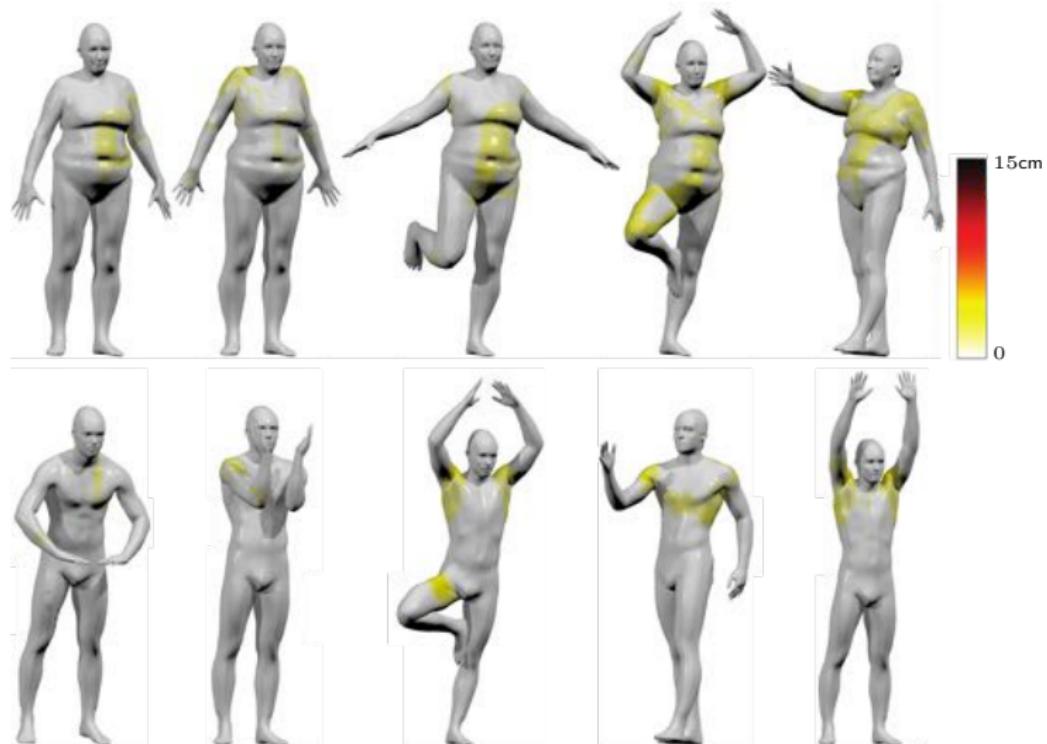
Pointwise correspondence error (geodesic distance from groundtruth)

Shape correspondence error: Anisotropic CNN



Pointwise correspondence error (geodesic distance from groundtruth)

Shape correspondence error: MoNet



Pointwise correspondence error (geodesic distance from groundtruth)

Shape correspondence visualization: MoNet



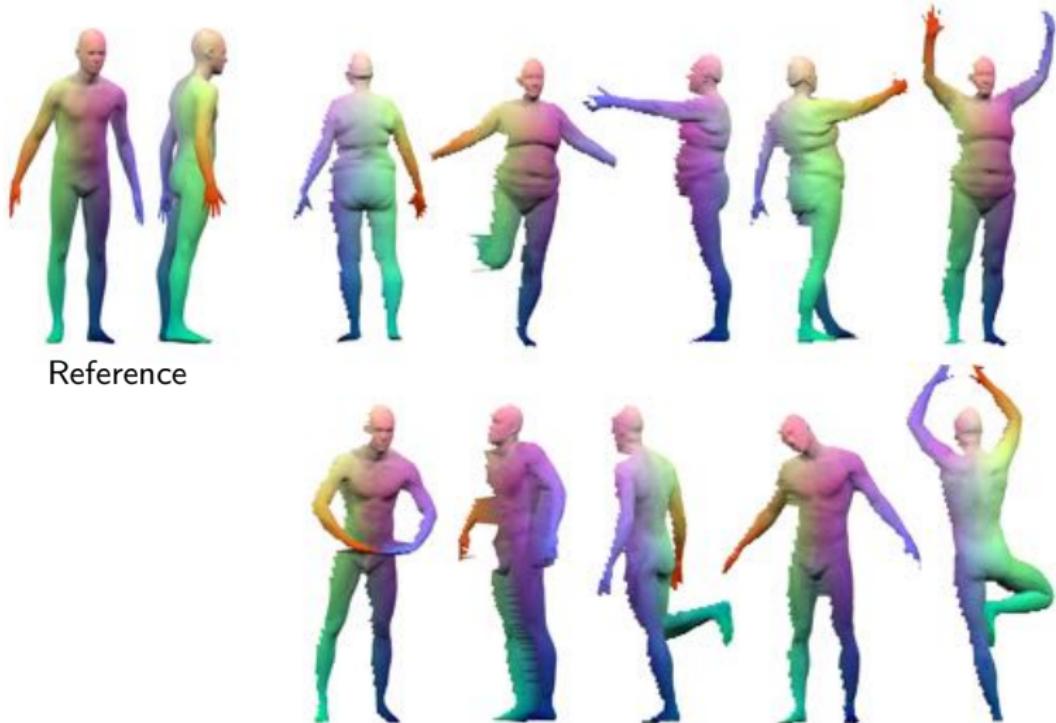
Texture transferred from reference to query shapes

Correspondence on range images: MoNet



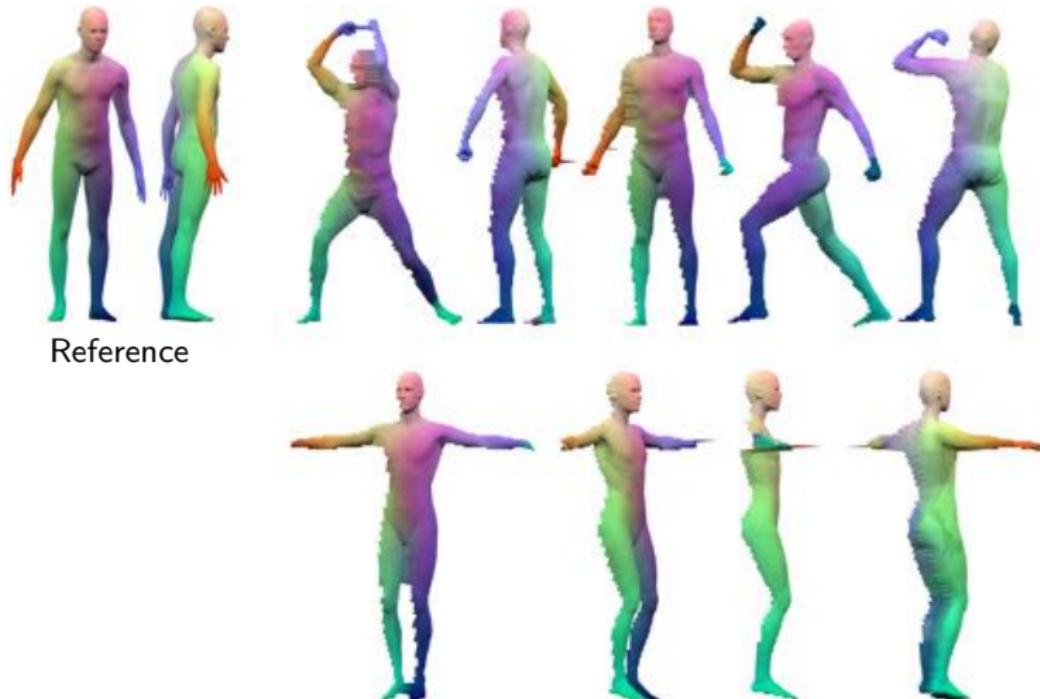
Pointwise correspondence error (geodesic distance from groundtruth)

Correspondence with MoNet: Range images



Correspondence visualization (similar colors encode corresponding points)

Correspondence with MoNet: Range images

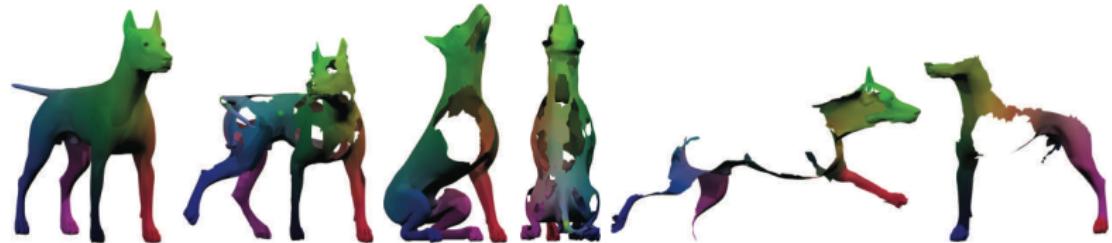


Correspondence visualization (similar colors encode corresponding points)

Partial correspondence with ACNN



Partial correspondence with ACNN

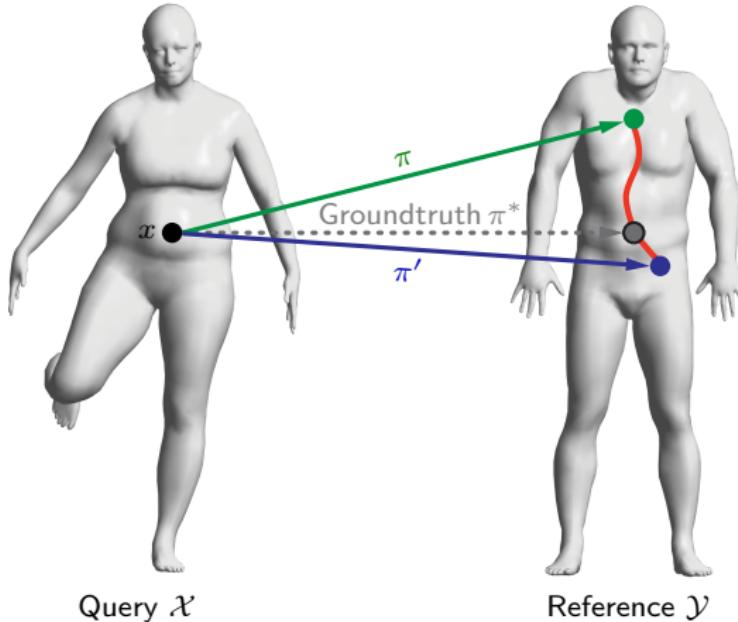


Correspondence



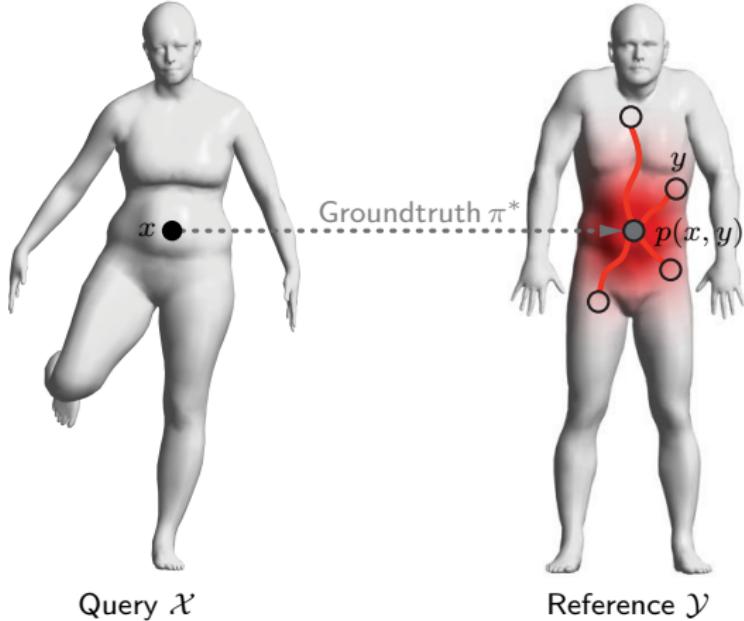
Correspondence error

Correspondence as classification problem, revisited



Classification cost considers equally correspondences
that deviate from the groundtruth (no matter how far)

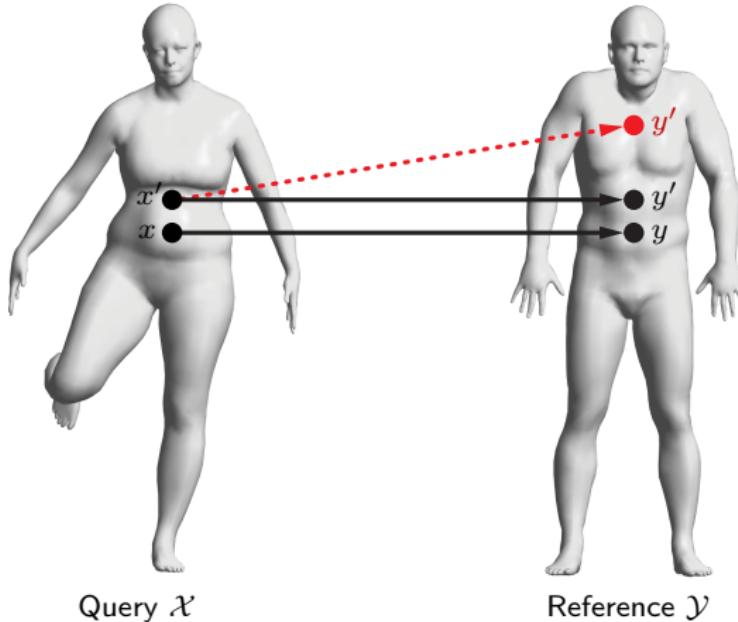
Soft correspondence error



Soft correspondence error = probability-weighted geodesic distance from the groundtruth

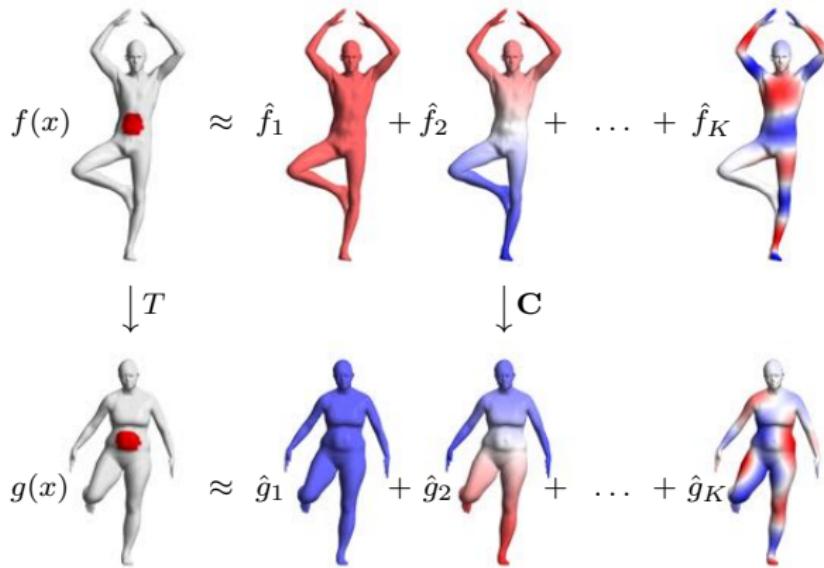
$$\bar{\epsilon}(x) = \int_{\mathcal{Y}} p(x, y) d_{\mathcal{Y}}(\pi^*(x), y) dy$$

Pointwise vs Structured learning



Nearby points x, x' on query shape are **not guaranteed** to map to nearby points y, y' on reference shape at **test time**

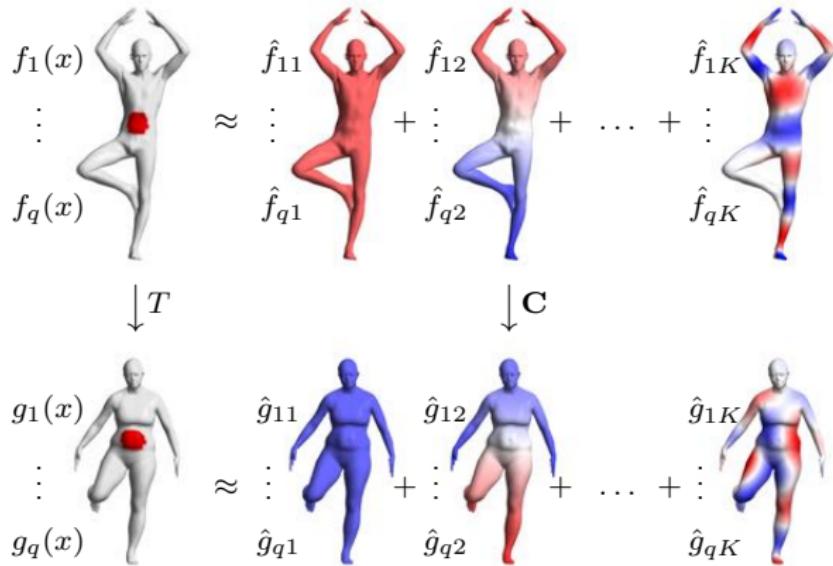
Functional maps: spectral domain



Functional correspondence $T =$ linear map \mathbf{C} between Fourier coefficients

$$\hat{\mathbf{g}}^\top = \hat{\mathbf{f}}^\top \mathbf{C}$$

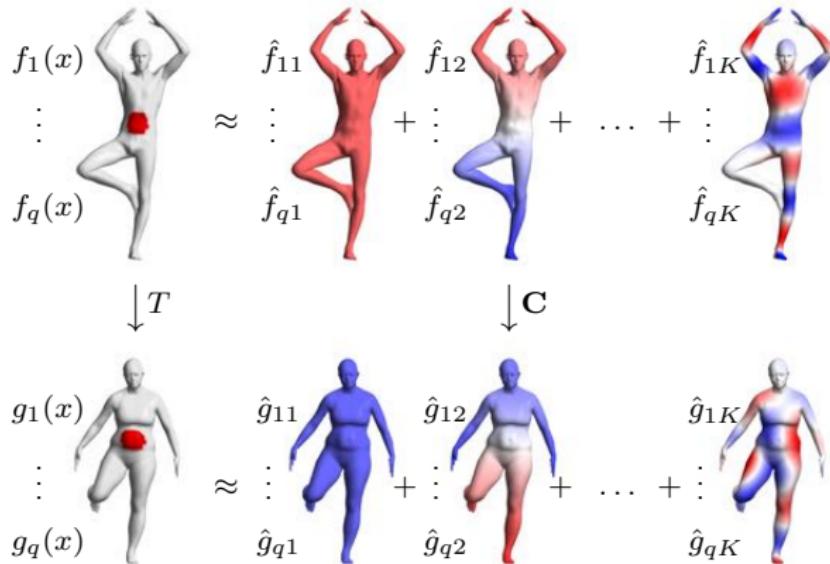
Functional maps: spectral domain



Recover correspondence from $q \geq k$ dimensional pointwise features

$$\begin{pmatrix} \hat{g}_{11} & \hat{g}_{12} & \dots & \hat{g}_{1K} \\ \vdots & \vdots & & \vdots \\ \hat{g}_{q1} & \hat{g}_{q2} & \dots & \hat{g}_{qK} \end{pmatrix} = \begin{pmatrix} \hat{f}_{11} & \hat{f}_{12} & \dots & \hat{f}_{1K} \\ \vdots & \vdots & & \vdots \\ \hat{f}_{q1} & \hat{f}_{q2} & \dots & \hat{f}_{qK} \end{pmatrix} \mathbf{C}$$

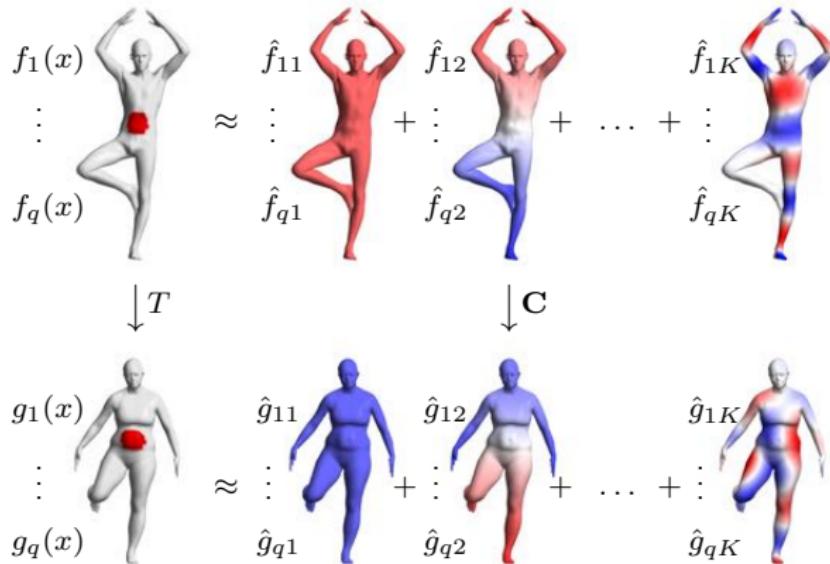
Functional maps: spectral domain



Recover correspondence from $q \geq k$ dimensional pointwise features

$$\hat{\mathbf{G}} = \hat{\mathbf{F}}\mathbf{C}$$

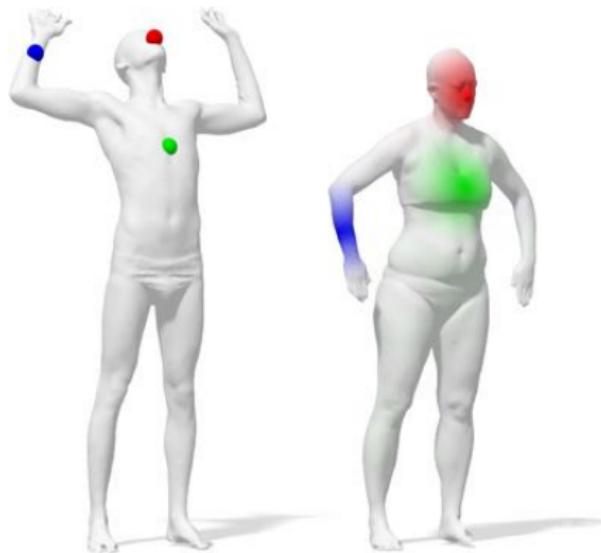
Functional maps: spectral domain



Recover correspondence from $q \geq k$ dimensional pointwise features

$$\mathbf{C}^* = \underset{\mathbf{C}}{\operatorname{argmin}} \|\hat{\mathbf{F}}\mathbf{C} - \hat{\mathbf{G}}\|_{\text{F}}^2$$

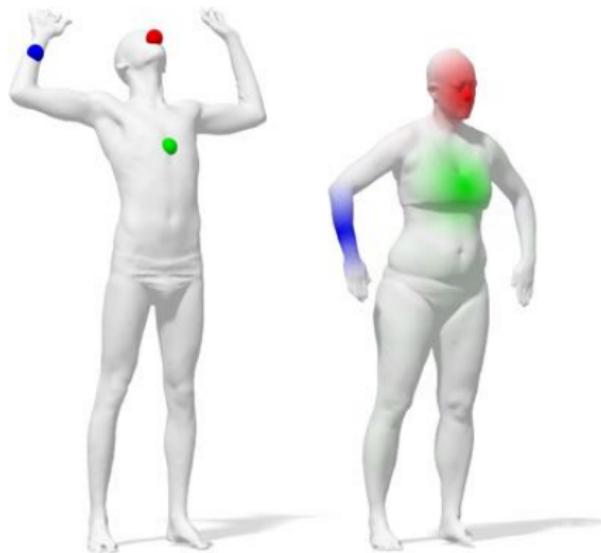
Functional maps: spatial domain



Rank- K approximation of spatial correspondence

$$\mathbf{T} \approx \boldsymbol{\Psi} \mathbf{C} \boldsymbol{\Phi}^\top$$

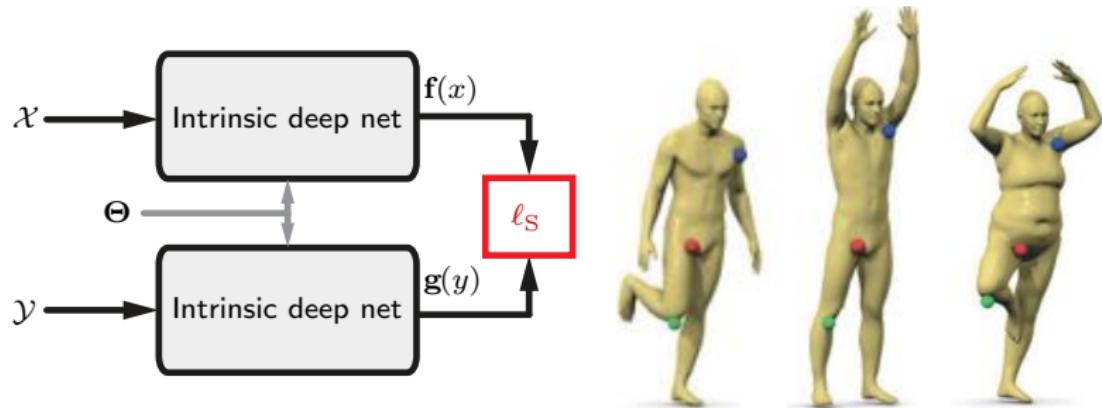
Functional maps: spatial domain



Probability $p(x, y)$ of point x mapping to y

$$\mathbf{P} \approx |\Psi \mathbf{C} \Phi^\top|_{\|\cdot\|}$$

Siamese metric learning



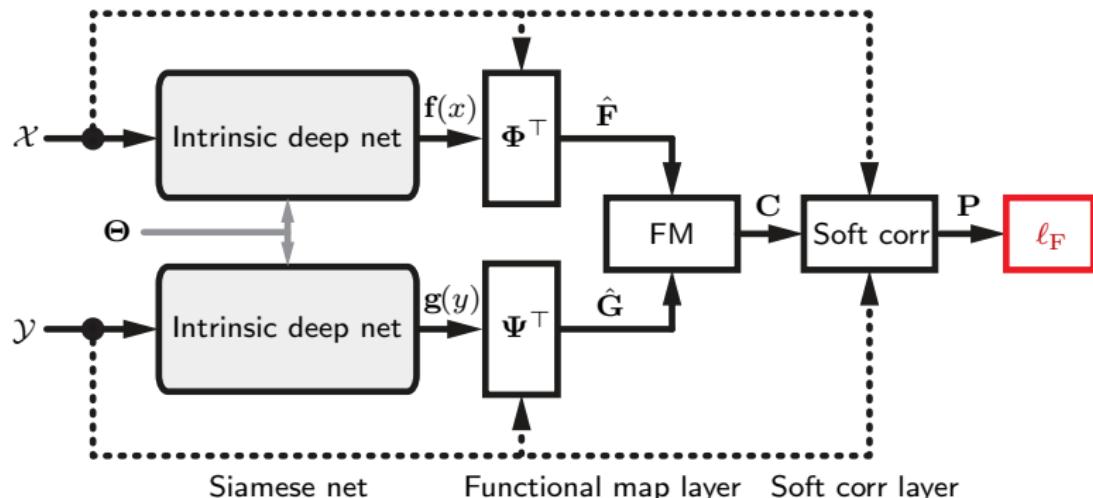
Siamese net

two net instances with shared parameters Θ

Poitwise feature cost

$$\begin{aligned}\ell_S(\Theta) = \gamma \sum_{x,x^+} & \| \mathbf{f}_\Theta(x) - \mathbf{f}_\Theta(x^+) \|_2^2 \\ & + (1 - \gamma) \sum_{x,x^-} [\mu - \| \mathbf{f}_\Theta(x) - \mathbf{f}_\Theta(x^-) \|_2^2]_+\end{aligned}$$

Structured correspondence with FMNet



Siamese net

two net instances with shared parameters Θ

Functional map layer

$$\mathbf{C}_\Theta^* = \underset{\mathbf{C}}{\operatorname{argmin}} \|\hat{\mathbf{F}}_\Theta \mathbf{C} - \hat{\mathbf{G}}_\Theta\|_F^2$$

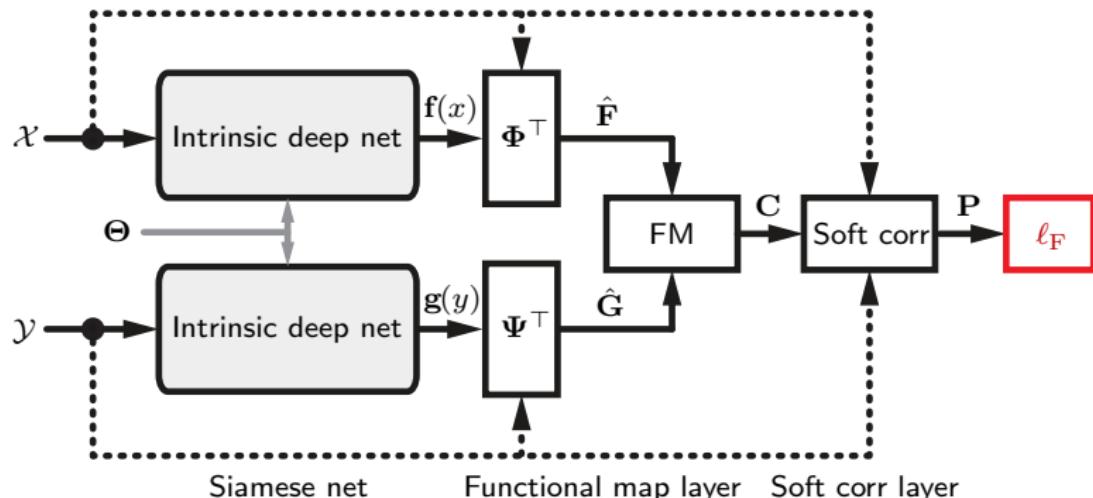
Soft correspondence layer

$$\mathbf{P}_\Theta = |\Psi \mathbf{C}_\Theta \Phi^\top|_{\|\cdot\|}$$

Soft error cost

$$\ell_F(\Theta) = \int_{\mathcal{X}} \int_{\mathcal{Y}} p_\Theta(x, y) d\mathcal{Y}(\pi^*(x), y) dx dy$$

Structured correspondence with FMNet



Siamese net

two net instances with shared parameters Θ

Functional map layer

$$\mathbf{C}_\Theta^* = \hat{\mathbf{F}}_\Theta^\dagger \hat{\mathbf{G}}_\Theta$$

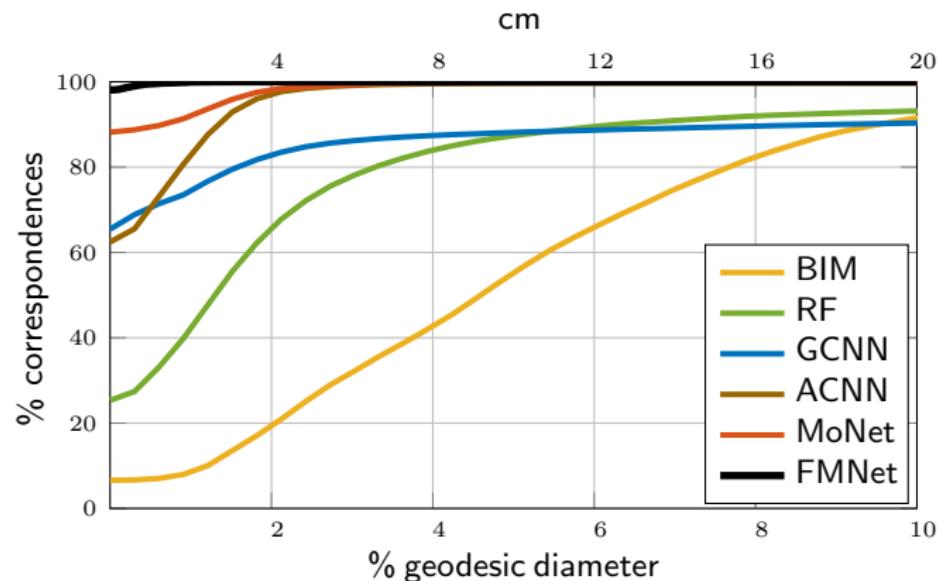
Soft correspondence layer

$$\mathbf{P}_\Theta = |\Psi \mathbf{C}_\Theta \Phi^\top|_{\|\cdot\|}$$

Soft error cost

$$\ell_F(\Theta) = \|\mathbf{P}_\Theta \circ \mathbf{D}_Y\|$$

Correspondence quality comparison



Correspondence evaluated using asymmetric Princeton benchmark
(training and testing: disjoint subsets of FAUST)

Methods: Kim et al. 2011 (BIM); Rodolà et al. 2014 (RF); Boscaini et al. 2015 (ADD); Masci et al. 2015 (GCNN);
Boscaini et al. 2016 (ACNN); Monti et al. 2016 (MoNet); Litany et al. 2017 (FMNet); data: Bogo et al. 2014 (FAUST);
benchmark: Kim et al. 2011

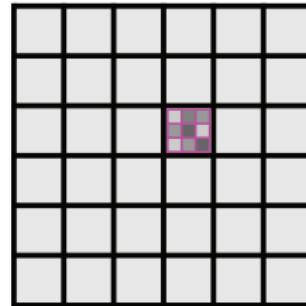
Different formulations of non-Euclidean CNNs



Spectral domain



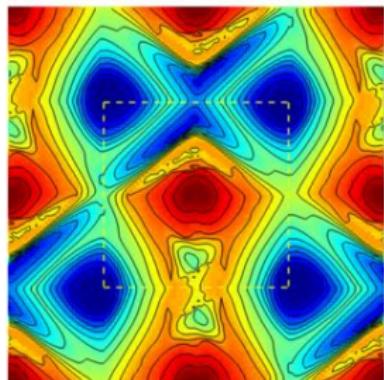
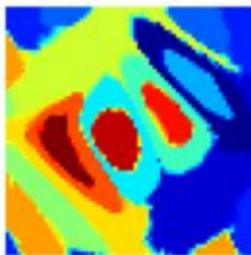
Spatial domain



Embedding domain

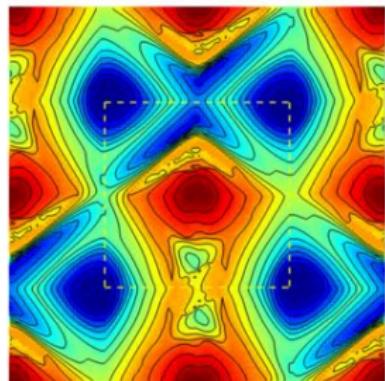
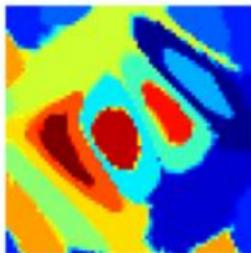
Global parametrization

Map the input surface to some [parametric domain](#) with shift-invariant structure



Global parametrization

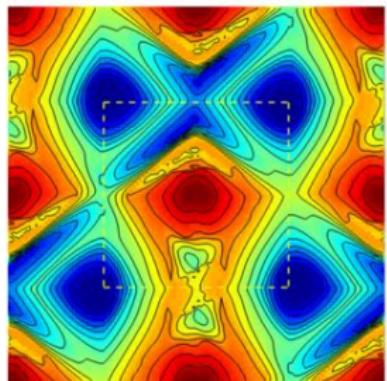
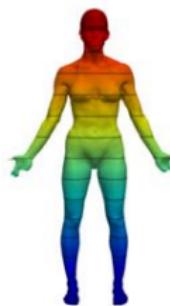
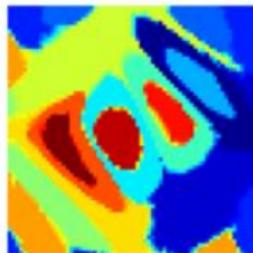
Map the input surface to some **parametric domain** with shift-invariant structure



- ☺ Allows to use standard CNNs (pull back convolution from the parametric space)

Global parametrization

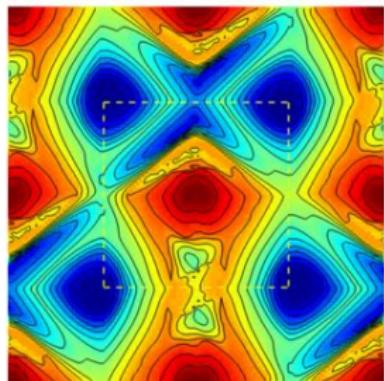
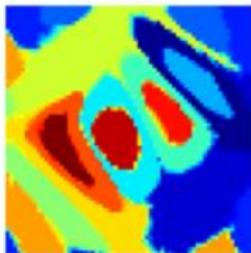
Map the input surface to some **parametric domain** with shift-invariant structure



- ☺ Allows to use standard CNNs (pull back convolution from the parametric space)
- ☺ Guaranteed invariance to some classes of transformations

Global parametrization

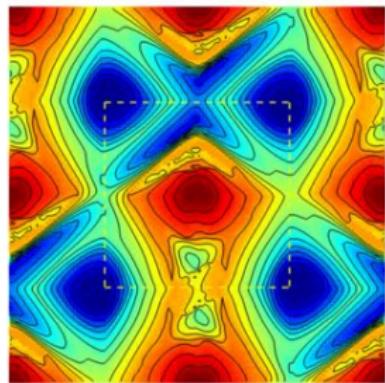
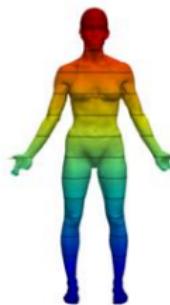
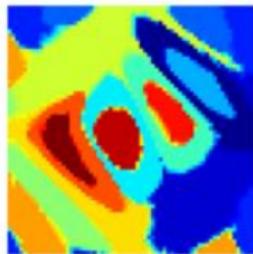
Map the input surface to some **parametric domain** with shift-invariant structure



- ☺ Allows to use standard CNNs (pull back convolution from the parametric space)
- ☺ Guaranteed invariance to some classes of transformations
- ☹ Parametrization may not be unique

Global parametrization

Map the input surface to some **parametric domain** with shift-invariant structure



- ☺ Allows to use standard CNNs (pull back convolution from the parametric space)
- ☺ Guaranteed invariance to some classes of transformations
- ☹ Parametrization may not be unique
- ☹ Embedding may introduce distortion

Translation invariance on manifolds

Translation on manifold = locally Euclidean translation

Translation invariance on manifolds

Translation on manifold = locally Euclidean translation = flow along a non-vanishing vector field

Translation invariance on manifolds

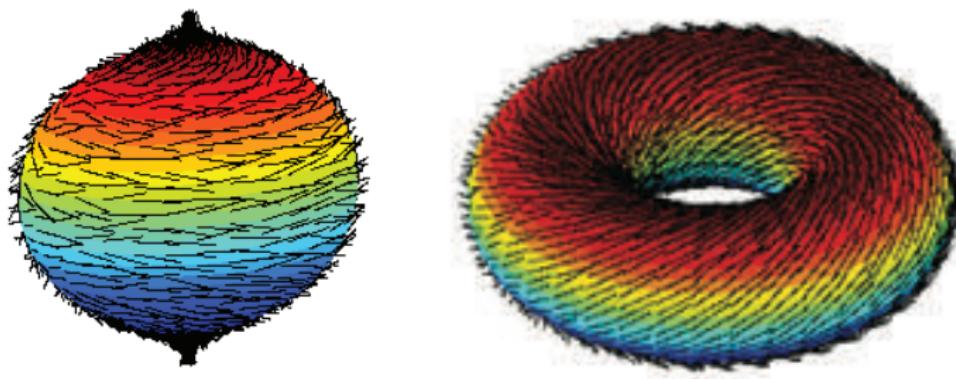
Translation on manifold = locally Euclidean translation = flow along a non-vanishing vector field

Poincaré-Hopf Theorem Non-vanishing vector field on a closed orientable compact 2-manifold implies manifold of genus 1 (torus)

Translation invariance on manifolds

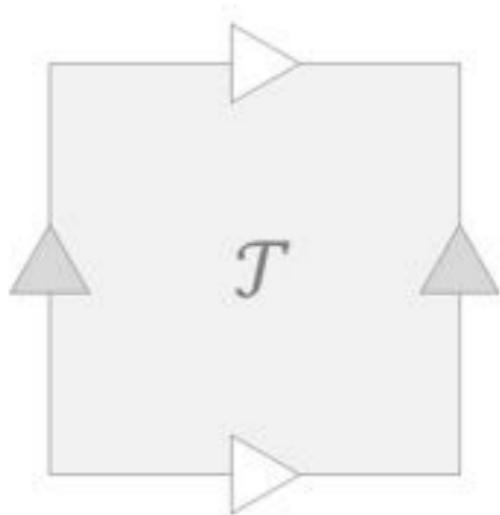
Translation on manifold = locally Euclidean translation = flow along a non-vanishing vector field

Poincaré-Hopf Theorem Non-vanishing vector field on a closed orientable compact 2-manifold implies manifold of genus 1 (torus)



'Hairy Ball Theorem' states that a sphere cannot be combed

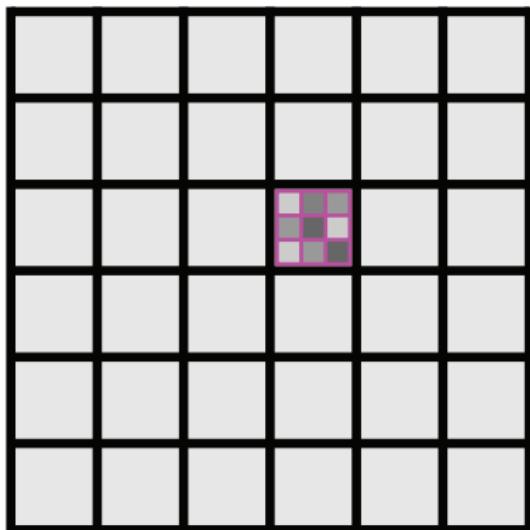
Translation invariance on the torus



Torus is the only closed orientable surface admitting a translation group

Convolution on torus

For any triplet of points on \mathcal{X} , construct [conformal homeomorphism](#) from the 4-cover \mathcal{X}^4 to \mathcal{T} using [orbifold-Tutte](#) method



Conformal zoom

- Embedding depends on the choice of the triplets of points

Conformal zoom

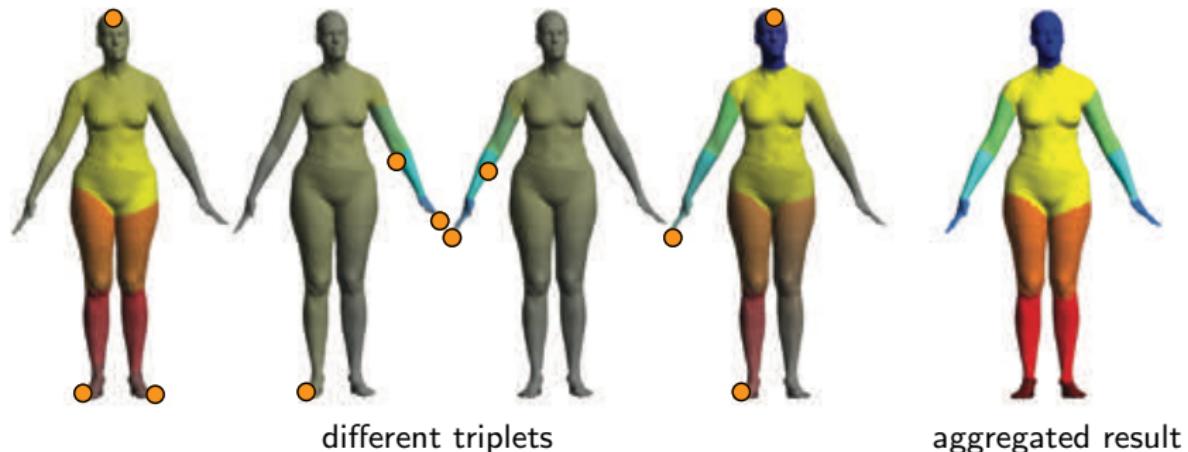
- Embedding depends on the choice of the triplets of points
- ‘Conformal zoom’ effect

Conformal zoom

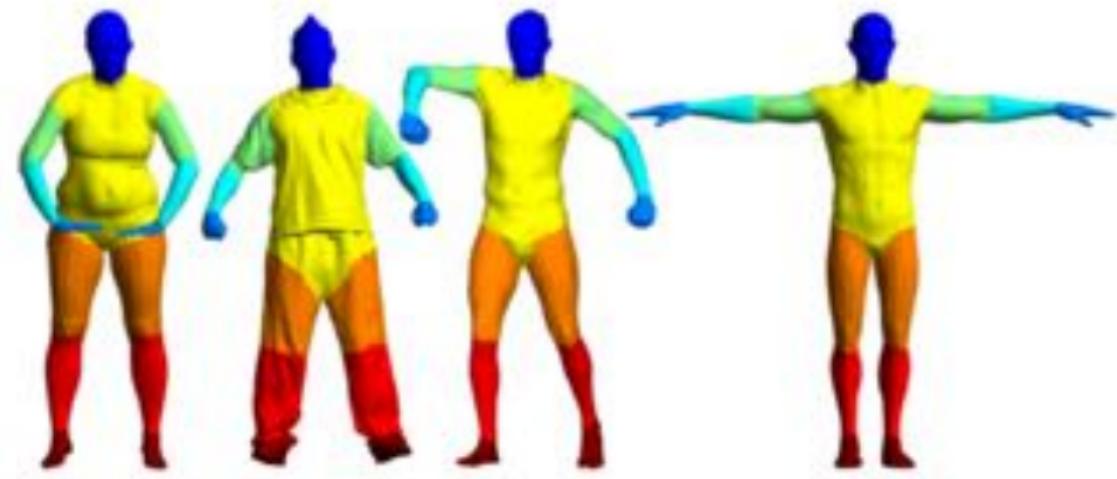
- Embedding depends on the choice of the triplets of points
- ‘Conformal zoom’ effect
- Choose multiple triples and aggregate results in training / test phase

Conformal zoom

- Embedding depends on the choice of the triplets of points
- ‘Conformal zoom’ effect
- Choose multiple triples and aggregate results in training / test phase

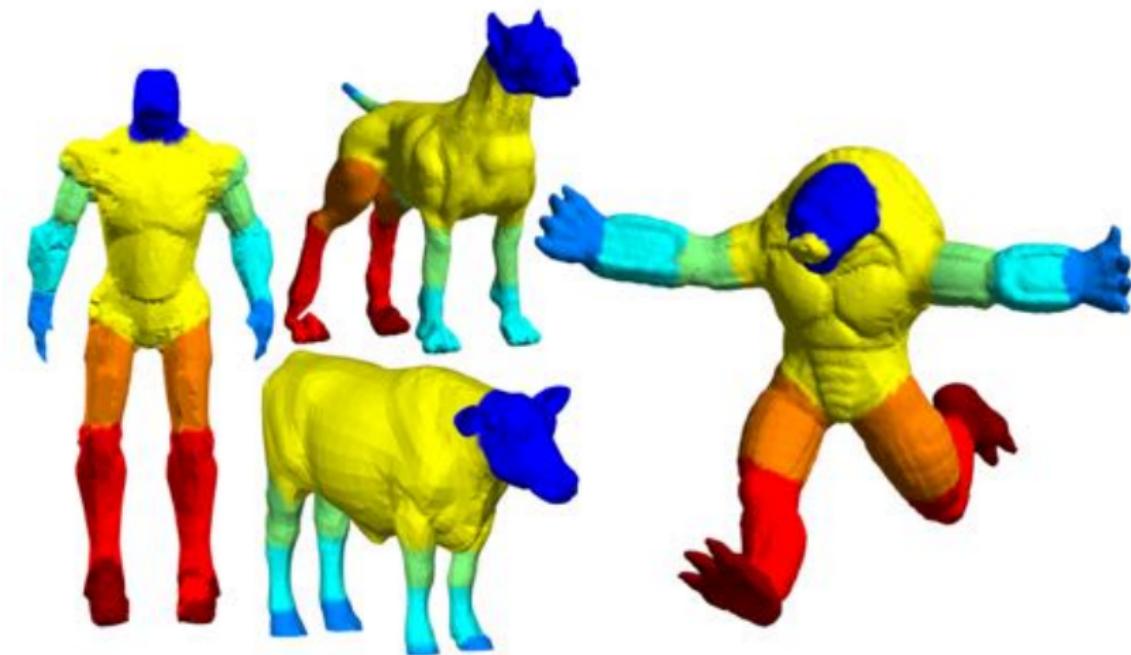


Example: shape segmentation with Toric CNN



Examples of shape segmentation obtained with Toric CNN

Example: shape segmentation with Toric CNN



Examples of shape segmentation obtained with Toric CNN